Universidad De Manila

Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic

A Capstone Project Presented to the College of
Computing Studies
Universidad de Manila

In Partial Fulfillment
of the Requirements for the Course
CAPSTONE 2

Brayan Axel L. Tubillas
Keith Ashly V. Sales
Reizen S. Arreola
Robin C. Ponce
Rizelle Carla N. Magpantay

Prof. Francis Arlando L. Atienza
Professor

October 2026

ABSTRACT

Dental screening at CodiSmile Dental Clinic relies on visual inspection, which makes early tartar, plaque, cavity, and gum inflammation hard to detect consistently, leaves no standardized visual record, and gives patients little finding-specific guidance. This study developed Toothwise, a web-based screening support system that detects these four conditions and healthy tooth in still intraoral photographs with a fused YOLOv8 and YOLOv12 ensemble, computes a 0–100 risk score and a low, medium, or high risk tier, and generates a personalized prevention plan with a K-Nearest Neighbors (KNN) classifier. The system was developed with an Agile methodology using React, FastAPI, and PostgreSQL, and deployed on cloud hosting. On 2,081 duplicate-free test images, YOLOv8, YOLOv12, and their ensemble each exceeded 80% on all six image-level measures for any dental condition and for cavity; box-level mAP@50 was about 50%. The KNN classifier reproduced the clinical risk rule with 92.8% accuracy on 2,000 unseen cases, and all 118 automated backend tests passed. The system was evaluated through two online forms answered on a four-point scale. The six ISO/IEC 25010 quality characteristics (security, privacy, usability, reliability, performance efficiency, and data protection), rated by 2 IT experts, obtained an overall weighted mean of 3.97, and the system's detection, results, and prevention-plan features, rated by 103 user respondents (98 patient respondents and 5 IT practitioners), obtained 3.48; both are interpreted as Highly Acceptable. The study concludes that Toothwise is an acceptable screening and documentation aid that supports, but does not replace, the dentist's diagnosis, and recommends more training data for weak classes such as plaque, the use of an intraoral camera for clearer images, a larger evaluation with dentists and clinic staff, and a controlled comparison of screening time and consistency.

Keywords: dental screening, YOLOv8, YOLOv12, ensemble object detection, K-Nearest Neighbors, prevention plan, ISO/IEC 25010

TABLE OF CONTENTS

CHAPTER I: INTRODUCTION

1.1 Background of the Study....9
1.2 Statement of the Problem....12
1.3 Objectives of the Study....14
1.3.1 General Objective....14
1.3.2 Specific Objectives....14
1.4 Scope, Delimitation and Limitation of the Study....15
1.5 Significance of the Study....17
1.6 Definition of Terms....18

CHAPTER II: REVIEW OF RELATED LITERATURE AND STUDIES

2.1 Related Literature....22
2.2 Related Studies....26
2.3 Synthesis....30

CHAPTER III: METHODOLOGY

3.1 Development Methodology....32
3.2 Project Planning....33
3.3 Roles and Responsibilities of the Five-Member Development Team....36
3.4 Requirements Analysis....37
3.5 System Design....39
3.5.1 Use Case Diagram....39
3.5.2 Entity Relationship Diagram (ERD)....41
3.5.3 Data Flow Diagram Level 0....42
3.5.4 Data Flow Diagram Level 1....44
3.5.5 Gantt Chart....48
3.6 Software Specifications....49
3.7 Hardware Specifications....52
3.8 System Development....53
3.9 System Testing....55
3.10 Evaluation Procedure....56

3.11 Full Current System Screenshots.....59

CHAPTER IV: RESULTS AND DISCUSSION

4.1 Results of the Trained Datasets.....70

4.2 Specialization and Technical Results.....81

4.3 Software Quality Evaluation Based on ISO/IEC 25010.....85

4.4 Speed and Reliability of the System.....94

CHAPTER V: SUMMARY OF FINDINGS, CONCLUSIONS, AND RECOMMENDATIONS

5.1 Summary of Findings.....99

5.2 Conclusions.....101

5.3 Recommendations.....102

References.....105

Appendix A. Letter of Approval from CodiSmile Dental Clinic.....109

Appendix B. Detailed Image-Level Results per Model.....111

Appendix C. Evaluator Feedback and Panel Consultation.....112

Appendix D. IT Expert Evaluation Results.....115

LIST OF TABLES

Table 1 Sources of Requirements and Technical Data....35
Table 2 Roles and Responsibilities of the Development Team....36
Table 3 Functional Requirements....37
Table 4 Non-Functional Requirements....38
Table 5 Toothwise Database Schema Data Dictionary....41
Table 6 Level 1 Process Descriptions....45
Table 7 Software Development Requirements....49
Table 8 Implementation Software Requirements....51
Table 9 Hardware Development Requirements....52
Table 10 Implementation Hardware Requirements....52
Table 11 Development Phase Activities....53
Table 12 Agile Iteration Summary....54
Table 13 Sections of the Evaluation Instruments....57
Table 14 Four-Point Likert Scale and Verbal Interpretation....58
Table 15 Source Datasets and Class Remapping to the Toothwise Taxonomy....70
Table 16 Merged Dataset Split....71
Table 17 Test Set Before and After Duplicate Removal....72
Table 18 Training Configuration of the Two Detection Models....72
Table 19 Image-Level Performance for Any Finding (Clean Test Set, Validation-Tuned Cut-offs)....75
Table 20 Image-Level Performance at the Deployed Cut-offs (Clean Test Set)....75
Table 21 Image-Level Performance per Condition, Fused Ensemble (Clean Test Set)....76
Table 22 Image-Level F1-Score per Condition for Each Model (Clean Test Set)....76
Table 23 Image-Level True and False Positives and Negatives (Clean Test Set)....77
Table 24 Box-Level (Standard YOLO) Performance of YOLOv8 and YOLOv12....78
Table 25 Inference Time per Photograph on CPU....79
Table 26 Risk Scoring Weights and Tier Thresholds....82
Table 27 KNN Tier Classification on Unseen Synthetic Cases....82
Table 28 Profile of the IT Expert Evaluators....85
Table 29 Distribution of User Respondents....86
Table 30 User Respondents per Survey Part....86
Table 31 Evaluation Results for Security....87
Table 32 Evaluation Results for Privacy....87
Table 33 Evaluation Results for Usability....88
Table 34 Evaluation Results for Reliability....89
Table 35 Evaluation Results for Performance Efficiency....89

Table 36 Evaluation Results for Data Protection....90
Table 37 Evaluation Results for AI Detection and Risk Classification....90
Table 38 Evaluation Results for Results Interface and Patient Records....91
Table 39 Evaluation Results for KNN-Based Personalized Prevention Plan....92
Table 40 Summary of Evaluation Results....93
Table 41 Automated Backend Test Results....95
Table 42 Lighthouse Audit Results....96
Table 43 Respondents' Assessment of the Current Manual Screening Process....97
Table 44 Respondents' Assessment of Screening Consistency and Time with Toothwise....97
Table 45 Image-Level Performance per Condition, YOLOv8 (Clean Test Set)....111
Table 46 Image-Level Performance per Condition, YOLOv12 (Clean Test Set)....111
Table 47 Frequency Distribution of the IT Experts' Ratings....115

LIST OF FIGURES

Figure 1 Agile Development Methodology....33
Figure 2 Use Case Diagram....40
Figure 3 Entity Relationship Diagram (ERD)....43
Figure 4 Data Flow Diagram Level 0....44
Figure 5 Data Flow Diagram Level 1 (Sheet 1 of 2: Processes 1.0, 2.0, 3.0, and 7.0)....46
Figure 6 Data Flow Diagram Level 1 (Sheet 2 of 2: Processes 4.0, 5.0, and 6.0)....47
Figure 7 Toothwise Development Gantt Chart....48
Figure 8 Public Landing Page....59
Figure 9 Data Privacy Notice....60
Figure 10 Staff Login Page....61
Figure 11 Patient Portal — Sign In....61
Figure 12 Reset Password (Staff)....61
Figure 13 Reset Password (Patient)....62
Figure 14 Change Password....62
Figure 15 Patient Portal — Sign Up....62
Figure 16 Dashboard....63
Figure 17 Patient List....64
Figure 18 Patient Details....64
Figure 19 New Screening (Step 1: Choose Image)....65
Figure 20 New Screening (Step 2: Review and Analyse)....65
Figure 21 Screening Results....66
Figure 22 Detected Conditions and Image-Level Evaluation....66
Figure 23 Staff Management (Dentist/Administrator)....67
Figure 24 Audit Log (Dentist/Administrator)....67
Figure 25 Patient Portal — Home....68
Figure 26 Patient Portal — Screening Details....69
Figure 27 PDF Screening Report....69
Figure 28 mAP@50 and mAP@50-95 per Training Epoch for YOLOv8 and YOLOv12....73
Figure 29 Letter Requesting Approval to Conduct the Study (page 1 of 2)....110
Figure 30 Letter Requesting Approval to Conduct the Study (page 2 of 2)....110
Figure 31 IT Expert 1 — Demonstration and Form Sent....112
Figure 32 IT Expert 1 — Feedback....112
Figure 33 IT Expert 2 — Form Submitted....113
Figure 34 IT Expert 2 — Feedback....113
Figure 35 Panel Consultation — Request for Review....114
Figure 36 Panel Consultation — Approval....114

Figure 37 IT Expert Form Results — Evaluator Profile.....116
Figure 38 IT Expert Form Results — Security, Privacy, and Usability.....116
Figure 39 IT Expert Form Results — Reliability, Performance Efficiency, and Data Protection 117
Figure 40 IT Expert Form Results — Comments and Recommendations.....117

CHAPTER I
INTRODUCTION

1.1 Background of the Study

CodiSmile Dental Clinic currently relies on manual visual examination and professional judgment to identify dental caries and other visible dental conditions from patient examinations and intraoral images. Although this approach is essential to dental practice, manually reviewing images may be time-consuming and can result in variations in assessment and documentation. To support the existing process, the proponents propose Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic, an AI-assisted system that analyzes intraoral images and detects possible dental conditions — tartar, plaque, cavity, gum inflammation, and healthy tooth — using the YOLOv12 and YOLOv8 object detection models. The system provides visual detection results, condition labels, descriptions, a computed risk classification, and preventive recommendations to assist dentists and improve patient understanding. Toothwise is intended to serve as a supplementary screening and documentation tool, while final interpretation and clinical decisions remain the responsibility of licensed dental professionals.

Dental caries, commonly referred to as tooth decay or cavities, remains one of the most prevalent non-communicable diseases in the world. The World Health Organization (2025) estimates that oral diseases affect nearly 3.7 billion people globally, with untreated dental caries in permanent teeth identified as the single most widespread health condition based on Global Burden of Disease 2021 data. Despite its prevalence, a significant portion of the population, particularly those in geographically isolated or economically disadvantaged communities, continues to face substantial barriers in accessing timely and adequate dental care. In the Philippines, Mendoza et al. (2025) note that an estimated 29 million Filipinos have untreated dental caries and that outpatient dental services were not covered by the PhilHealth benefit package at the time of their analysis. The financial cost

of consultations, the distance to dental facilities, and low oral health literacy frequently result in delayed care-seeking, by which point cavities have often advanced beyond their most treatable stages. Untreated caries does not just cause physical pain; Almajed et al. (2023) found that children with untreated caries had higher rates of poor oral-health-related quality of life.

A critical dimension of the current dental care challenge lies in the inconsistency of manual cavity screening practices. Wang et al. (2026), in a review on the clinical translation of AI-based caries diagnosis, noted that traditional diagnostic methods suffer from limited sensitivity and high variability between examiners, particularly for early-stage lesions. This inter-examiner variability is particularly problematic in high-volume clinic settings where practitioner fatigue and time pressure can further compromise screening reliability. Research has shown that deep learning models are capable of detecting dental cavities with accuracy that is on par with, or sometimes even better than, trained clinicians, as reported in the reviews by Schwendicke et al. (2019) and Mohammad-Rahimi et al. (2022). This makes AI a genuinely useful tool for supplementing human judgment during screenings.

Standard digital cameras and smartphone cameras have also changed the way dental screening can be supported outside specialized clinical equipment. Instead of relying purely on what a dentist can see with a mirror and a light, or requiring dedicated intraoral camera hardware, clear intraoral photographs can now be captured using widely available consumer devices. Park et al. (2022), for example, improved caries classification on intraoral camera photographs by first segmenting the tooth surfaces, and studies such as those by Moharrami et al. (2025) and Askar et al. (2021) show that RGB photographs captured under ordinary imaging conditions are detailed enough to support reliable artificial intelligence-based detection of dental conditions. It is this still-image, standard-camera modality that serves as the primary data source for Toothwise during patient screenings at CodiSmile Dental Clinic, since the system is designed to work from

photographs taken with an ordinary device camera, such as a smartphone, and uploaded to the system, rather than requiring a dedicated intraoral camera.

The AI model powering the primary detection module of this system is YOLOv12, which was introduced by Tian et al. (2025) as an attention-centric object detection framework. Compared to previous versions in the YOLO family, YOLOv12 uses area attention mechanisms that help it focus selectively on spatially relevant regions, something that happens regularly in dental images where carious lesions appear as small, low-contrast areas across multiple tooth surfaces simultaneously. Its combination of speed and accuracy makes it well-suited for efficient use during clinical consultations, where screening results need to be generated quickly enough to fit within a standard patient visit.

To evaluate whether YOLOv12's architectural improvements translate into meaningful gains for this specific screening task, this study also trains and evaluates YOLOv8, an earlier and widely adopted model in the YOLO family (Jocher et al., 2023). YOLOv8 has been applied in several prior dental imaging studies with strong results; de Brito et al. (2025) used YOLOv8 together with EfficientNet-B0 to diagnose caries in interproximal radiographic images of young permanent molars, and other work has applied YOLOv8's anchor-free detection mechanism to interproximal caries detection with improved accuracy and speed over earlier YOLO versions. This makes it a stable, well-benchmarked baseline for object detection tasks. While YOLOv12 and YOLOv8 are formally compared against each other using precision, recall, mAP, and inference speed on a held-out labeled test set, both models are also deployed together in the operational system, where their independent detection outputs are fused into a single clinical detection set. This dual approach allows the study to identify which architecture performs better on this specific task while still taking advantage of both models' detections in the deployed screening workflow.

Beyond detection, Toothwise also generates a personalized prevention plan for each patient based on their screening results. This is done using the K-Nearest Neighbors

(KNN) algorithm, a supervised learning method originally introduced by Cover and Hart (1967), which classifies a new case by comparing it against the most similar cases in a reference dataset. In this study, KNN classifies each screening into a risk tier (low, medium, or high) based on the patient's detected findings — condition type, number of affected areas, and severity indicators derived from the detection output — and this risk classification, together with condition-specific rules, is used to compose an ordered, patient-friendly prevention plan, while the 0–100 risk score itself comes from a separate rule-based formula. Since no dentist-labeled cases exist before deployment, the KNN is first trained on synthetic cases derived from the clinical severity rule, and each tier the dentist confirms or corrects is added as a real case for retraining.

While AI-assisted dental screening has grown internationally, its application in the Philippine setting remains an emerging area of research. Wu et al. (2025), in collaboration with Ateneo de Manila University's Laboratory for Intelligent Visual Environments, applied a deep learning model to identify tooth and sinus structures in dental panoramic radiographs, illustrating the early but growing interest in AI-assisted dental diagnostics locally. This study seeks to build on that emerging interest by applying AI-assisted screening directly within a clinic setting through Toothwise.

This project was created in partnership with CodiSmile Dental Clinic, which serves as the deployment site and clinical reference point throughout the study. The goal is not to take the place of the dentist but to give them a more efficient and objective tool to work with. With Toothwise integrated into the clinic's workflow, screenings can be documented more consistently, and patients can walk away with a clearer understanding of their oral health through the system's risk classification and personalized prevention plan features.

1.2 Statement of the Problem

Dental screening at CodiSmile Dental Clinic relies on visual inspection. Early-stage tartar, plaque, cavity, and gum inflammation are small, look similar to one another,

and vary with lighting, angle, and stage, so visual assessment is subjective and findings can differ between examinations. The clinic has no standardized visual record of findings to compare across visits, and patients are rarely given finding-specific preventive guidance. Undetected conditions can progress to pain, tooth loss, and costlier treatment, which makes consistent early screening a challenge of the method itself. Any tool introduced to address this must also be fast and dependable enough not to slow down the clinic's daily work.

Specifically, this study addresses the following problems:

1. Early-stage tartar, plaque, cavity, and gum inflammation, as well as healthy teeth, are difficult to detect and distinguish consistently through visual-only screening, and the clinic has no objective, measured means of identifying these conditions in intraoral photographs or of estimating a patient's risk level from them.
2. Screening findings are not kept in a standardized visual record, so dentists cannot readily compare findings across visits, and patients have no way to review their own screening results and history.
3. Patients rarely receive timely, finding-specific preventive guidance after screening; the advice they receive is often generic rather than based on their own findings and risk level.
4. The acceptability of Toothwise in terms of the six ISO/IEC 25010 quality characteristics evaluated in this study (security, privacy, usability, reliability, performance efficiency, and data protection) has not yet been established among IT experts and patient respondents.
5. It has not yet been established whether an AI-assisted screening system such as Toothwise can complete a screening quickly and run reliably enough for daily use in a clinic.

1.3 Objectives of the Study

1.3.1 General Objective:

To design, develop, and evaluate Toothwise, a dental screening support system that uses YOLO-based detection on still intraoral photographs, a KNN-based risk classification and recommendation feature, to provide consistent and accessible screening support for licensed dentists and patients at CodiSmile Dental Clinic.

1.3.2 Specific Objectives:

1. To train YOLOv12 and YOLOv8 on annotated still intraoral images to detect tartar, plaque, cavity, gum inflammation, and healthy tooth, compute an associated risk score and risk classification (low, medium, high) from the detected findings, and evaluate the two models and their fused ensemble using image-level accuracy, precision, recall, specificity, F1-score, and area under the ROC curve (AUROC) as the primary measures and the standard YOLO measures (precision, recall, mean average precision, and inference speed) as secondary measures.
2. To develop a results interface that saves every screening as a standardized visual record showing bounding boxes, condition labels, plain-language descriptions, and risk classification, with a per-patient record of findings across visits accessible to dentists and, through a self-service patient portal, to patients themselves.
3. To develop a KNN-based feature that generates a personalized, finding-specific prevention plan from the detected findings and computed risk classification after every screening.
4. To evaluate Toothwise using six ISO/IEC 25010 quality characteristics with IT experts and patient respondents.

5. To evaluate the speed and reliability of Toothwise in terms of screening processing time, automated functional testing, and website performance and accessibility.

1.4 Scope, Delimitation and Limitation of the Study

Scope

This study covers the design, development, training, and evaluation of Toothwise, an AI-assisted dental screening support system for CodiSmile Dental Clinic. Using still photographs taken with a standard device camera, such as a smartphone, and uploaded to the system as HEIC/HEIF, JPEG, PNG, or WebP files, Toothwise analyzes the patient's teeth and gums to detect five classes: tartar, plaque, cavity, gum inflammation, and healthy tooth. YOLOv12 and YOLOv8 models are trained on publicly available annotated dental image datasets and evaluated in terms of image-level accuracy, precision, recall, specificity, F1-score, and AUROC as the primary measures, and the standard YOLO measures (precision, recall, mean average precision (mAP), and inference speed) as secondary measures, on a held-out test set from which photographs duplicated in the training and validation sets were removed; both models are also used together in the deployed system, with their outputs fused into a single clinical detection set. After training, the system is tested at CodiSmile Dental Clinic to assess how it performs in an actual clinical environment. Detection results are displayed through a results interface that shows bounding boxes, condition labels, plain-language descriptions, and a computed risk classification for licensed dentists, and findings are saved in a per-patient record that can be reviewed across visits. Based on the detected findings, a recommendation feature that uses the K-Nearest Neighbors (KNN) algorithm classifies the screening into a risk tier and generates a personalized prevention plan for each patient. The system also generates a downloadable PDF screening report for clinical recordkeeping and maintains an audit log of sensitive staff actions for administrative review. Toothwise is a browser-based web application; for the study it is deployed on cloud hosting (the React frontend on Vercel,

the FastAPI backend on Render, and the PostgreSQL database on Neon), so clinic staff and patients can open it from any device with a modern web browser and an Internet connection. The intended users are licensed dentists and clinic personnel at CodiSmile Dental Clinic; in addition, patients may register for a self-service portal to view their own screening history and past findings. The system is evaluated by IT experts and patient respondents through an online questionnaire based on six ISO/IEC 25010 quality characteristics, and its screening speed, reliability under automated testing, and website performance and accessibility are measured.

Delimitations

Toothwise is a support tool and not a replacement for a licensed dentist. The dentist at CodiSmile Dental Clinic always has the final say on any diagnosis or treatment decision. The prevention plan is meant to guide the patient and is not a formal medical prescription, so it does not include medication, treatment procedures, or dosages. The system works only with still photographs that are uploaded as image files (for example, a photograph taken with a smartphone camera); it does not analyze live video, does not process X-rays or any other radiographic input, and does not require specialized intraoral camera hardware. It is limited to the five classes named above and does not detect oral cancer, bone loss, impacted teeth, orthodontic issues, or other conditions that surface-level visual detection cannot reliably identify. The comparison of detection models is limited to YOLOv12 and YOLOv8 and does not include other detection architectures. The system is developed and tested only at CodiSmile Dental Clinic and does not integrate with billing or appointment scheduling systems. The evaluation covers software performance, quality, and acceptability, and does not measure long-term clinical outcomes such as changes in caries rates over time.

Limitations

The accuracy of the detection models depends on image quality, lighting, camera positioning, and the patient's cooperation when the photograph is taken, and may vary

more than it would with specialized intraoral camera hardware since a standard device camera is used instead. Because the models are trained on public datasets, performance in the clinic may differ from training results if the clinic's captured images look different from the dataset images. The quality of the KNN-generated risk classification and recommendations depends on its reference data, which starts as synthetic cases and grows as the dentist confirms tiers. Findings may not generalize to other clinics, camera devices, or patient populations.

The quality of the KNN-generated risk classification and recommendations depends on its reference data, which starts as synthetic cases and grows as the dentist confirms tiers. Findings may not generalize to other clinics, camera devices, or patient populations. Finally, the study deployment uses free hosting tiers: the backend sleeps after 15 minutes without requests, so the first request after a quiet period can take about a minute while the service starts again, and the backend has 512 MB of memory, which is why the deployed models run on ONNX Runtime and are loaded one at a time for each screening.

1.5 Significance of the Study

This study is expected to benefit the following groups:

Patients. Patients of CodiSmile Dental Clinic may benefit from a more consistent screening experience, in which findings such as tartar, plaque, cavities, gum inflammation, and healthy tooth are shown visually, explained in plain language, and accompanied by a risk classification. The personalized prevention plan gives them specific guidance based on their own findings, which can support earlier preventive action and better awareness of their oral health, and the patient portal allows them to review their own screening history at any time.

Dentists and CodiSmile Dental Clinic. Toothwise serves as a screening support tool that gives licensed dentists a consistent, image-based reference during consultations, while diagnosis and treatment decisions remain with them. The per-patient record of

images, findings, and risk classifications gives the clinic a standardized documentation of screenings that can be reviewed and compared across visits, which the clinic currently lacks, and the PDF reporting and audit logging features support clinical recordkeeping and administrative oversight.

IT Practitioners and Software Developers. The comparison of YOLOv12 and YOLOv8, the use of KNN for risk classification and recommendations, and the ISO/IEC 25010 evaluation provide a documented example of building and assessing an AI-based health support system in a real clinic setting.

Future Researchers. Students and researchers interested in AI in dentistry can use the methodology, dataset preparation, model comparison, and system design of this study as a basis for further work, such as expanding the detectable conditions, testing other clinics, or studying long-term clinical outcomes.

The Research Community. There are relatively few local studies that apply AI-based detection to dental screening in the Philippine setting. This project contributes to that growing body of work and may serve as a reference for similar initiatives at other clinics across the country.

1.6 Definition of Terms

For clarity and common understanding, the following terms are defined based on how they are used in this study.

Area Under the ROC Curve (AUROC). A threshold-independent measure of how well a model separates photographs that show a condition from photographs that do not, from 50% (no better than chance) to 100% (perfect separation).

Calculus (Tartar). A hardened deposit of mineralized plaque that forms on the tooth surface and along the gumline. In this study, it refers to one of the five conditions detected by the Toothwise system.

Dental Caries (Cavity). The decay and structural breakdown of tooth enamel and underlying tissue caused by bacterial activity. In this study, it refers to one of the five conditions detected by the Toothwise system.

F1-score. The harmonic mean of precision and recall, which summarizes both in a single value.

Gum Inflammation (Gingival Inflammation). A visible reddening, swelling, or irritation of the gum tissue, often an early indicator of gingivitis. In this study, it refers to one of the five conditions detected by the Toothwise system.

Image-Level Evaluation. In this study, the primary way of scoring the detection models: each photograph is judged by whether the model correctly reports that a condition is present or absent in it, regardless of the exact position of the bounding box. It reflects the screening question that Toothwise answers for the dentist.

Inference Speed. The amount of time a trained detection model takes to process an input image and produce a detection result. In this study, it is one of the metrics used to compare the performance of YOLOv12 and YOLOv8.

K-Nearest Neighbors (KNN). A supervised machine learning algorithm introduced by Cover and Hart (1967), which classifies a new case by identifying and comparing it to the most similar cases in a reference dataset. In this study, KNN is used to classify each screening into a risk tier and generate a personalized prevention plan based on a patient's detected findings.

Mean Average Precision (mAP). A standard evaluation metric for object detection models that summarizes a model's precision and recall performance across all detected classes. In this study, it is one of the metrics used to compare YOLOv12 and YOLOv8.

ONNX Runtime. An open-source engine for running machine learning models saved in the Open Neural Network Exchange (ONNX) format. In this study, the deployed

system uses ONNX exports of the trained YOLOv8 and YOLOv12 models so that both can run on a small cloud server without PyTorch.

Plaque. A soft, sticky film of bacteria that forms on the surface of teeth. In this study, it refers to one of the five conditions detected by the Toothwise system.

Precision. A performance metric that measures the proportion of correctly identified detections out of all detections made by the model. In the image-level evaluation, it is the share of photographs reported to show a condition that actually show it. In this study, it is used to evaluate and compare the performance of YOLOv12 and YOLOv8.

Recall. A performance metric that measures the proportion of actual conditions in an image that were correctly identified by the model. In the image-level evaluation, also called sensitivity, it is the share of photographs showing a condition that the model correctly reports. In this study, it is used to evaluate and compare the performance of YOLOv12 and YOLOv8.

Risk Score. A numerical value from 0 to 100, classified as low, medium, or high, generated by the system from a patient's detected findings to indicate the overall severity of a screening result. In this study, the risk score is computed by a rule-based formula.

Specificity. The share of photographs without a given condition that the model correctly reports as not showing it; the complement of the false-alarm rate.

Toothwise. The AI-assisted dental screening support system developed in this study, which uses YOLOv12 and YOLOv8 for detection and KNN for risk classification and generating personalized prevention plans.

YOLOv8 (You Only Look Once, version 8). An object detection model architecture in the YOLO family, used in this study as a baseline for comparison against YOLOv12 in detecting dental conditions.

YOLOv12 (You Only Look Once, version 12). An attention-centric object detection model architecture introduced by Tian et al. (2025), used in this study as the primary detection model for identifying dental conditions from intraoral images.

CHAPTER II

REVIEW OF RELATED LITERATURE AND STUDIES

This chapter presents the related literature and studies that have been gathered by the proponents. It introduces many ideas linked to the topic, which includes the research's principal focus. The study's synthesis is offered to analyze the previous research that the researchers have gathered.

2.1 Related Literature

Convolutional Neural Networks for Dental Image Diagnostics. Schwendicke et al. (2019) carried out a scoping review of 36 studies, published from 2015 to 2019, that applied convolutional neural networks (CNNs) to dental images. Most of the studies detected, segmented, or classified anatomical structures such as teeth, jaw bone, and skeletal landmarks, and only a few targeted pathologies such as caries. The authors noted that CNNs have shown the potential to detect structures and pathologies on medical images with accuracy similar to, or even higher than, that of professionals, but they also warned of possible bias in the underlying algorithms and of limited robustness and generalizability. These cautions carry directly into Toothwise, where the models are evaluated on a held-out test split and all detected findings at CodiSmile Dental Clinic are reviewed by a licensed dentist before being communicated to any patient. The review also shows that the field has concentrated on detection and segmentation tasks, which is the gap that the personalized prevention plan feature of Toothwise was built to address.

Deep Learning for Caries Detection: A Systematic Review. Mohammad-Rahimi et al. (2022) reviewed 42 studies that used machine learning and deep learning to detect caries across different image types, including intraoral photographs. Their main takeaway was that deep learning, particularly convolutional neural networks, worked reliably for identifying both early and more advanced carious lesions. They also raised a concern that is very relevant to this project: there are not many large, standardized, and publicly available dental datasets out there, which makes it harder to train models that

perform well across different clinical settings. For Toothwise, this limitation was addressed by combining publicly available annotated datasets with data augmentation techniques during model training. The studies they reviewed covered classification, object detection, and segmentation only, with image, object, or pixel accuracy between 68% and 99%, and the authors concluded that deep learning models can serve as an assistant for decisions on the presence or absence of carious lesions. None of the reviewed models went on to produce a recommendation, which reinforces why the prevention plan feature of Toothwise fills a genuine gap in the existing body of work.

Deep Learning for Detecting Dental Plaque and Gingivitis from Oral Photographs. Moharrami et al. (2025) conducted a systematic review evaluating deep learning models specifically for detecting dental plaque and gingivitis from RGB intraoral photographs, including images captured through standard consumer devices rather than specialized dental equipment. Across the studies they reviewed, models showed stronger and more consistent performance for plaque detection than for gingivitis, where the certainty of evidence was rated lower due to the more subtle and gradual visual presentation of early gingival inflammation. This finding is directly relevant to Toothwise, since gum inflammation is one of the five conditions the system is trained to detect from standard-camera photographs; it signals that this class may require more careful dataset curation and potentially larger sample sizes than the other conditions to reach comparable performance during training.

Oral Screening of Dental Calculus, Gingivitis, and Dental Caries through Segmentation on Intraoral Photographic Images. Liu et al. (2024) developed a deep learning segmentation system, Oral-Mamba, capable of simultaneously detecting three conditions — dental caries, dental calculus, and gingivitis — from a single set of intraoral photographic images, achieving accuracy rates above 0.81 for each condition. This study is one of the few in the literature that approaches multi-condition detection from a single intraoral image source rather than treating each condition as a separate, isolated detection task. It directly supports the feasibility of the multi-class detection approach used in

Toothwise, which similarly trains a single model to identify several distinct conditions — tartar, plaque, cavity, and gum inflammation — from one still-image source.

Oral Health Philippines Country Profile. The World Health Organization (2022) released a country-specific oral health profile for the Philippines as part of its Global Oral Health Status Report. The numbers are difficult to ignore. Dental caries was listed as one of the most prevalent health concerns in the country, with high rates of untreated cases across all age groups. The report also pointed out something that many Filipinos already know from experience: access to dental care is not evenly distributed. People in rural and community settings have significantly fewer options, and the limited dental workforce makes it harder to reach those who need care the most. This profile is the most authoritative reference for understanding the national oral health landscape that Toothwise seeks to contribute to. Deploying an artificial intelligence-assisted screening system at CodiSmile Dental Clinic is a concrete step toward addressing the access and consistency gaps that the WHO profile describes.

Pediatrician Knowledge, Attitude and Practices on Children's Oral Health in a Tertiary Public Hospital in the Philippines. Oliveros-Villarico et al. (2023) looked at the role of pediatricians in preventing dental caries among Filipino children, a role that often goes undiscussed. Their survey of 122 pediatricians at a tertiary public hospital found that most had high oral health knowledge (67.8%) and a positive attitude toward caries prevention (85.0%), but that their reported practices, such as caries risk assessment and preventive oral health education, varied. The authors describe dental caries as a tolerated and overlooked chronic disease in the Philippines, with untreated caries in the primary teeth of 85.18% of five-year-olds. For this project, this literature helps explain why a tool like Toothwise matters at the clinic level: knowledge alone has not translated into consistent early detection, and an AI-assisted screening system that flags dental conditions during a regular consultation at CodiSmile Dental Clinic can help catch problems before they get worse.

Guiding Responsible AI in Healthcare in the Philippines. Sarmiento et al. (2025) reviewed the barriers to the safe and effective deployment of AI in Philippine healthcare and drew on lessons from other Southeast Asian nations, the European Union, and the United States to propose directions for a national strategy. They noted that the Philippines lags behind fellow ASEAN member states in governmental AI readiness, and they pointed to structural constraints, such as connectivity, data stewardship, and workforce capacity, together with governance requirements, such as bias and safety monitoring, transparency, and accountability, that must accompany any deployment. These points shaped several decisions in how Toothwise was designed, particularly the requirement that all AI detection results at CodiSmile Dental Clinic are reviewed and confirmed by a licensed dentist before they are communicated to patients, that every sensitive action is written to an audit log, and that patient data is stored securely with authenticated access in line with the Data Privacy Act of 2012.

An Overview of Recommendation Techniques and Their Applications in Healthcare. Yue et al. (2021) reviewed the major recommendation approaches used in healthcare contexts, namely content-based, collaborative filtering, and hybrid methods, and mapped their application across five scenarios: dietary recommendation, lifestyle recommendation, training recommendation, decision-making support for patients and physicians, and disease-related prediction. Their review positions K-Nearest Neighbors-based approaches as a foundational technique underlying several of these healthcare recommendation scenarios, particularly where recommendations are generated by classifying a case according to its similarity to previously seen cases. This framework directly informs how the KNN-based risk classification and prevention plan feature of Toothwise was conceptualized, positioning it within the "decision-making support for patients" category identified in their review, an application area that, per their analysis, remains comparatively underexplored relative to dietary and lifestyle recommendation systems.

2.2 Related Studies

Detecting White Spot Lesions on Dental Photography Using Deep Learning. Askar et al. (2021) worked on detecting white spot lesions in intraoral dental photos captured with standard consumer cameras, using deep learning. White spot lesions are early stage demineralization, essentially the very beginning of a cavity before it becomes a visible hole. What made this study particularly useful is that it was done with a small dataset, which is a realistic constraint for many research projects, including this one. Despite that limitation, the model still managed to detect subtle surface level changes that would be easy to miss in a routine visual check. For Toothwise, this matters because it shows that even with limited annotated data and images from standard, non-specialized cameras, deep learning can pick up early, crucial changes from photographs, which is the same image modality the system uses at CodiSmile Dental Clinic. It also gave the team confidence that training a YOLOv12 model for early cavity detection is achievable even when the training dataset is not massive and specialized imaging hardware is not required.

Ultralytics YOLO (Version 8.0.0). Jocher et al. (2023) released YOLOv8 as an open source framework under the Ultralytics platform. Compared to older YOLO versions, YOLOv8 came with a cleaner training pipeline, an anchor free detection head, and support for multiple tasks like detection, segmentation, and classification all in one package. It became one of the most widely used detection frameworks in both academic and applied settings because it was relatively easy to work with while still delivering strong results. Understanding YOLOv8 matters for this project because YOLOv12 was built on the same foundation. The training tools, dataset formats, and evaluation methods used in YOLOv8 largely carry over to YOLOv12, so much of what Jocher et al. (2023) established applies directly to how the Toothwise models were trained and tested. YOLOv8 also serves as the direct comparison baseline against YOLOv12 in this study's model evaluation.

Accuracy in Diagnosing Caries in Young Permanent Molars Using Interproximal Radiographic Imaging and Validation by Artificial Intelligence. de Brito et al. (2025) compared the caries diagnoses of calibrated dentists on interproximal radiographs of the first permanent molars of children aged 6 to 9 with an AI pipeline in which YOLOv8 located the teeth and a pre-trained EfficientNet-B0 model classified each tooth as carious or sound. Agreement between the examiners was excellent (97.4%, kappa 0.88); YOLOv8 detected the carious teeth with 91% accuracy and 98% precision, and EfficientNet-B0 classified teeth with and without caries lesions with 89% accuracy. This study provides evidence that YOLOv8 is a credible baseline for caries-related detection tasks, supporting its selection as the comparison model against YOLOv12 in this study.

Deep Learning in Oral Hygiene: Automated Dental Plaque Detection via YOLO Frameworks and Quantification Using the O'Leary Index. Ramírez-Pedraza et al. (2025) trained and compared several YOLO model variants, including YOLOv11, specifically for detecting dental plaque in varied imaging conditions, using the O'Leary Index as a clinical benchmark for plaque quantification. Their results showed that YOLO-based models could reliably distinguish plaque at different stages of maturity, though newer, less mineralized plaque was harder to detect than older, more visually distinct plaque, since it tends to blend with surrounding gingival tissue. This finding is directly relevant to Toothwise, since it identifies a specific detection challenge, distinguishing early plaque from gum tissue, that the model must account for during training and dataset annotation for the plaque class.

Diagnostic Accuracy of Dental Caries Detection Using Ensemble Deep Learning with Intraoral Camera Images. Kang et al. (2024) took a different approach from most dental AI studies by using an ensemble of deep learning models rather than a single model for caries detection. By combining multiple detectors, they were able to reduce false positives while keeping sensitivity high, which is especially important in a clinical setting where unnecessary alarms can make dentists lose trust in the system. Their work confirmed that photographic intraoral images, rather than radiographs, are a

sufficient and practical input for reliable AI-based detection, validated against annotations made by an experienced dentist. This finding supports the still-image, standard-camera approach used in Toothwise, and gave the team confidence that deep learning for photographic dental screening is not just theoretically possible but has already been shown to work efficiently in a clinical environment similar to CodiSmile Dental Clinic.

AI-Assisted Detection of Interproximal, Occlusal, and Secondary Caries on Bite-Wing Radiographs: A Single-Shot Deep Learning Approach. Karakuş et al. (2024) applied YOLOv8 to the task of detecting different types of dental caries from bitewing radiographs. One thing that stood out in their findings was how well YOLOv8 handled small and overlapping lesions, which is a tricky problem in dental imaging because caries can be close together or partially hidden. Even though Toothwise uses standard photographic images rather than radiographs, the core detection challenge is similar. The model needs to find and correctly label multiple small, carious areas within a single image. The performance of YOLOv8 in this study, combined with the architectural improvements introduced in YOLOv12, provided the team with a strong technical basis for believing that YOLOv12 can handle the multi-condition detection demands of Toothwise.

YOLOv12: Attention-Centric Real-Time Object Detectors. Tian et al. (2025) introduced YOLOv12 in a preprint that gained significant attention in the computer vision community. The paper presents a new approach to the YOLO architecture centered around area attention mechanisms, which allow the model to focus on relevant regions of an image rather than processing everything equally. The results were impressive. The YOLOv12 N configuration achieved 40.6% mean average precision at an inference latency of just 1.64 milliseconds on a T4 GPU, outperforming comparable YOLOv10 and YOLOv11 models. For a system like Toothwise, where detection results need to be produced quickly enough to fit within a standard dental consultation, that combination of accuracy and speed is exactly what is needed. This paper is the primary technical reference for the detection model used in this project, and its findings directly justify

choosing YOLOv12 for evaluation against YOLOv8 for deployment at CodiSmile Dental Clinic. At the time of this review, no widely published study had applied YOLOv12 specifically to multi-condition dental screening from standard photographic images, positioning Toothwise among the earliest efforts to explore this architecture in that context.

Feedback Collection and Nearest-Neighbor Profiling for Recommendation Systems in Healthcare Scenarios. António et al. (2024) developed a physical activity recommendation system for patients in rehabilitation using a K-Nearest Neighbors sampling strategy combined with feedback collection, generating activity recommendations by identifying patients with the highest similarity to a given case based on their health profile. Their system was shown to outperform traditional collaborative filtering and purely human-constrained recommendation methods in producing recommendations suited to a patient's specific therapeutic needs, physical capacity, and stage of recovery. This is one of the few studies to apply a KNN-based approach directly to generating personalized, condition-specific health guidance rather than general product or content recommendations, and it provides a close technical parallel for how Toothwise's KNN-based feature classifies a screening into a risk tier and composes a prevention plan by matching a patient's findings against similar synthetic and dentist-confirmed cases.

Bridging the Gap or Widening the Divide: A Call for Capacity-Building in Artificial Intelligence for Healthcare in the Philippines. Gutierrez and Viacrusis (2023) took a critical look at AI readiness in Philippine healthcare. When they asked GPT-4, GPT-3, Bard, and Claude about isoniazid dosing for Filipinos compared with other ethnic groups, three of the four models produced convincing but false content. The authors argued that regions with lower medical research output risk bias and misrepresentation in AI-generated content, and they called on Filipino clinicians and researchers to begin capacity-building, including locally representative data, quality assurance, and early AI education in medical schools, before relying on such tools in clinical practice. This study

speaks directly to a design decision in Toothwise. The system does not allow AI detection results to be shown to patients without going through a dentist first. At CodiSmile Dental Clinic, the dentist reviews every output before it becomes part of the clinical record or is communicated to the patient, which reflects the caution that Gutierrez and Viacrusis called for.

2.3 Synthesis

The literature and studies reviewed in this chapter come together to build a clear and well-supported case for the approach that Toothwise has taken. Schwendicke et al. (2019) and Mohammad-Rahimi et al. (2022) showed through separate reviews that deep learning models are genuinely reliable for dental caries detection, often performing at the level of trained clinicians or better. Moharrami et al. (2025) and Liu et al. (2024) extended this evidence beyond caries alone, showing that plaque, calculus, and gingivitis can also be reliably detected from standard intraoral photographs, with the work of Liu et al. in particular demonstrating that multiple conditions can be detected from a single image source using one model, directly supporting the five-class detection approach used in Toothwise. Kang et al. (2024) and Askar et al. (2021) further confirmed that photographic images captured without specialized intraoral camera hardware are a sufficient and practical input modality for AI-based detection in an actual clinical setting, closely paralleling the still-image, standard-camera acquisition approach of Toothwise at CodiSmile Dental Clinic.

The most important gap identified across the reviewed works is that existing dental AI systems detect conditions but stop there; the photograph-based reviews by Moharrami et al. (2024, 2025) cover detection only. None of the detection-focused studies reviewed generated recommendations or patient-facing guidance after detection. At the same time, Yue et al. (2021) and António et al. (2024) show that KNN-based recommendation approaches have been successfully applied in other healthcare domains, such as physical rehabilitation, but remain comparatively unexplored in dental screening

specifically. Toothwise fills this gap by embedding a KNN-based risk classification and recommendation engine into a multi-condition detection pipeline, pairing YOLO-based detection with an actionable, personalized prevention plan, a combination not found together in the literature reviewed.

On the technical side, Tian et al. (2025) established YOLOv12 as a highly accurate and efficient real-time object detection architecture, while Jocher et al. (2023), de Brito et al. (2025), Ramírez-Pedraza et al. (2025), and Karakuş et al. (2024) documented the strong and well-validated performance of YOLOv8 specifically across caries and plaque detection tasks, supporting its selection as this study's comparison baseline. Askar et al. (2021) further showed that deep learning-based detection is achievable even under limited dataset conditions, supporting the training strategy used in this project.

The Philippine context literature from the WHO (2022) and Oliveros-Villarico et al. (2023) grounds all of this in a concrete public health reality. Dental conditions are widespread among Filipinos, they often go undetected until they have already progressed, and access to professional dental care remains out of reach for many due to cost and geography. Sarmiento et al. (2025) and Gutierrez and Viacrusis (2023) shaped responsible design decisions around human oversight, data privacy, and clinical transparency. What no existing work reviewed has done is bring together YOLO-based multi-condition detection from standard photographic images, a KNN-based risk classification and personalized prevention plan generator, and a full-stack clinical system deployed at an actual Philippine dental clinic. That is the specific contribution of Toothwise.

CHAPTER III
METHODOLOGY

This chapter presents the methodology used in the development of Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic. It describes the development methodology, the composition and responsibilities of the five-member development team, the project planning and requirements analysis process, the major system design diagrams required for the capstone examination, and the development, testing, and screenshot documentation of the current system.

3.1 Development Methodology

Toothwise was developed using the Agile Software Development Life Cycle (Agile SDLC). Agile was selected because the system integrates several interconnected components, including a React.js web interface, a FastAPI REST backend, a PostgreSQL database, a dual-model YOLOv8/YOLOv12 detection pipeline, a KNN-based risk classification and recommendation engine, authentication and authorization, audit logging, a separate patient self-service portal, and a cloud deployment of all three tiers. These components benefit from incremental development, continuous testing, regular review, and repeated refinement rather than a single sequential build.

Instead of completing the entire system in one long development cycle, the project was carried out through manageable iterations in which requirements were prioritized, features were implemented, tested, reviewed, and improved before the next set of features was developed. Each iteration followed the same general cycle of planning, requirements analysis, system design, development, testing, evaluation, and refinement, illustrated in Figure 1.

The cycle in Figure 1 was applied repeatedly across the project's documented iterations. Planning and requirements analysis identified the features required for the

clinic workflow; system design translated those requirements into frontend routes, API endpoints, and database structures; development implemented the corresponding React, FastAPI, and machine-learning components; testing verified the implemented behavior; evaluation compared the result against the intended requirements; and refinement corrected any issues found before the next iteration began.

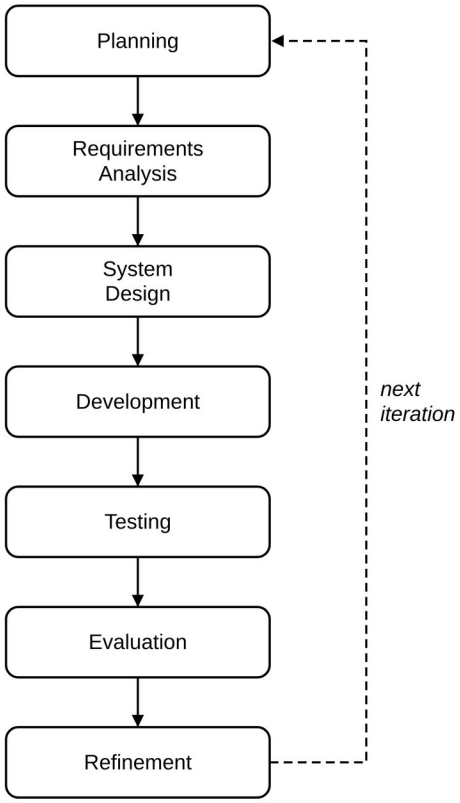


Figure 1 Agile Development Methodology.

3.2 Project Planning

The Toothwise project was planned around the screening and record-keeping workflow of CodiSmile Dental Clinic. The following considerations guided the planning of the system:

- **Project Objectives** — to provide CodiSmile Dental Clinic with an AI-assisted dental screening system that detects tartar, plaque, cavity, gum inflammation, and healthy tooth from a still photograph, estimates a risk level, and generates a

recommendation/prevention plan, while also supporting day-to-day clinic record-keeping.

- **Project Scope** — the system covers staff authentication, patient management, dental screening and AI detection from uploaded still images, risk assessment, recommendation generation, password recovery, clinical notes, PDF reporting, audit logging, a self-service patient portal, a public AI assistant, and cloud deployment.
- **Target Users** — CodiSmile Dental Clinic staff (dentist, receptionist, and assistant roles) and the clinic's patients through the patient portal.
- **Clinic Workflow** — a patient is registered or already on file, a still intraoral photograph taken with a standard device camera is uploaded during a screening, the system produces detection, risk, and recommendation results for the dentist to review, and the dentist may confirm the risk tier, add clinical notes, or generate a PDF report.
- **Functional Requirements** — detailed in Section 3.4.
- **Non-Functional Requirements** — detailed in Section 3.4, covering security, privacy, usability, reliability, performance efficiency, and data protection.
- **AI Screening Requirements** — the system was planned to use two independent detection models (YOLOv8 and YOLOv12) whose outputs are fused into a single clinical detection set, rather than relying on a single model.
- **Database Requirements** — a relational PostgreSQL schema was planned to store staff accounts, patients, screenings and their photographs, detections, recommendations, clinical notes, and audit logs with referential integrity between related records.
- **Security Requirements** — JWT-based authentication, hashed passwords, rate-limited login endpoints, secure password recovery, authenticated (non-public)

image access, and audit logging were planned from the outset given the clinical and personal nature of the data involved.

- **Deployment Requirements** — the system had to run on free cloud hosting without a payment card, which limits the backend to 512 MB of memory and a disk that is cleared whenever the service restarts.

The requirements and technical decisions behind this planning were drawn directly from the current Toothwise project package rather than from separate interview transcripts or survey forms, since the project does not maintain those as separate dated records. Table 1 summarizes the actual sources used.

Table 1 Sources of Requirements and Technical Data

Source	Use in the Study
System and code review	The frontend routes, pages, API client, FastAPI routers, schemas, security module, ML modules, and database models were reviewed to identify actual implemented functions and data flows.
Database and schema review	The PostgreSQL schema, SQLAlchemy models, and migration scripts were reviewed to verify entities, relationships, enumerated values, detection sets, stored screening photos, and audit logs.
Testing and debugging records	The automated backend test suite (118 test cases) and the project's version history document API contract reconciliation, enum and image-path fixes, smoke testing, security validation, and feature verification.
AI evaluation procedure	The evaluate_image_level.py utility computes the primary image-level measures (accuracy, precision, recall, specificity, F1-score, and AUROC per condition), the evaluate_models.py utility computes the secondary box-level measures (precision, recall, F1, mAP50, mAP50-95, and average inference time), and the check_test_leakage.py utility removes test photographs that also appear in the training or validation sets. The trained YOLOv8 and YOLOv12 weights are installed, and the measured results are reported in Chapter IV.
Deployment configuration	The Render blueprint (render.yaml), the Vercel configuration, the backend Dockerfile, and the deployment guide (docs/DEPLOYMENT.md) were reviewed to document the hosting environment and its limits.

3.3 Roles and Responsibilities of the Five-Member Development Team

Toothwise was developed by a five-member capstone development team. Table 2 presents each member's assigned role and corresponding responsibilities based on the work represented in the current project package. A member may carry more than one responsibility.

Table 2 Roles and Responsibilities of the Development Team

Group Member	Assigned Role	Responsibilities
Arreola, Reizen S.	Project Manager / System Analyst / Programmer	Coordinated the activities of the development team, monitored the project schedule, translated clinic workflow requirements into system requirements, and ensured that project deliverables remained aligned with the objectives of the Toothwise system; developed and integrated the YOLOv8/YOLOv12 dental image detection pipeline and the KNN recommendation engine.
Magpantay, Rizelle Carla N.	Backend Developer	Implemented and integrated the FastAPI REST API, authentication and authorization logic, and database operations and connected the frontend with the backend API; resolved programming and integration issues identified during testing.
Sales, Keith Ashly V.	Frontend Developer	Implemented the React.js user interface, including the dashboard, patient and dentist management screens, screening workflow, patient portal, and the visual design and layout of the system.
Tubillas, Brayan Axel L.	Database Developer / Documentation Lead	Designed and maintained the PostgreSQL database schema and relationships; and conducted functional, security, and integration testing of the system; responsible for organizing, compiling, and standardizing all documents into a cohesive capstone paper that complies with formatting, style, and academic guidelines.
Ponce, Robin C.	UI-UX Designer	Responsible for defining information architecture, and creating accessible design system for Toothwise, to ensure the capstone

Group Member	Assigned Role	Responsibilities
		deliverable delivers an intuitive, user-centered experience.

The specific division of responsibilities above reflects the major components of the current system — project coordination and requirements, backend/API development, frontend/interface development, database/documentation, and interface design work.

3.4 Requirements Analysis

The requirements of Toothwise were analyzed from the intended clinic workflow and the functions represented in the current system and its API. Toothwise has two authenticated user groups: clinic staff (dentist, receptionist, and assistant roles) and patients using the self-service portal. Tables 3 and 4 present the functional and non-functional requirements that are currently implemented in the application.

Table 3 Functional Requirements

ID	Requirement	Specification
FR-01	Staff authentication	Authenticate staff (dentist, receptionist, assistant) by email and password and issue a JWT bearer token for API access; an optional “Remember me” keeps the device signed in for 30 days.
FR-02	Role-based access	Distinguish dentist, receptionist, assistant, and patient scopes and restrict functions.
FR-03	Patient management	Create, view, update, and delete patient records and view screening history and risk trends.
FR-04	Dental screening	Upload a still intraoral photograph (HEIC/HEIF, JPEG, PNG, or WebP, up to 15 MB) taken with a standard device camera in a single multipart screening request and receive detection, risk, recommendation, and screening-record results.
FR-05	AI detection	Run the configured YOLOv8 and YOLOv12 detectors, retain raw model outputs, and produce a fused clinical detection set.
FR-06	Risk assessment	Calculate a 0–100 risk score with a rule-based formula and classify the screening as LOW, MEDIUM, or HIGH.
FR-07	Recommendation	Use the KNN recommender and condition-specific rules

ID	Requirement	Specification
	generation	to generate an ordered prevention plan.
FR-08	Screening review	View screening images, fused detections, risk information, recommendations, clinical notes, and model-comparison information; let the dentist confirm or correct the risk tier.
FR-09	PDF report	Generate an authenticated PDF screening report containing patient/screening information, findings, prevention plan, and clinical notes.
FR-10	Patient portal	Patients register, log in (with optional “Remember me”), and view their latest results, a risk-over-time chart, their next check-up date, their screening history, and screening details.
FR-11	Audit logging	Record relevant authentication and create/update/delete actions for administrative review.
FR-12	Health monitoring	Expose a health endpoint that checks database connectivity without exposing internal errors.
FR-13	Database backup	Provide a pg_dump-based backup utility that creates timestamped, retained backups.
FR-14	Photo persistence	Keep a database copy of every screening photograph and restore the file automatically when the host's disk has been cleared.
FR-15	AI assistant	Answer general dental-care and clinic questions on the public website and patient portal without reading patient records.
FR-16	Password recovery and change	Patients who forget their password can set a new one after confirming their email, birthdate, and mobile number; staff passwords are reset by the dentist from the Staff page; signed-in users can change their password.

Table 4 Non-Functional Requirements

ID	Quality Area	Specification
NFR-01	Security and privacy	A Data Privacy Notice (Republic Act No. 10173) must be read and agreed to before the system can be used; passwords hashed with bcrypt; protected operations require JWT authentication; login, registration, and password-reset endpoints rate-limited to 5 requests per minute per IP; a failed password reset always returns the same message, so it cannot reveal which accounts exist or which detail was wrong, and every reset attempt is written to the audit log.

ID	Quality Area	Specification
NFR-02	Privacy	Screening images are not publicly exposed; authenticated and ownership-checked image endpoints are used instead of a public upload route.
NFR-03	Usability	Browser-based workflows for dashboard monitoring, patient records, screening, dentist management, audit review, and patient self-service.
NFR-04	Reliability	Backend returns a generic response for unhandled errors while recording diagnostic detail server-side; screening photos survive host restarts through their database copies.
NFR-05	Performance Efficiency	YOLOv8 and YOLOv12 inference times are recorded per screening; model-comparison views report runtime and detection statistics; the deployed backend runs both models through ONNX Runtime within 512 MB of memory.
NFR-06	Data Protection	Relational constraints, enumerated status/risk values, foreign keys, and transactions preserve consistent, valid records.

3.5 System Design

This section presents the major design diagrams required for the capstone preliminary examination: the Use Case Diagram, the Entity Relationship Diagram, the Data Flow Diagrams at Level 0 and Level 1, and the development Gantt Chart. Each diagram was produced from the actual implemented system roles, database schema, and documented development sequence.

3.5.1 Use Case Diagram

Figure 2 presents the use case diagram of Toothwise. It shows the four system roles – Dentist (Administrator), Receptionist, Dental Assistant, and Patient – and the functions each role can perform, verified against the implemented routes and authorization checks. A screening always includes uploading an intraoral image, and the dentist may extend the review of findings by confirming the risk tier, adding a clinical note, or generating a PDF report. Any signed-in user may also reset a forgotten password or change it, which extends Authenticate User. The diagram represents system user roles, not the development team in Section 3.3.

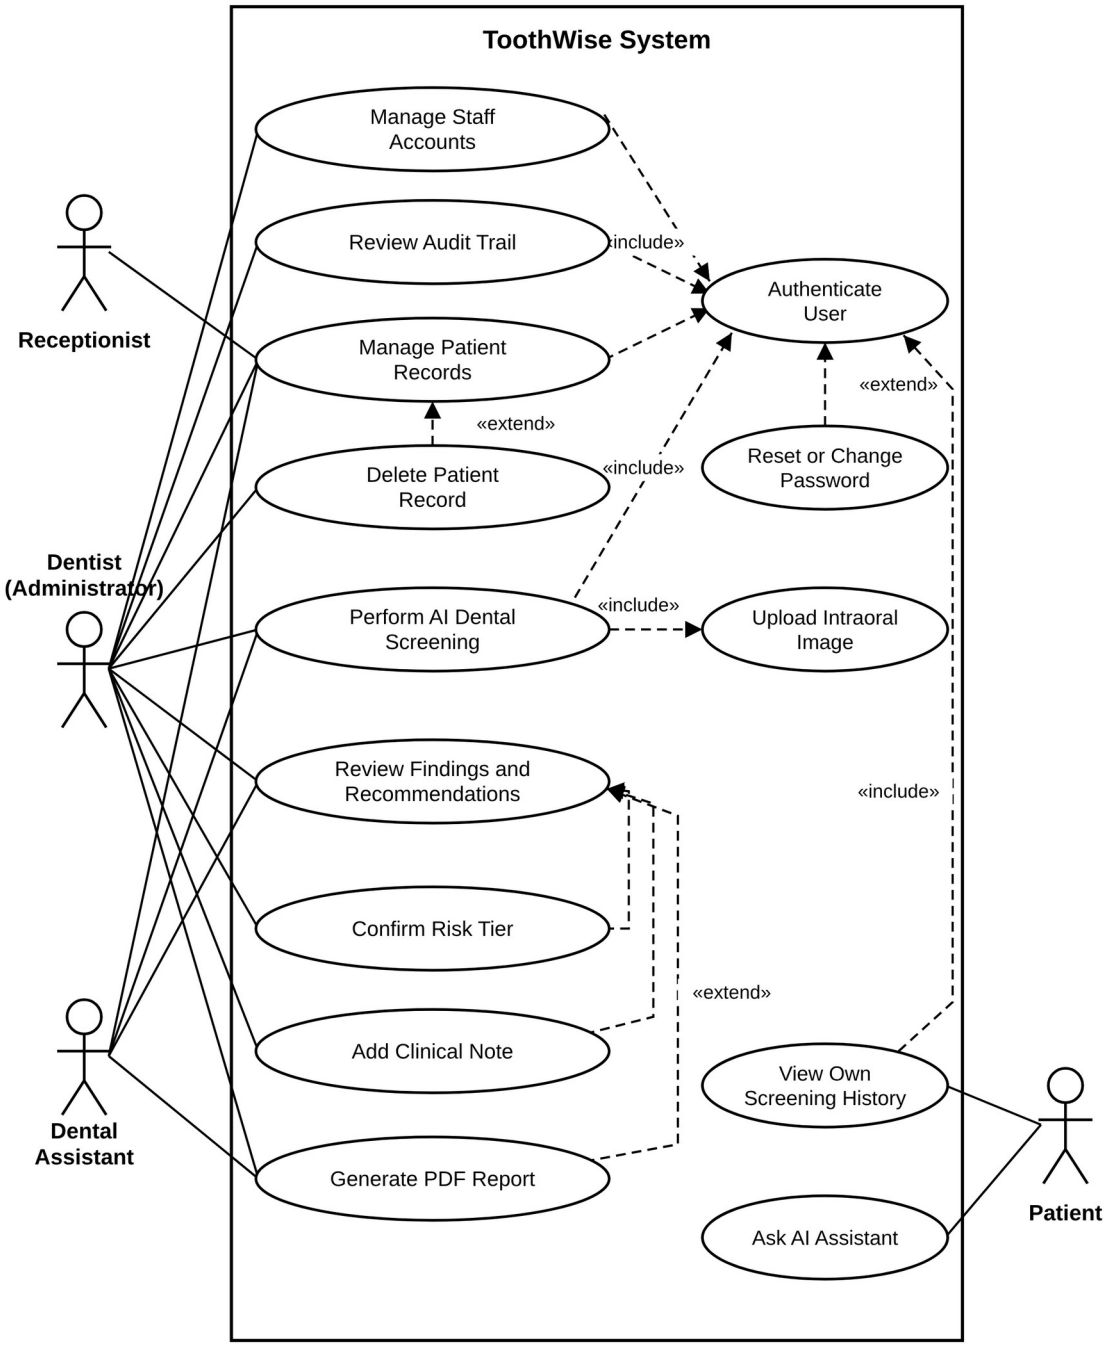


Figure 2 Use Case Diagram.

3.5.2 Entity Relationship Diagram (ERD)

Figure 3 presents the Entity Relationship Diagram of the Toothwise PostgreSQL database, verified against the current database/schema.sql and backend/app/models.py. The database contains eight tables: dentists, patients, screenings, screening_images, detections, recommendations, clinical_notes, and audit_logs, related through foreign keys as shown below. The screening_images table was added for the cloud deployment: it keeps a copy of each screening photograph in the database, because the hosting service clears its local disk whenever it restarts. Every field of each table is listed in the diagram.

Table 5 consolidates the eight entities of the Toothwise database into a single database schema, summarizing each table's primary key, purpose, classification, and relationships to the other entities.

Table 5 Toothwise Database Schema Data Dictionary

Table	Primary Key	Role/ Classification	Purpose	Related Tables
dentists	id	Dimension Table	Stores staff accounts (dentist, receptionist, or assistant) used for login and access control.	screenings, clinical_notes, audit_logs
patients	id	Dimension Table	Stores patient demographic information and, optionally, patient-portal login credentials.	screenings
screenings	id	Fact Table	Central record of one AI-assisted dental image screening event, including risk level and model timing.	patients, dentists, detections, recommendations, clinical_notes, screening_images
screening_images	screening_id	Fact Detail Table	Database copy of a screening photograph, used to restore the image file after the host's disk is cleared.	screenings
detections	id	Fact Detail Table	Stores each bounding-box finding produced by the YOLOv8/YOLOv12 ensemble for a screening (5-label taxonomy).	screenings

Table	Primary Key	Role/ Classification	Purpose	Related Tables
recommendations	id	Fact Detail Table	Stores the risk classification and prevention/advice text generated by the KNN recommender for a screening.	screenings
clinical_notes	id	Fact Detail Table	Stores dentist-authored clinical notes attached to a screening.	screenings, dentists
audit_logs	id	System Table	Append-only record of sensitive staff actions, kept for administrative review.	dentists

3.5.3 Data Flow Diagram Level 0

Figure 4 presents the Level 0 (context) Data Flow Diagram of Toothwise, drawn in Gane and Sarson notation. It shows the system as a single process exchanging data with its four external entities – Dentist (Administrator), Receptionist, Dental Assistant, and Patient – verified against the authenticated routes each role can reach.

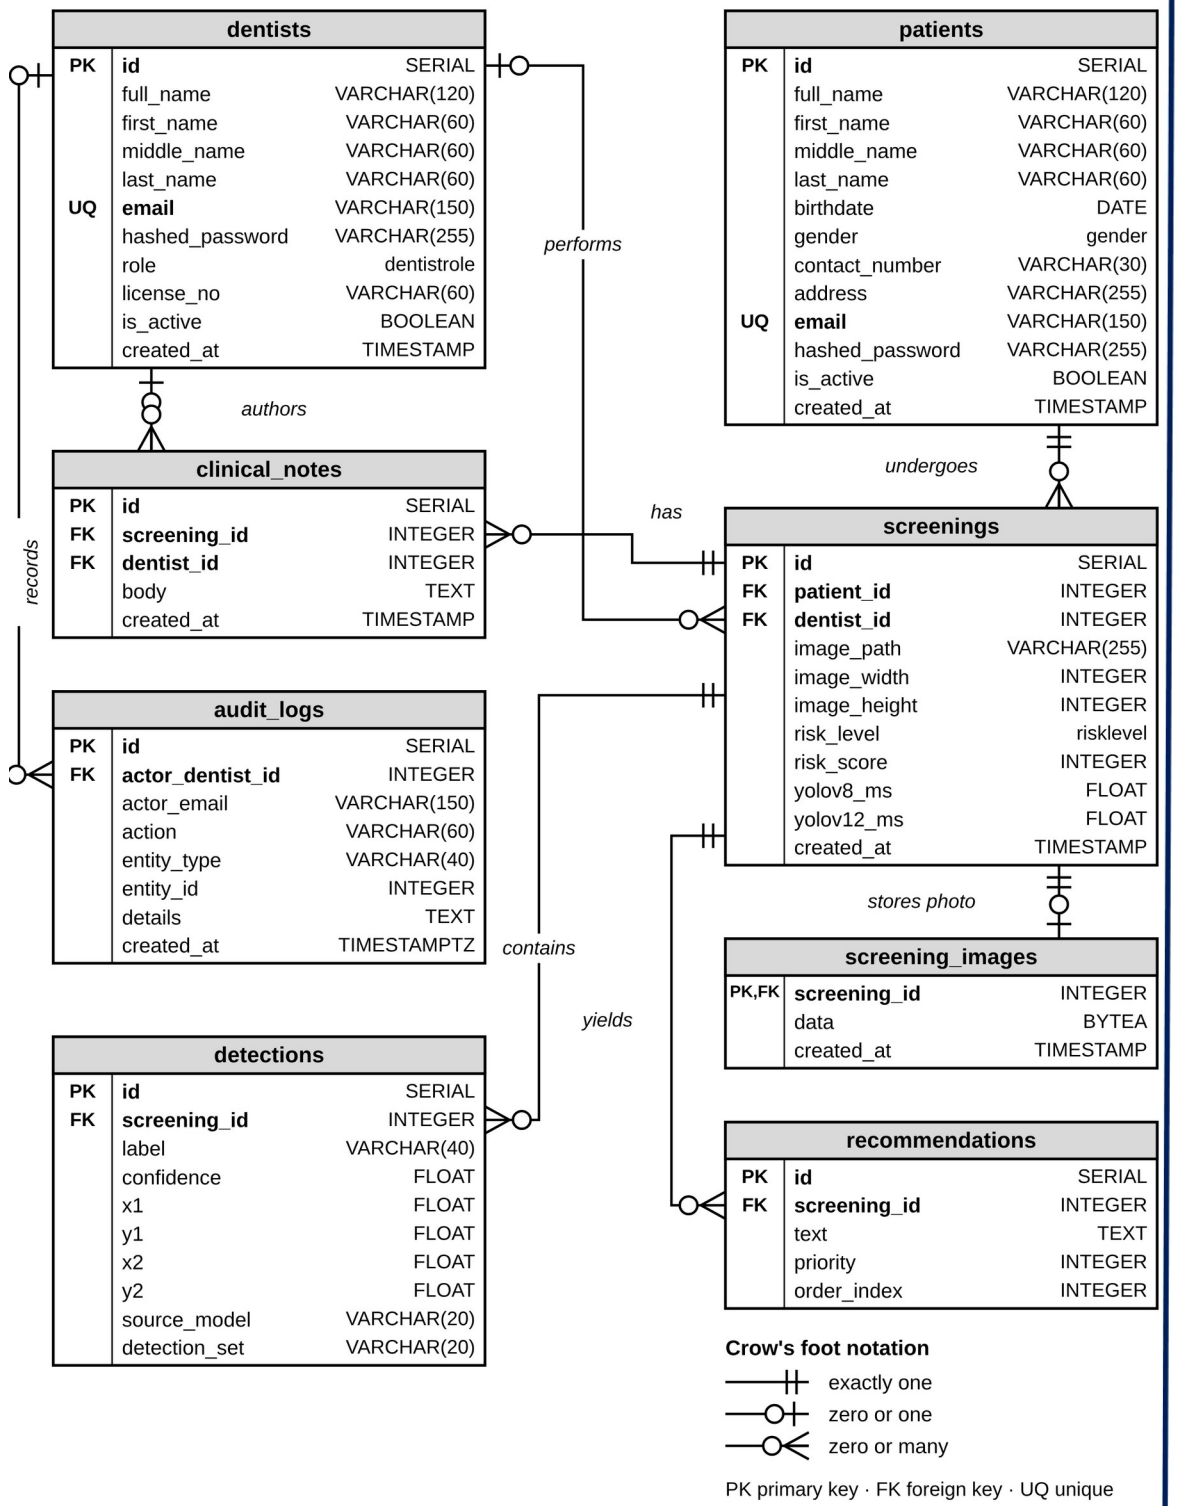


Figure 3 Entity Relationship Diagram (ERD).

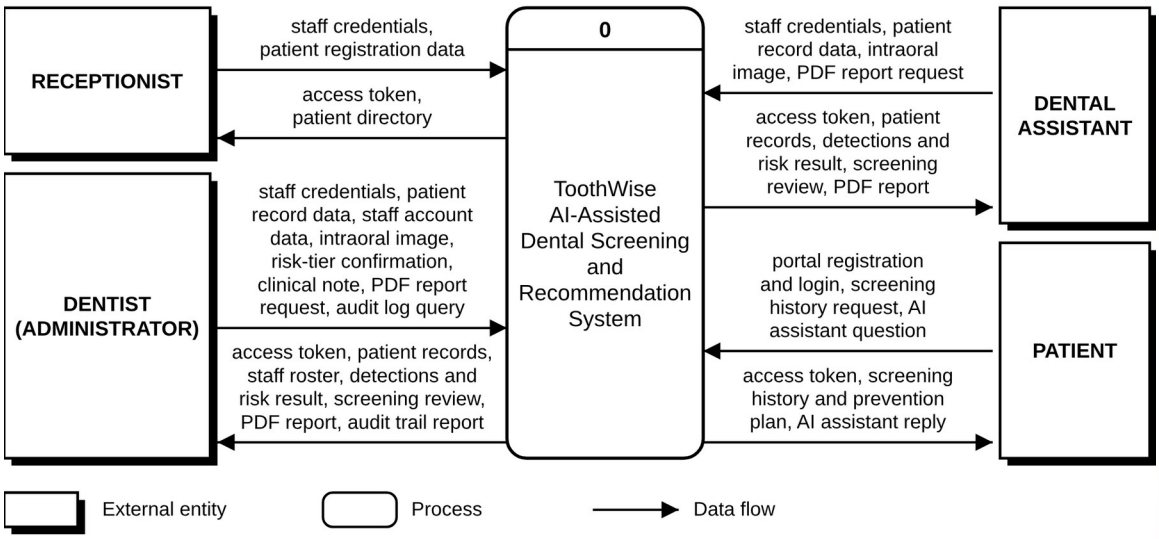


Figure 4 Data Flow Diagram Level 0.

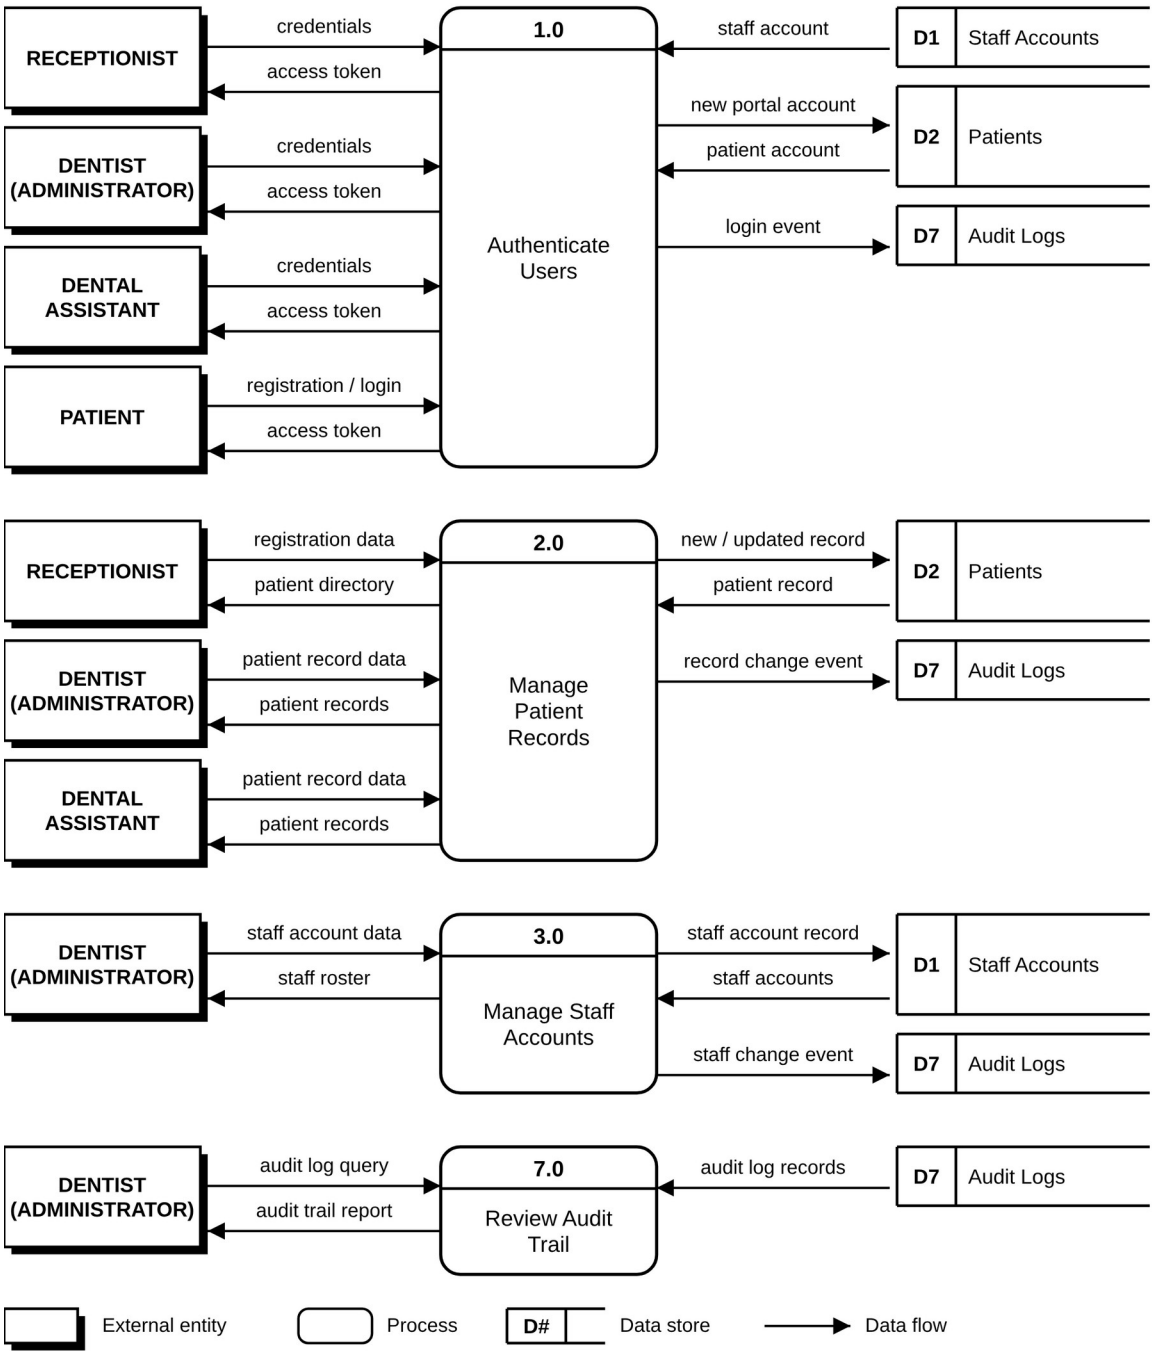
3.5.4 Data Flow Diagram Level 1

Figure 5 decomposes the Toothwise system, in Gane and Sarson notation, into seven major processes: 1.0 Authenticate Users, 2.0 Manage Patient Records, 3.0 Manage Staff Accounts, 4.0 Perform AI Dental Screening (YOLOv8/YOLOv12 detection, risk scoring, and KNN recommendation), 5.0 Review Findings and Generate Report, 6.0 Serve Patient Portal and AI Assistant, and 7.0 Review Audit Trail. So that it remains legible in print, the diagram is drawn on two sheets: Sheet 1 covers processes 1.0, 2.0, 3.0, and 7.0, and Sheet 2 covers processes 4.0, 5.0, and 6.0. Each process is drawn in its own band with its external entities on the left and its data stores on the right, so no data flows cross. The seven data stores, D1 to D7, correspond to the database tables in Figure 3 (D3 Screenings covers both the screenings and screening_images tables), and the external data flows match the Level 0 diagram in Figure 4.

Table 6 describes each Level 1 process, with its main inputs and outputs and the data stores it reads or writes.

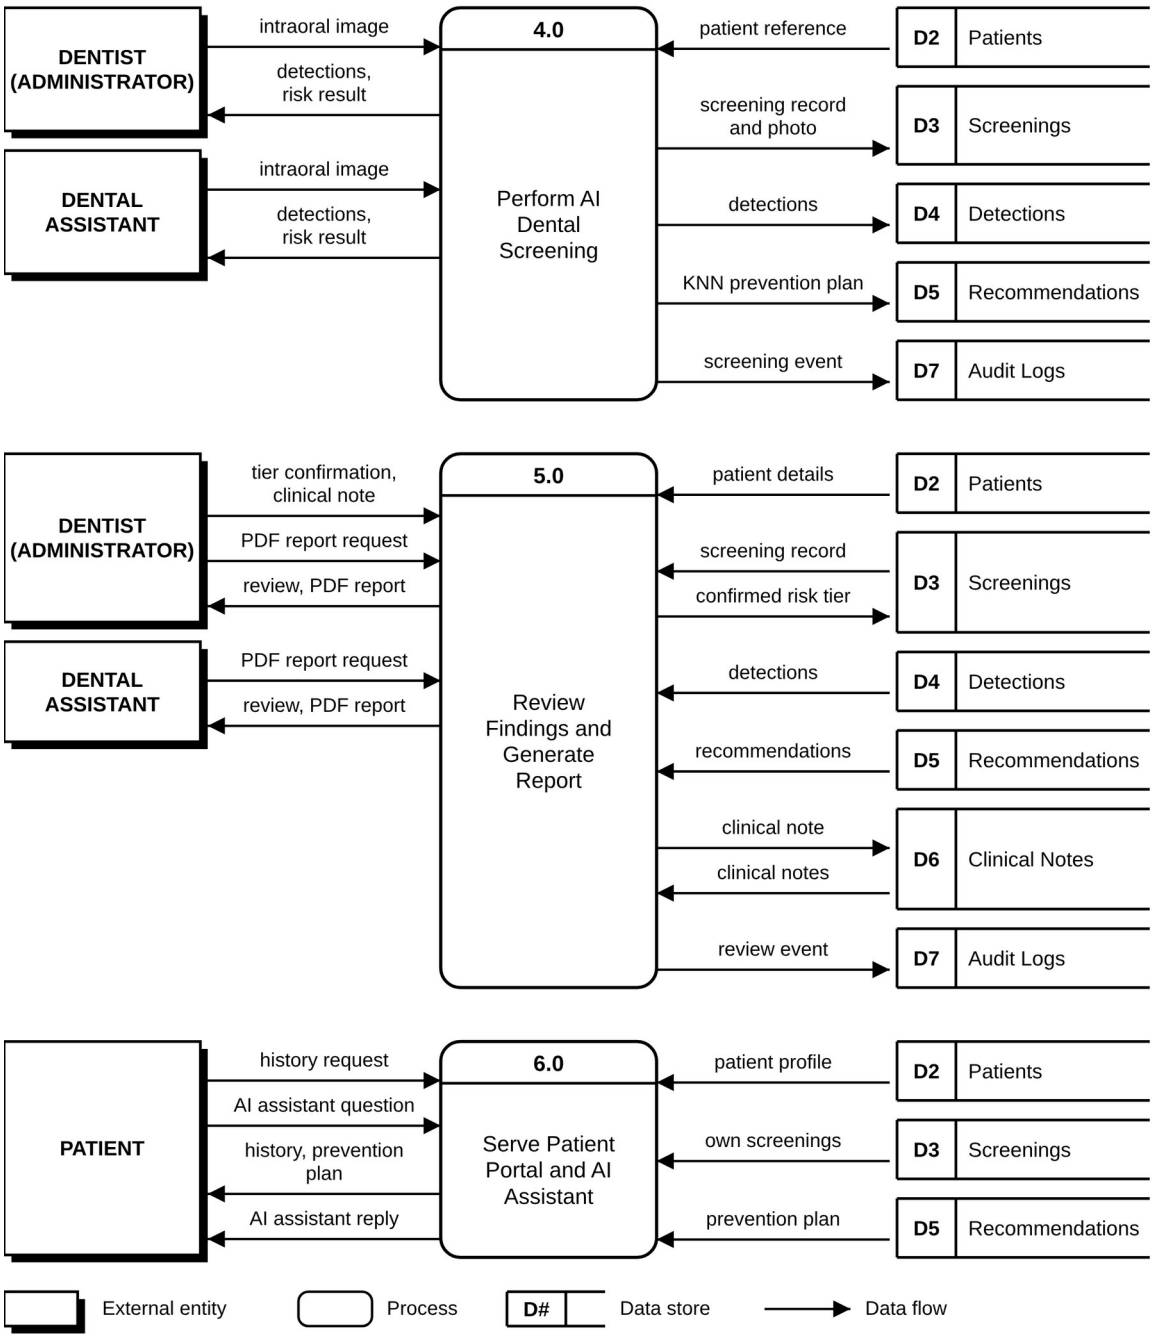
Table 6 Level 1 Process Descriptions

Process	Description	Main inputs	Main outputs and data stores
1.0 Authenticate Users	Checks staff and patient credentials and issues a JWT access token for the matching scope (30 days with “Remember me”); registers new patient-portal accounts; resets a forgotten password after checking the details on file, and changes passwords.	Credentials; portal registration; reset and change-password requests	Access token; reads D1 and D2, writes new portal accounts and new passwords to D1 and D2 and login and reset events to D7
2.0 Manage Patient Records	Creates, lists, and updates patient records; only the dentist can delete one.	Patient registration and record data	Patient directory and records; D2; record change events to D7
3.0 Manage Staff Accounts	Lets the dentist add, edit, reset, deactivate, and delete staff accounts and assign the dentist, receptionist, or assistant role.	Staff account data	Staff roster; D1; staff change events to D7
4.0 Perform AI Dental Screening	Validates the uploaded photograph, runs YOLOv8 and YOLOv12, fuses their detections, computes the risk score and tier, and builds the KNN prevention plan.	Intraoral image; patient reference from D2	Detections and risk result; writes D3, D4, and D5; screening event to D7
5.0 Review Findings and Generate Report	Shows the screening for review, records the dentist's tier confirmation and clinical notes, and produces the PDF report.	Tier confirmation, clinical note, report request; D2 to D6	Screening review and PDF report; confirmed tier to D3; notes to D6; review event to D7
6.0 Serve Patient Portal and AI Assistant	Shows patients only their own screenings, latest results, risk trend, and prevention plans and answers general questions through the AI assistant.	History request, assistant question; D2, D3, D5	Screening history, prevention plan, assistant reply
7.0 Review Audit Trail	Lets the dentist search and filter the audit log.	Audit log query; D7	Audit trail report



Entities and data stores are repeated next to each process so that no data flows cross.

Figure 5 Data Flow Diagram Level 1 (Sheet 1 of 2: Processes 1.0, 2.0, 3.0, and 7.0).



Entities and data stores are repeated next to each process so that no data flows cross.

Figure 6 Data Flow Diagram Level 1 (Sheet 2 of 2: Processes 4.0, 5.0, and 6.0).

3.5.5 Gantt Chart

Figure 7 presents the Toothwise development schedule from July 12 to October 30, organized into the five Agile iterations summarized in Table 12. Each bar shows the working period of one task, the black band shows the span of each iteration, and the diamonds mark the 30%, 60%, and 100% system completion milestones that the documentation followed. The final iteration includes the cloud deployment of the system and the move of the deployed models to ONNX Runtime.

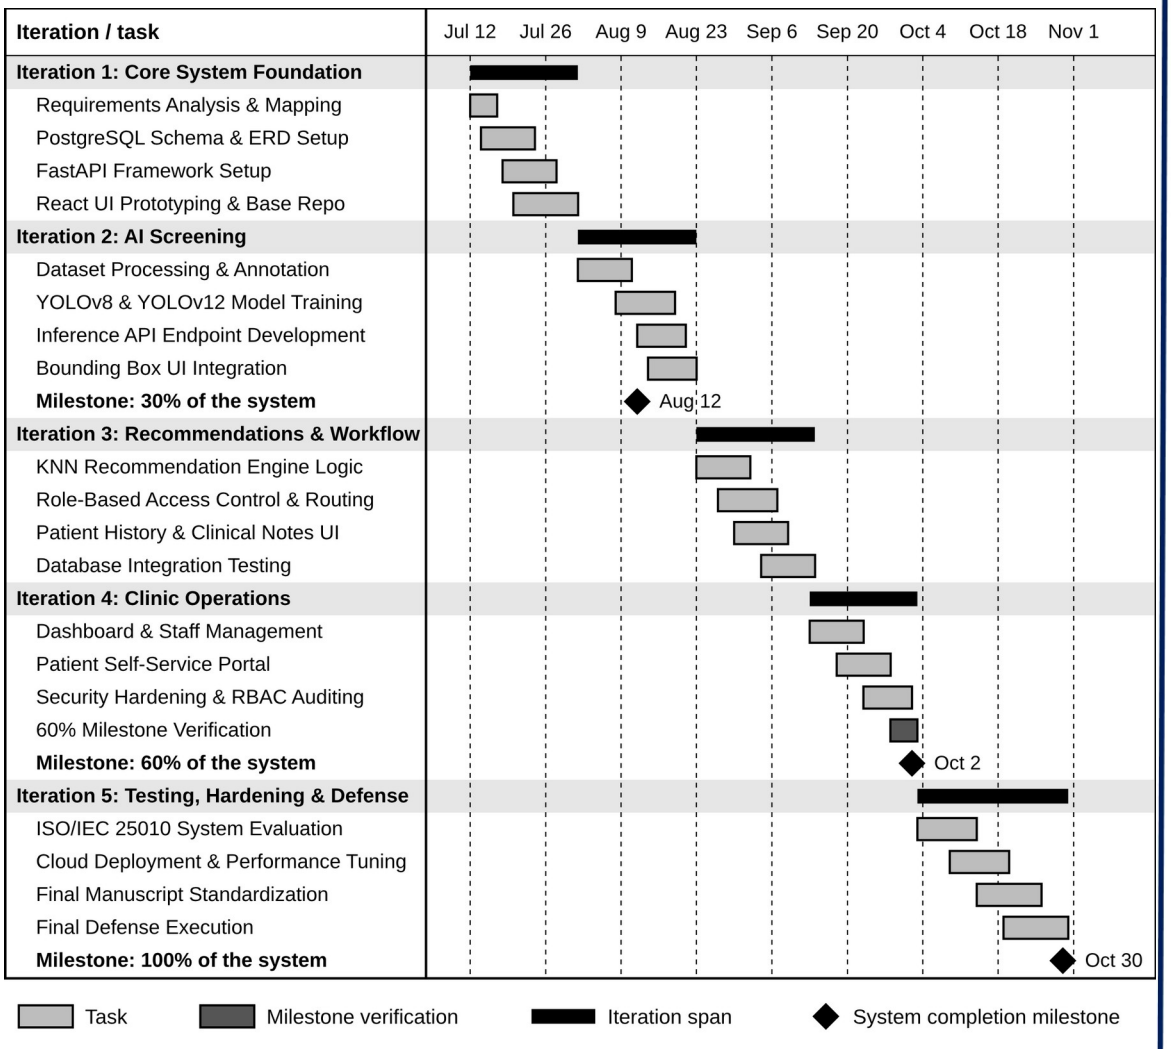


Figure 7 Toothwise Development Gantt Chart.

3.6 Software Specifications

The software requirements of Toothwise cover the programming languages, frameworks, libraries, services, and development tools represented in the current project package, distinguishing development-environment resources (Table 7) from implementation/deployment-environment resources (Table 8).

Table 7 Software Development Requirements

Software	Specification	Description
Operating System	Windows 10 or 11, 64-bit	Primary development environment for the frontend, backend, database administration, and system testing.
Runtime / Backend Language	Python 3.11	Used for the FastAPI backend, AI inference, KNN-based recommendation engine, database integration, testing, and utility scripts.
Frontend Language	JavaScript (ES6+)	Used to implement the React.js user interface and client-side application logic.
Frontend Framework	React 18.2.0 (Create React App, react-scripts 5.0.1)	Provides the component-based web interface for dentists, administrators, clinic staff, and the patient portal.
Routing	React Router DOM 6.21.1	Provides client-side routing for login, dashboard, patient records, screening, staff management, audit logs, and the patient portal.
HTTP Client	Axios 1.6.5	Handles authenticated communication between the React frontend and the FastAPI REST API.
UI Feedback and Icons	React Hot Toast 2.4.1, Lucide React 0.383.0	Provide success and error notifications and the interface icons.
Image Upload	Browser file input (drag and drop)	Accepts HEIC/HEIF, JPEG, PNG, or WebP photographs up to 15 MB (HEIC is previewed in the browser and stored as a full-resolution JPEG); on phones and tablets the same control can open the device camera.
Charting	Recharts 2.10.3	Renders the patient risk trend and dashboard trend visualizations.

Software	Specification	Description
Backend Framework	FastAPI 0.111.0, Uvicorn 0.30.1	Implements the REST API, request validation, authentication dependencies, routing, health checks, and integration with AI and PostgreSQL.
ORM / Database Driver	SQLAlchemy 2.0.30 and psycopg 3.1.19	Provide ORM-based database access and PostgreSQL connectivity.
Validation / Configuration	Pydantic 2.7.4 and pydantic-settings 2.3.4	Validate API payloads and load environment-based configuration.
AI Framework (training)	Ultralytics 8.4.116 with PyTorch	Trains and runs the YOLOv8 and YOLOv12 object-detection models during development.
AI Runtime (deployment)	ONNX Runtime 1.31.0	Runs the ONNX exports of both models in production, without PyTorch.
Computer Vision	OpenCV 4.10.0.84 (headless), Pillow 10.3.0, NumPy 1.26.4	Support image processing, validation, letterboxing, and preparation of uploaded still images.
Recommendation / ML	scikit-learn 1.5.0, joblib 1.4.2	Support the KNN risk classification and recommendation model and persistence of the trained recommender artifact.
Database	PostgreSQL 15	Stores patient, staff, screening, photo, detection, recommendation, clinical note, and audit-log records.
Authentication / Security	python-jose 3.3.0, bcrypt 4.2.0, SlowAPI 0.1.9	Provide JWT authentication, password hashing, password reset tokens, and rate limiting.
PDF Reporting	ReportLab 4.2.2	Generates authenticated PDF screening reports.
AI Assistant	Google Gemini API (gemini-3.6-flash)	Answers general dental-care and clinic questions for the public AI assistant.
Testing	pytest 8.2.2, HTTPX 0.27.0	Support automated and API-level testing of the FastAPI application.
Version Control	Git and GitHub	Used for source-code versioning, collaboration, and project traceability; pushes to the main branch redeploy the hosted system.
IDE	Visual Studio Code	Primary editor for frontend, backend, database, and AI development.

Software	Specification	Description
Package Managers	npm and pip	Manage JavaScript and Python project dependencies.

Table 8 Implementation Software Requirements

Software	Specification	Description
Frontend Hosting	Vercel (Hobby plan)	Serves the React production build over HTTPS and runs the /api/chatbot serverless function for the AI assistant.
Backend Hosting	Render free web service, Python 3.11.9, Singapore region	Runs the FastAPI backend with Uvicorn; it is rebuilt automatically after every push and checked through /api/health.
Database Hosting	Neon serverless PostgreSQL (AWS Asia Pacific, Singapore)	Stores persistent clinical and administrative records, including the database copies of screening photos.
Client Platform	Any operating system with a modern web browser	Dentists, clinic staff, and patients access the system without installing a separate application.
Browser	Google Chrome 110+ or Mozilla Firefox 110+	Supported browser environment for the React.js interface.
Authentication	JWT Bearer tokens	Separates staff and patient access scopes and protects authenticated API routes.
AI Runtime	YOLOv8 + YOLOv12 ensemble on ONNX Runtime (CPU)	Runs both detection models and fuses their outputs. Each model is loaded only while it runs so the service fits in 512 MB; its boxes, labels, and confidences are identical to the PyTorch checkpoints (checked on 22 test images).
Recommendation Runtime	KNN, K = 5	Classifies the screening into a risk tier used to compose the prevention plan.
Upload Limit	15 MB; HEIC/HEIF, JPEG, PNG, or WebP	Limits the size and type of uploaded screening images before processing.
Photo Storage	Local upload folder plus database copy (STORE_IMAGES_IN_DB)	Keeps screening photos available after the free host restarts and clears its disk.
Alternative Deployment	Docker image	For hosts with at least about 1.5 GB of

Software	Specification	Description
	(backend/Dockerfile) with PyTorch	memory, such as a Hugging Face Docker Space or a paid Render plan.
API Documentation	FastAPI Swagger UI / ReDoc	Provides developer-facing documentation for the REST API.

3.7 Hardware Specifications

Table 9 Hardware Development Requirements

Hardware	Specification	Description
Processor	Intel Core i7 or AMD Ryzen 7, 8 cores, 2.5 GHz or above	Recommended for concurrent frontend, backend, image-processing, and model-development workloads.
RAM	16 GB minimum; 32 GB recommended	Supports model training, dataset preprocessing, database services, and development tools.
GPU	NVIDIA GPU with CUDA support, at least 6 GB VRAM	Recommended for practical YOLO training and accelerated inference during model development (the final models were trained on a Google Colab NVIDIA Tesla T4).
Storage	512 GB SSD or higher	Stores source code, datasets, model weights, Python environments, PostgreSQL data, and generated artifacts.
Display	1920 × 1080 or higher	Supports development, image review, detection overlays, and responsive interface testing.
Network	Stable Internet connection, 10 Mbps or higher	Supports dependency installation, dataset access, repository synchronization, and cloud deployment.

Table 10 Implementation Hardware Requirements

Hardware	Specification	Description
Server	Cloud web service (Render free tier): shared CPU, 512 MB RAM	Runs the FastAPI API and the YOLO and KNN inference;

Hardware	Specification	Description
		measured at about 210 MB of memory at rest and 400 MB at peak during a screening.
Client Workstation	Computer, tablet, or phone with at least 4 GB RAM and a modern browser	Used by dentists, clinic staff, and patients to access the web application.
Device Camera	Smartphone or digital camera capable of capturing still images	Takes the patient's still intraoral photograph, which is then uploaded to the screening page; no specialized intraoral camera hardware is required.
Internet Connection	Broadband or mobile data, 10 Mbps or higher recommended	Connects the clinic's devices to the cloud-hosted frontend, backend, and database.
Uninterruptible Power Supply	600 VA minimum recommended	Keeps the clinic workstation and router running during short power interruptions.

3.8 System Development

Toothwise was developed iteratively, following the six development phases in Table 11, applied repeatedly across the five documented iterations in Table 12 rather than as a single sequential build.

Table 11 Development Phase Activities

Phase	Activities
1. Planning and Requirements	Define the clinic workflow, user roles, screening objectives, AI detection needs, risk assessment, recommendation output, patient portal functions, security requirements, and reporting needs.
2. System Analysis	Translate the requirements into frontend routes, FastAPI endpoints, database entities, AI modules, authentication dependencies, and data flows.
3. System Design	Design the three-layer architecture, relational schema, JWT authentication model, detection-set structure, risk classification and recommendation flow, patient ownership checks, and cloud deployment.
4. Development	Implement React pages and API services, FastAPI routers and schemas, SQLAlchemy models, PostgreSQL schema, YOLO detector

Phase	Activities
	integration, KNN risk classifier and recommender, reports, audit logs, and patient portal.
5. Integration and Testing	Reconcile frontend requests with the actual backend contract, validate database enum labels and image paths, smoke-test API workflows, test security controls, and verify the frontend build.
6. Deployment and Maintenance	Prepare the Neon PostgreSQL database, the Render backend service (which rebuilds the ONNX models from the repository at build time), the Vercel frontend build, environment secrets, database copies of screening photos, backups, and health monitoring. Continue maintenance through versioned changes, automatic redeploys on each push, and database migration scripts.

Table 12 Agile Iteration Summary

Iteration / Focus	Major Activities	Output
Iteration 1 – Core System Foundation	Initial React.js frontend, FastAPI backend, PostgreSQL schema, authentication, and core patient/dentist structures.	Working application foundation and authenticated API.
Iteration 2 – AI Screening	Integration of still-image upload, YOLO detection workflow, detection storage, risk scoring, and screening results.	End-to-end dental screening workflow.
Iteration 3 – Recommendations and Clinical Workflow	KNN risk classification and recommendation logic, prevention plans, clinical notes, screening details, and PDF reporting.	Expanded clinical screening and recommendation workflow.
Iteration 4 – Clinic Operations	Dashboard functions, staff management, audit logs, and patient portal features.	Broader clinic and patient workflows.
Iteration 5 – Testing, Hardening, and Deployment	API contract reconciliation, database and enum corrections, image access controls, rate limiting, validation, security headers, frontend integration fixes, cloud deployment on Vercel, Render, and Neon, ONNX Runtime inference, and database copies of screening photos.	Hardened, verified, and deployed system.

Each iteration in Table 12 was carried through the full cycle shown earlier in Figure 1: the corresponding requirements were planned and analyzed, the affected parts of the system were designed, the frontend and backend components were developed and

integrated, the iteration's features were tested, the result was evaluated against the intended requirement, and any issue found was corrected before the next iteration began. The Gantt chart in Figure 7 visualizes this same sequence.

The system was implemented as a three-tier web application: a React.js single-page frontend, a FastAPI REST backend exposing versioned JSON endpoints, and a PostgreSQL relational database. The AI detection pipeline (YOLOv8 and YOLOv12) and the KNN-based risk classification and recommendation engine were developed as backend modules invoked synchronously by the screening endpoint. Each tier is deployed on its own free cloud service: the frontend on Vercel, the backend on Render, and the database on Neon. The complete technical specification of this architecture is presented in Section 4.2.

3.9 System Testing

Toothwise was tested through functional, integration, and security testing conducted against the running backend using FastAPI's TestClient, in an automated suite of 118 test cases (Section 4.4.2), supplemented by manual and code-review verification for checks that are not easily exercised through an automated client alone.

To confirm that Toothwise satisfies its intended functional requirements, a systematic software testing phase was conducted. This process evaluated module-level correctness, system-wide component interaction, and user acceptance, while specifically validating the accuracy of the YOLOv8/YOLOv12 detection pipeline and the KNN risk classification and recommendation engine in producing appropriate detections, risk classifications, and prevention plans based on patient screening data.

The system was evaluated using three structured testing tiers: Unit Testing (utilizing a white-box approach), Integration Testing (utilizing a gray-box approach), and User Acceptance Testing (utilizing a black-box approach). Structured test case matrices recorded the function under test, input parameters, expected behavior, actual results, and overall pass/fail status.

3.10 Evaluation Procedure

This section describes how the acceptability of Toothwise was evaluated (the fourth problem of the study, with supporting information for the fifth): who the respondents were, the instrument they answered, how the data were gathered, and how the responses were treated. The results are presented in Sections 4.3 and 4.4.

3.10.1 Respondents and Data Gathering

Two evaluation forms were prepared in Google Forms and answered online: the IT Expert Evaluation Form, on the six ISO/IEC 25010 quality characteristics, and the user survey, on the system's main features. Participation was voluntary: each form opened with a statement explaining the purpose of the study and the handling of the data, and only respondents who agreed to it could continue. The IT experts were sent a video demonstration of Toothwise, the deployed system's address with test accounts, and the form link. In the user survey, each respondent first stated their group and then answered only the parts meant for that group. The user survey was answered from October 19 to 28, 2026 with 103 complete submissions, and the IT Expert Evaluation Form by 2 IT experts.

The study used purposive sampling, a non-probability technique in which respondents are chosen because they fit the purpose of the study (Etikan et al., 2016). Two groups were sampled. IT experts had to work in information technology and were personally invited by the researchers, because the software quality evaluation calls for technical knowledge; two took part. User respondents had to be able to judge the system from a user's point of view, as a dental patient or as an IT practitioner, and had to agree to the consent statement. A minimum of 50 user respondents was set, and 103 took part (98 patient respondents and 5 IT practitioners). Because the respondents were not randomly selected, the results describe the respondents of this study and are not generalized to all patients or all IT practitioners. Tables 28, 29, and 30 show the respondents.

3.10.2 Research Instrument

Two instruments were used. The IT Expert Evaluation Form asks for the evaluator's profile, then has 30 statements, five for each of the six quality characteristics named in the fourth problem, and three open-ended questions on the system's technical strengths, the problems found, and recommended improvements. The user survey has 14 statements on the three main features, 12 statements on the current manual process and on screening consistency and time (supporting information for the fifth problem), a software quality part for clinic staff and IT practitioners, and three open-ended questions on what respondents liked, the difficulties they met, and the improvements they suggest (Table 13).

Table 13 Sections of the Evaluation Instruments

Instrument and section	Item codes	Items	Problem addressed
IT Expert Form — Security	SE1–SE5	5	Problem 4 (quality characteristic)
IT Expert Form — Privacy	PR1–PR5	5	Problem 4 (quality characteristic)
IT Expert Form — Usability	US1–US5	5	Problem 4 (quality characteristic)
IT Expert Form — Reliability	RE1–RE5	5	Problem 4 (quality characteristic)
IT Expert Form — Performance Efficiency	PE1–PE5	5	Problem 4 (quality characteristic)
IT Expert Form — Data Protection	DP1–DP5	5	Problem 4 (quality characteristic)
User Survey — AI Detection and Risk Classification	AD1–AD5	5	Problems 1 and 4
User Survey — Results Interface and Patient Records	RI1–RI5	5	Problems 2 and 4
User Survey — KNN-Based Personalized Prevention Plan	PP1–PP4	4	Problems 3 and 4
User Survey — Current Manual Screening Process	MS1–MS7	7	Problem 1 (current process)
User Survey — Screening Consistency and Time	SC1–SC5	5	Problem 5 (supporting)

Instrument and section	Item codes	Items	Problem addressed
User Survey — Software Quality (Part B3)	SE1–DP3	20	Problem 4 (supplementary)
Total		76	

3.10.3 Rating Scale and Interpretation

Each statement is rated on a four-point Likert scale with no neutral option, so every respondent states either agreement or disagreement. The weighted mean of each statement and section is interpreted using the ranges in Table 14, which divide the 1-to-4 range into four equal intervals of 0.75.

Table 14 Four-Point Likert Scale and Verbal Interpretation

Scale	Mean range	Verbal interpretation	Acceptability
4	3.26 – 4.00	Strongly Agree	Highly Acceptable
3	2.51 – 3.25	Agree	Acceptable
2	1.76 – 2.50	Disagree	Slightly Acceptable
1	1.00 – 1.75	Strongly Disagree	Not Acceptable

3.10.4 Statistical Treatment

The weighted mean of each statement is computed as $WM = \Sigma(f \times x) / N$, where f is the number of respondents who chose a rating, x is the rating (1 to 4), and N is the number of respondents who answered the statement. The mean of a section is the average of its statement means, and the overall mean is the average of the section means. The standard deviation (SD) of the ratings shows how closely the respondents agreed with one another: the smaller the SD, the more uniform the ratings.

3.10.5 Ethical Considerations

All respondents gave their informed consent before answering. The user survey asked for the respondent's name and the IT Expert Evaluation Form asked for it as an optional field; names are not reported anywhere in this paper, and results are presented only as group statistics. No patient records or photographs from CodiSmile Dental Clinic

were collected for the evaluation, and all data were handled in accordance with the Data Privacy Act of 2012 (Republic Act No. 10173). The screening images shown in this paper are demonstration images.

3.11 Full Current System Screenshots

This section documents the current Toothwise system as implemented, one screen per numbered figure, in the order a user would encounter them: the public landing page, the data privacy notice, the staff and patient sign-in pages, password recovery and change, the dashboard, patient management, the screening workflow and results, staff management, the audit log, the rest of the patient portal, and the PDF screening report. The screenshots were captured from a running instance of the current system in light mode, using demonstration accounts, demonstration patients, and sample intraoral photographs taken from CodiSmile Dental Clinic's own published before-and-after images; no patient records are shown.

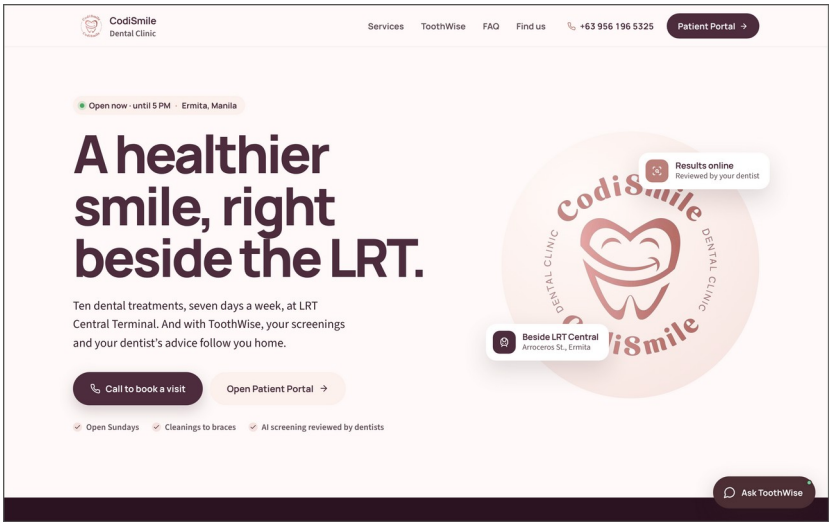


Figure 8 Public Landing Page. *This is the first page guests see. It introduces CodiSmile Dental Clinic and Toothwise, gives the clinic's location and opening hours, links to the patient portal, and opens the AI assistant.*

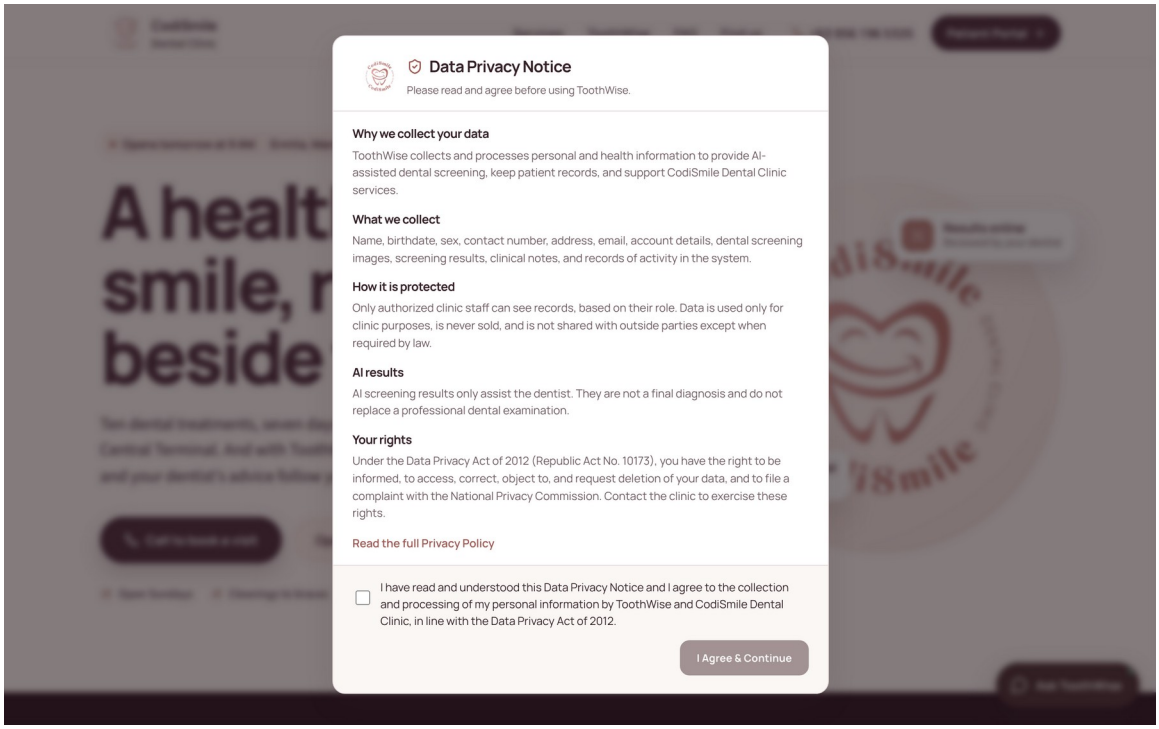


Figure 9 Data Privacy Notice. *Shown before anyone can use the system, once per browser session. It explains why and what data are collected, how they are protected, that AI results only assist the dentist, and the user's rights under the Data Privacy Act of 2012 (Republic Act No. 10173). The user must tick the agreement before continuing.*

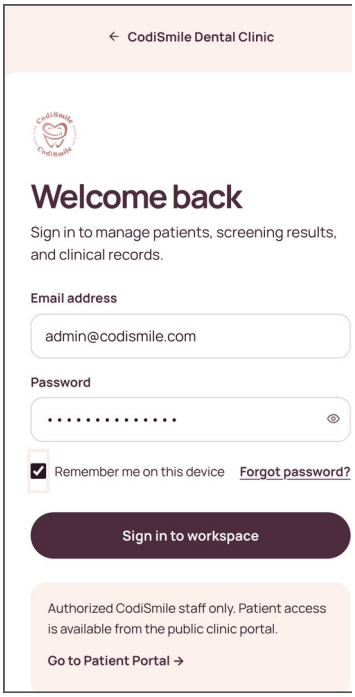


Figure 10 Staff Login Page. *Staff sign-in with "Remember me on this device" (30 days).*

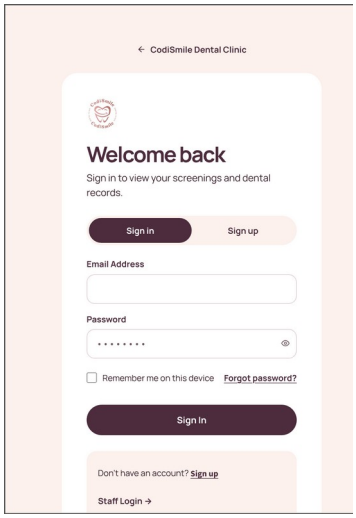


Figure 11 Patient Portal — Sign In. *Patients sign in to see only their own records.*

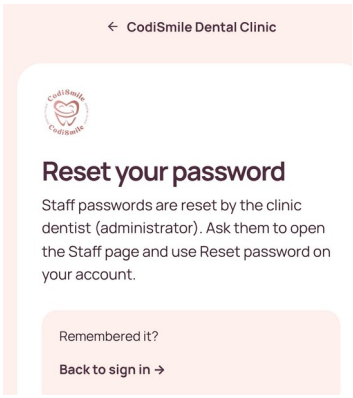


Figure 12 Reset Password (Staff). *Staff passwords are reset by the dentist from the Staff page, so the page directs staff to the dentist.*

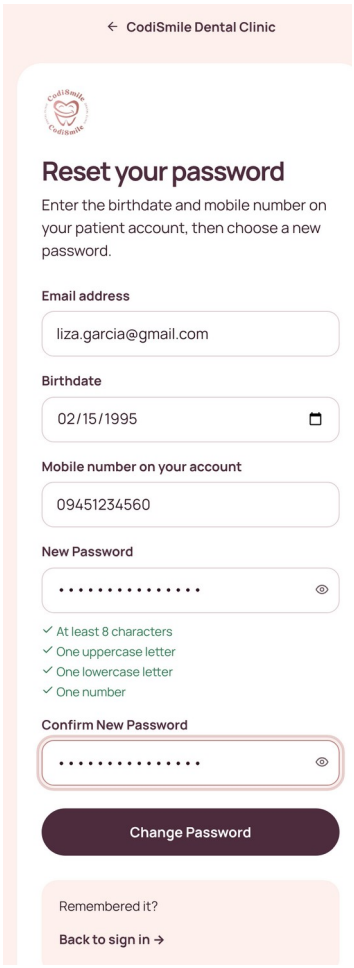


Figure 13 Reset Password (Patient). *Patients confirm their email, birthdate, and mobile number, then set a new password.*

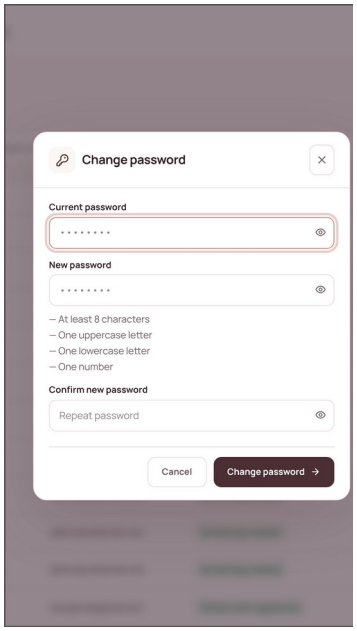


Figure 14 Change Password. *Opened from the sidebar or portal header; needs the current password.*

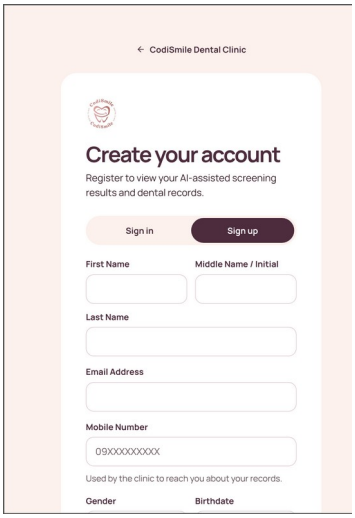


Figure 15 Patient Portal — Sign Up. *New patients create an account with their contact details and a password.*

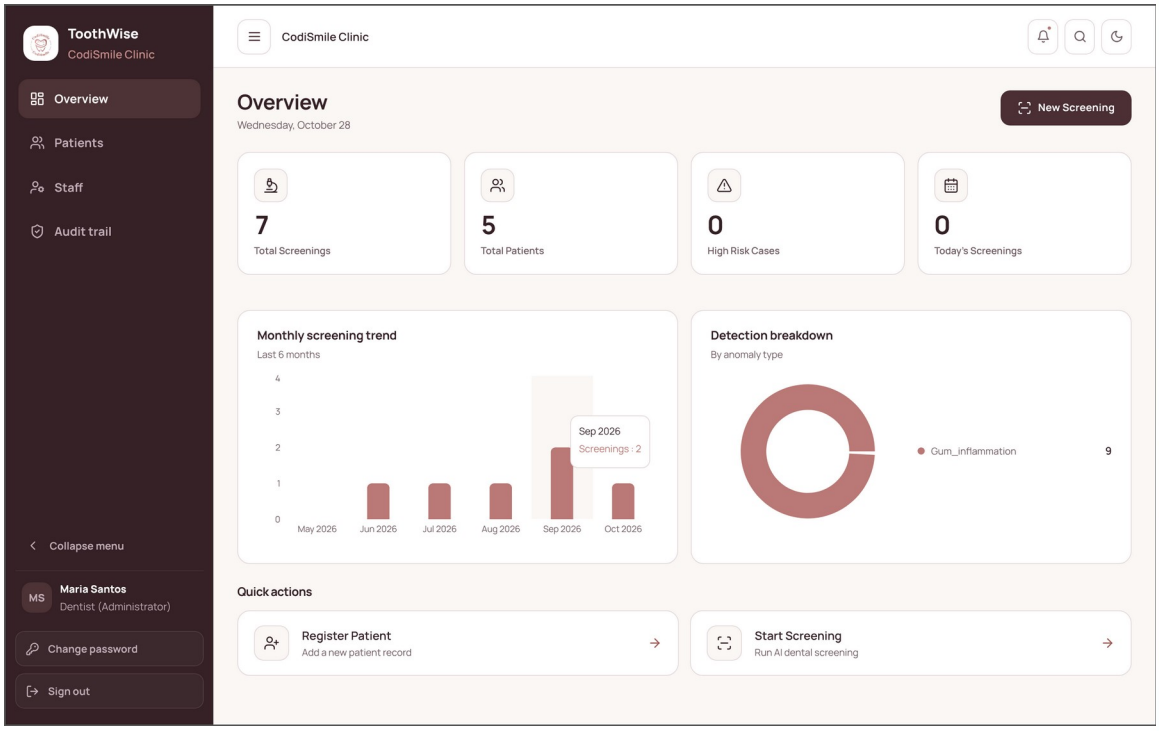


Figure 16 Dashboard. *The first screen staff see after signing in. It summarizes total screenings, patients, high-risk cases, and today's screenings, and shows the monthly screening trend and the breakdown of detected conditions.*

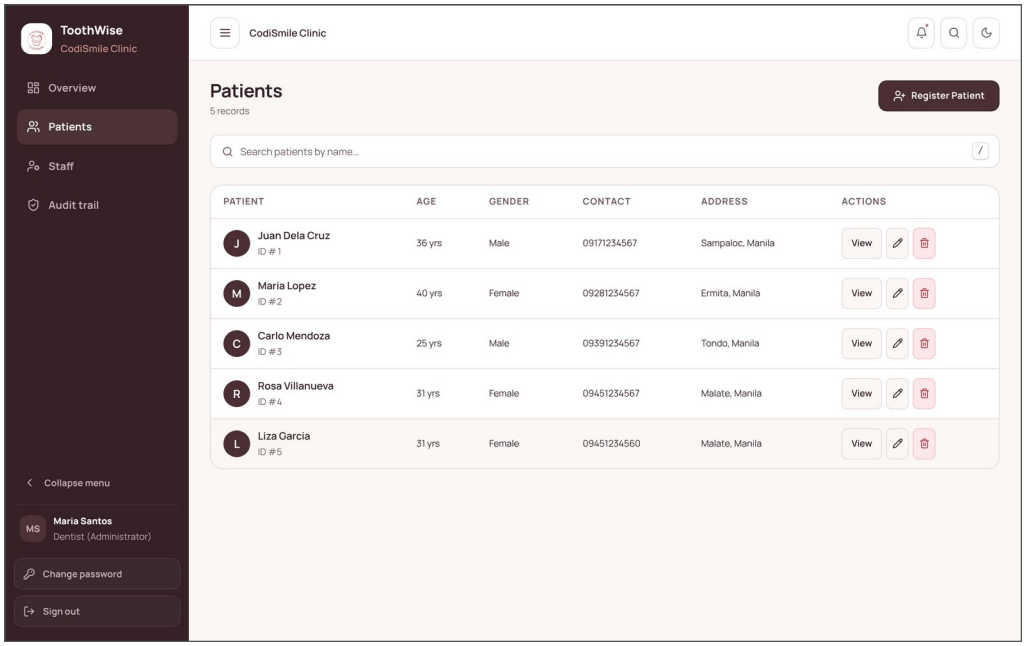


Figure 17 Patient List. *Lists every registered patient with age, gender, contact number, and address. Staff can search, register a patient, open a record, edit it, or (as dentist) delete it.*

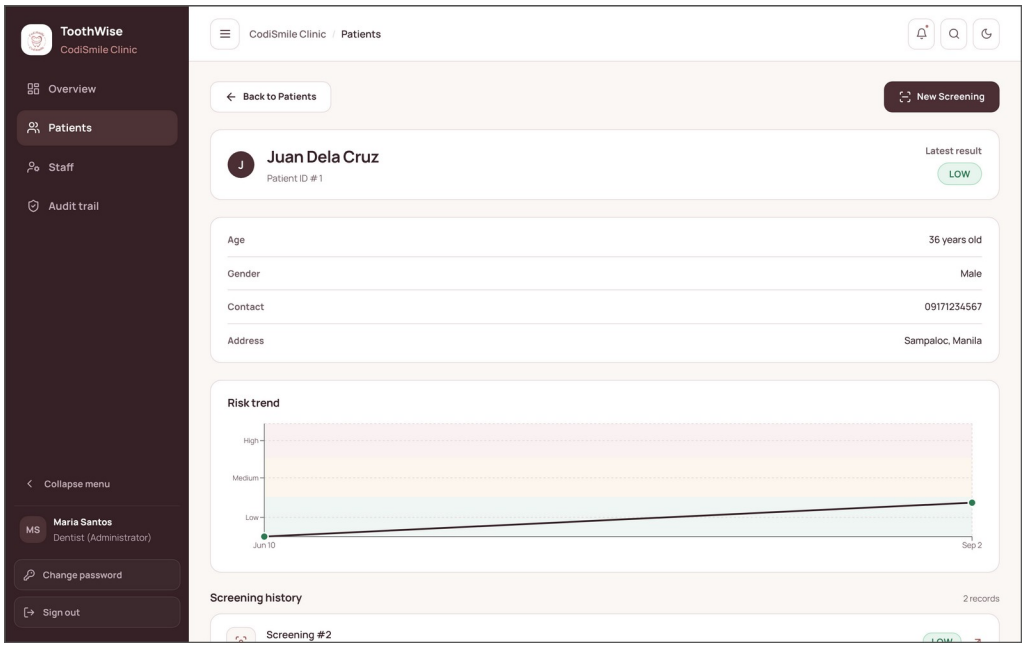


Figure 18 Patient Details. *The record of one patient: demographic details, the latest risk result, the risk trend across visits, and the screening history.*

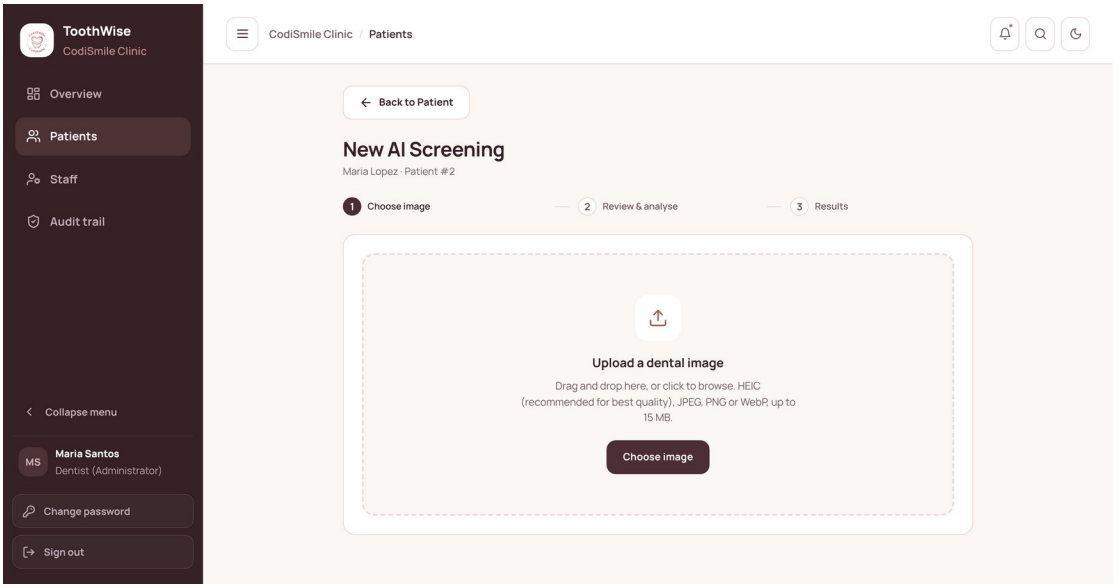


Figure 19 New Screening (Step 1: Choose Image). *Staff drag and drop or browse for the patient's intraoral photograph (HEIC, JPEG, PNG, or WebP, up to 15 MB). HEIC, the default camera format of iPhones, is recommended because it keeps the most detail.*

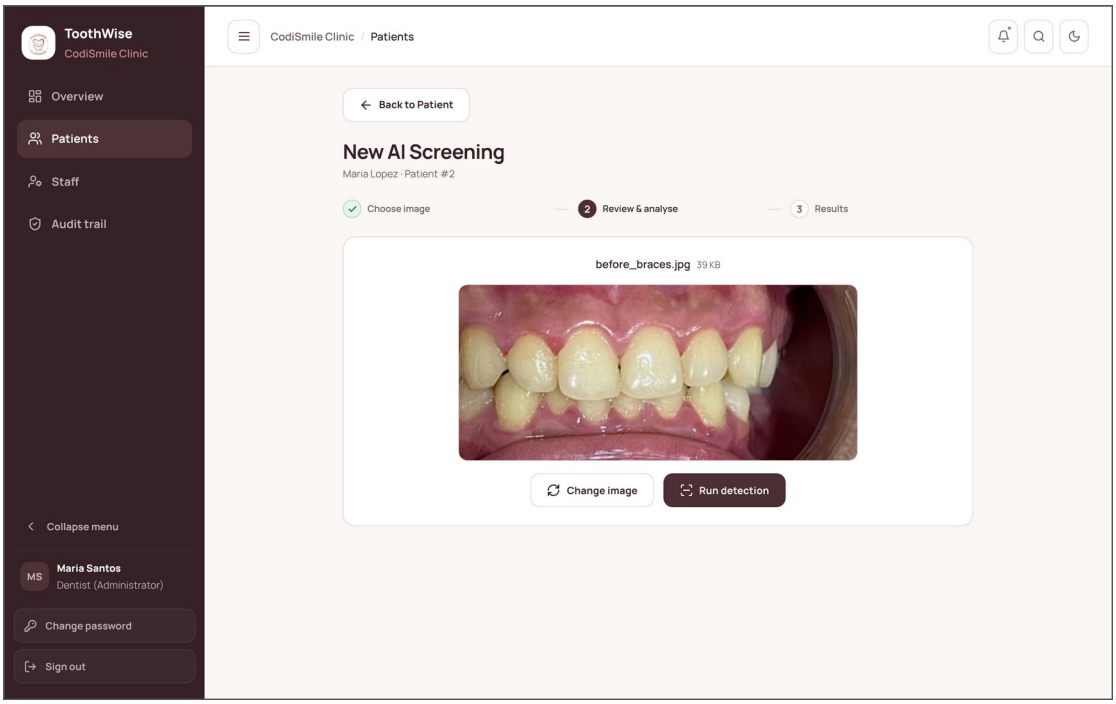


Figure 20 New Screening (Step 2: Review and Analyse). *The selected photograph is shown for checking before the YOLOv8 and YOLOv12 ensemble is run.*

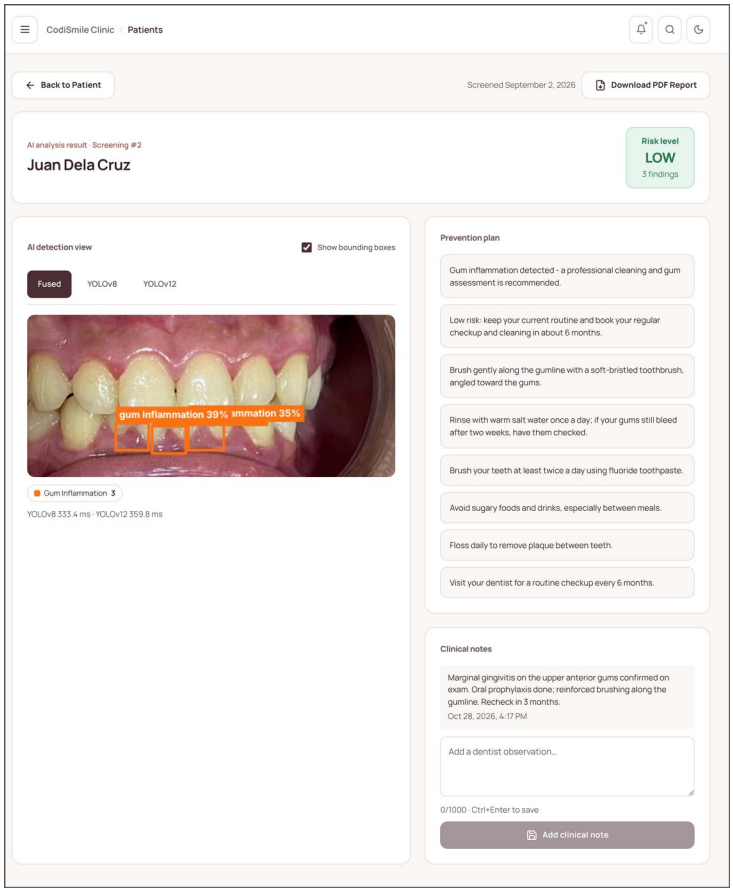


Figure 21 Screening Results. *The analysis of the uploaded photograph: color-coded bounding boxes for the fused, YOLOv8, and YOLOv12 detections (here, gum inflammation), the risk level, the prevention plan, the dentist's clinical notes, and the button that downloads the PDF report.*

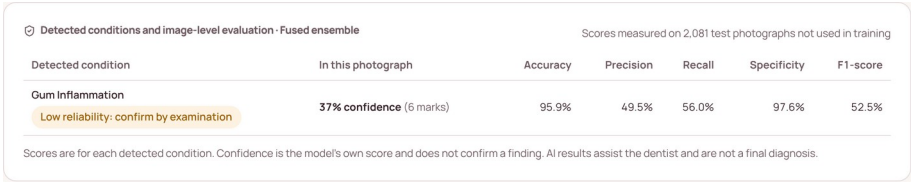


Figure 22 Detected Conditions and Image-Level Evaluation. *Shown below every screening result for the model being viewed (fused, YOLOv8, or YOLOv12). Each detected condition is listed with its highest confidence and number of marks in this photograph, next to the model's measured per-photograph scores for that condition on the 2,081 clean test photographs (Chapter IV). A condition with an F1-score below 60% is flagged as low reliability, to be confirmed by examination.*

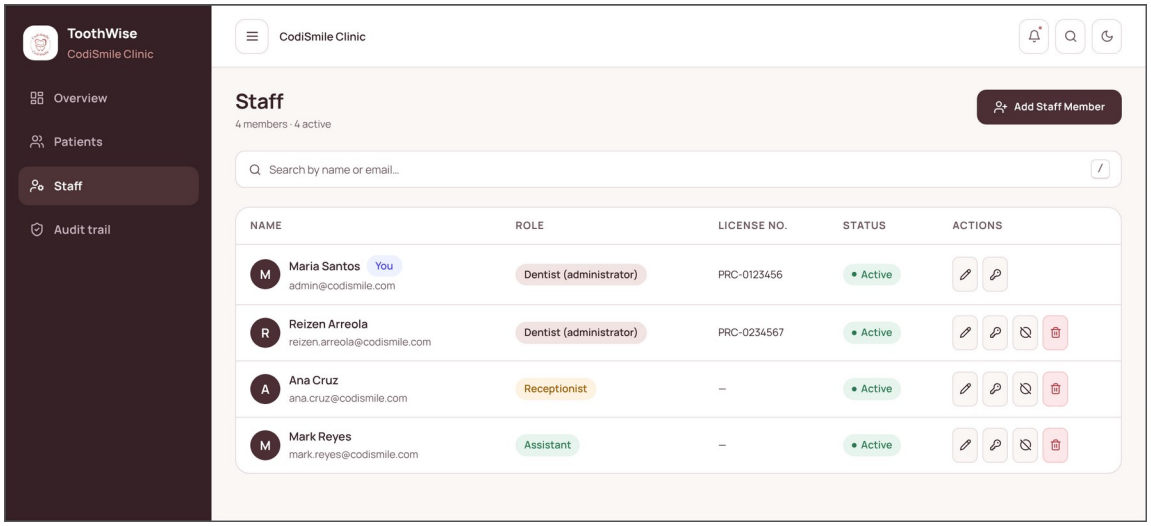


Figure 23 Staff Management (Dentist/Administrator). *Lets the dentist add staff members, assign the dentist, receptionist, or assistant role, edit details, reset passwords, and deactivate or delete accounts.*

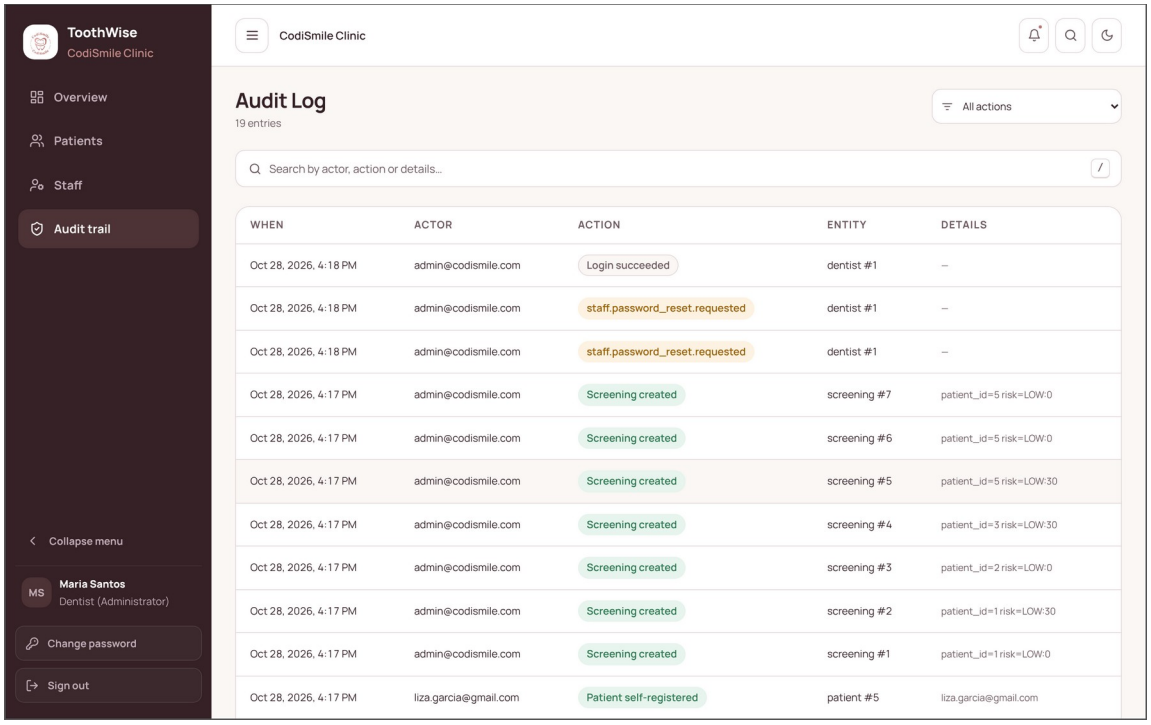


Figure 24 Audit Log (Dentist/Administrator). *A permanent, filterable timeline of who signed in, reset or changed a password, and created, changed, or deleted records, and when.*

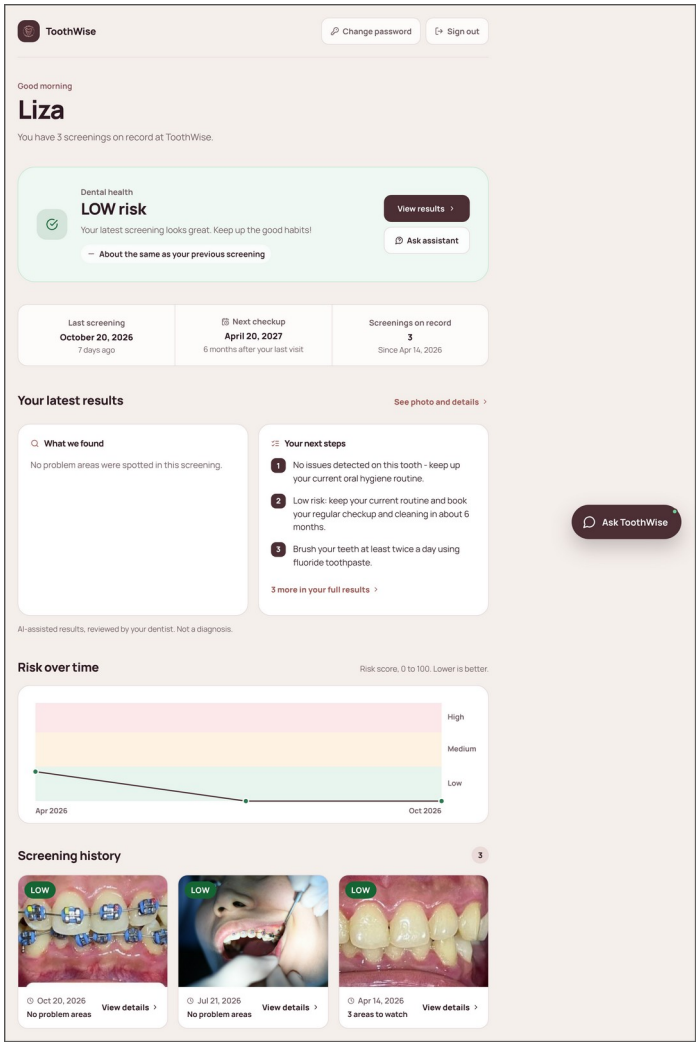


Figure 25 Patient Portal — Home. *Latest risk level and change since the last visit, next check-up date, latest findings and next steps, risk over time, and screening history.*

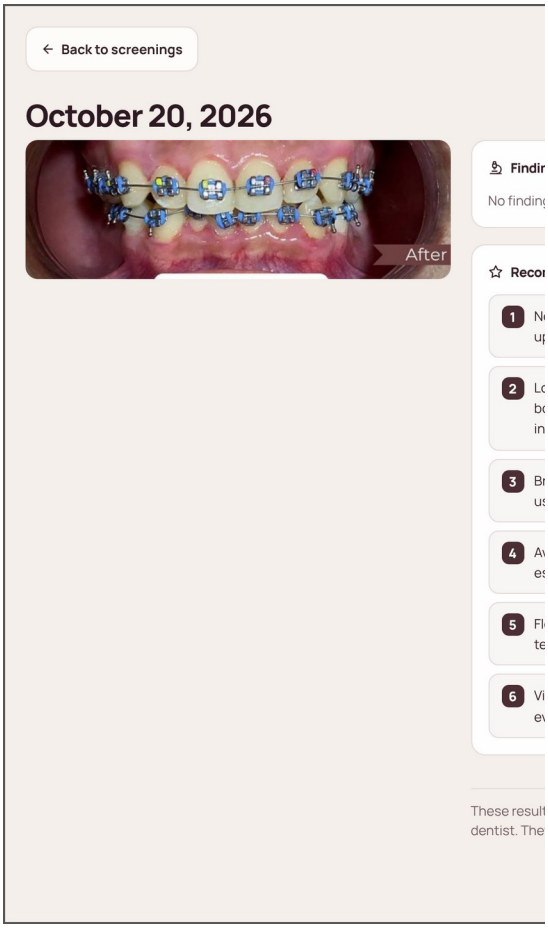
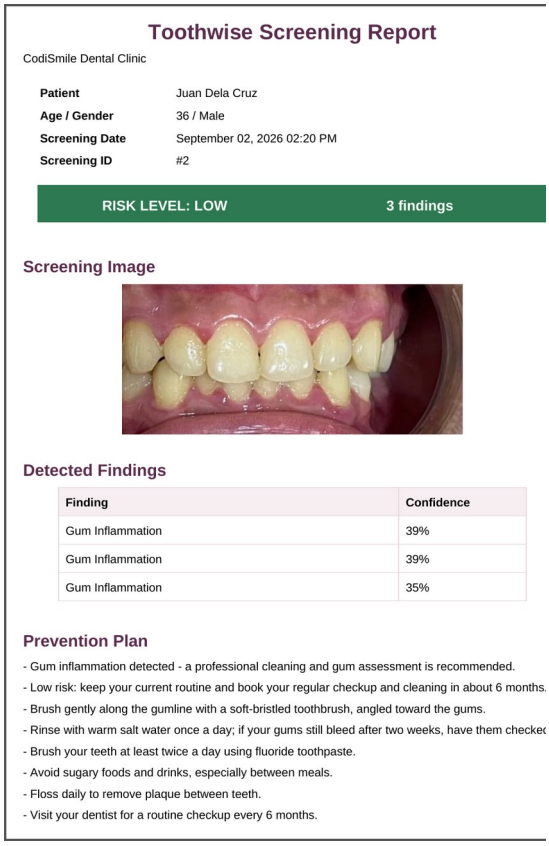


Figure 26 Patient Portal — Screening Details. *One screening: the photograph with its findings and the prevention plan.*



Toothwise Screening Report

CodiSmile Dental Clinic

Patient	Juan Dela Cruz
Age / Gender	36 / Male
Screening Date	September 02, 2026 02:20 PM
Screening ID	#2

RISK LEVEL: LOW — 3 findings

Screening Image

Detected Findings

Finding	Confidence
Gum Inflammation	39%
Gum Inflammation	39%
Gum Inflammation	35%

Prevention Plan

- Gum inflammation detected - a professional cleaning and gum assessment is recommended.
- Low risk: keep your current routine and book your regular checkup and cleaning in about 6 months.
- Brush gently along the gumline with a soft-bristled toothbrush, angled toward the gums.
- Rinse with warm salt water once a day; if your gums still bleed after two weeks, have them checked
- Brush your teeth at least twice a day using fluoride toothpaste.
- Avoid sugary foods and drinks, especially between meals.
- Floss daily to remove plaque between teeth.
- Visit your dentist for a routine checkup every 6 months.

Figure 27 PDF Screening Report. *First page of the report: patient details, risk level, image, findings, and prevention plan.*

CHAPTER IV

RESULTS AND DISCUSSION

This chapter presents the results of the study. It follows the five problems in Chapter I: how well the trained models detect the five dental conditions (Section 4.1, Problem 1), how the finished system records findings and gives prevention advice (Section 4.2, Problems 2 and 3), how the respondents rated the system (Section 4.3, Problem 4), and how fast and reliable the system is (Section 4.4, Problem 5). Section 4.3 is based on the answers of 103 respondents.

4.1 Results of the Trained Datasets

This section answers the first problem: how well can the trained models tell whether a photograph of the teeth shows tartar, plaque, cavity, gum inflammation, or a healthy tooth, and which model, YOLOv8 or YOLOv12, does it better? The main results are scored per photograph (Sections 4.1.4 and 4.1.5). The standard YOLO scores follow as supporting results (Section 4.1.6).

4.1.1 Dataset Preparation

The models learned from four free, public collections of dental photographs in which the conditions had already been marked by hand. Each collection named its conditions differently, so every condition was matched to one of the five Toothwise classes, and conditions outside the study (for example, missing teeth, crooked teeth, gaps, ulcers, and missing tooth buds) were left out. Some collections outlined each condition as a shape instead of a box, so those outlines were turned into boxes, which is the format both models use. Table 15 shows how each collection's conditions were matched.

Table 15 Source Datasets and Class Remapping to the Toothwise Taxonomy

Source dataset (export)	Original classes	Mapped to Toothwise class	Dropped classes
dadad12 (v3)	calculus, cavities, maligned, missing,	calculus → tartar; cavities → cavity; plaque → plaque	maligned, missing, spacing, stains

Source dataset (export)	Original classes	Mapped to Toothwise class	Dropped classes
	plaque, spacing, stains		
DenPal (v8)	Caries, Missing Tooth, Tartar	Caries → cavity; Tartar → tartar	Missing Tooth
oral-diseases (v2)	calculus, caries, gingivitis, hypodontia, tooth_discolation, ulcer	calculus → tartar; caries → cavity; gingivitis → gum_inflammation	hypodontia, tooth_discolation, ulcer
Teeth 2 (v1)	Caries, Cavity, Healthy	Caries → cavity; Cavity → cavity; Healthy → healthy_tooth	None

Source: class lists and remap printed by the YOLOv12 training notebook (train_dental_yolov12_5class_2.ipynb).

The combined collection was split into three parts (Table 16): a training set that the models learn from, a validation set used to check progress during training, and a test set that is kept apart for the final evaluation. Photographs that showed none of the five conditions were kept, because they teach the models not to mark healthy areas by mistake.

Table 16 Merged Dataset Split

Split	Images	Background images (no target label)	Share of images
Training	40,271	10,791	83.0%
Validation	5,299	2,833	10.9%
Test	2,930	1,067	6.0%
Total	48,500	14,691	100.0%

Label rows written during the merge: 86,028 bounding boxes and 7,457 polygons converted to boxes (93,485 in total). Test background count taken from the rebuilt test split.

Removing duplicate photographs. Some of the photo collections were taken partly from the same sources, so some test photographs were copies of training photographs. A model tested on a photograph it has already learned from would look better than it really is. Every test photograph was therefore compared with every training and validation photograph, looking for identical files, edited copies of the same original photograph, and near-identical images. Of the 2,930 test photographs, 774 (26.4%) were found in the training set and 336 (11.5%) in the validation set. All 849 of them were

removed, leaving a **clean test set of 2,081 photographs**. All results in this section use the clean test set. Table 17 shows the test set before and after.

Table 17 Test Set Before and After Duplicate Removal

Photographs showing	Full test set	Clean test set
Healthy tooth	39	34
Cavity	1,190	850
Gum inflammation	297	84
Plaque	147	98
Tartar	558	345
Any of the four conditions	1,826	1,204
No labeled condition (background)	1,067	845
Total images	2,930	2,081
Labeled instances	6,865	3,733

An image can contain more than one condition, so the condition rows do not add up to the total.

Gum inflammation lost the most photographs: 213 of its 297 test photographs were copies, so only 84 remain.

4.1.2 Training Configuration

Both models were trained on the same data with the same settings, so that the comparison between them is fair. Training ran on Google Colab with an NVIDIA Tesla T4 graphics card. Table 18 lists the settings.

Table 18 Training Configuration of the Two Detection Models

Parameter	YOLOv8	YOLOv12
Model scale	YOLOv8m (medium; depth 0.67, width 0.75)	YOLOv12m (medium)
Trainable parameters	25,859,215	19,662,111
Pretrained starting weights	Yes (COCO)	Yes (COCO)
Classes	5	5
Epochs	60	60
Image size	640 × 640 px	640 × 640 px
Batch size	12	12

Parameter	YOLOv8	YOLOv12
Optimizer	auto (selected by Ultralytics)	auto (MuSGD, lr 0.01, momentum 0.9)
Early-stopping patience	15 epochs	15 epochs
Random seed	0	0
Ultralytics version	8.4.146	8.4.160
Training hardware	Google Colab, NVIDIA Tesla T4 (15 GB)	Google Colab, NVIDIA Tesla T4 (15 GB)
Checkpoint date (UTC)	2026-09-10	2026-09-23

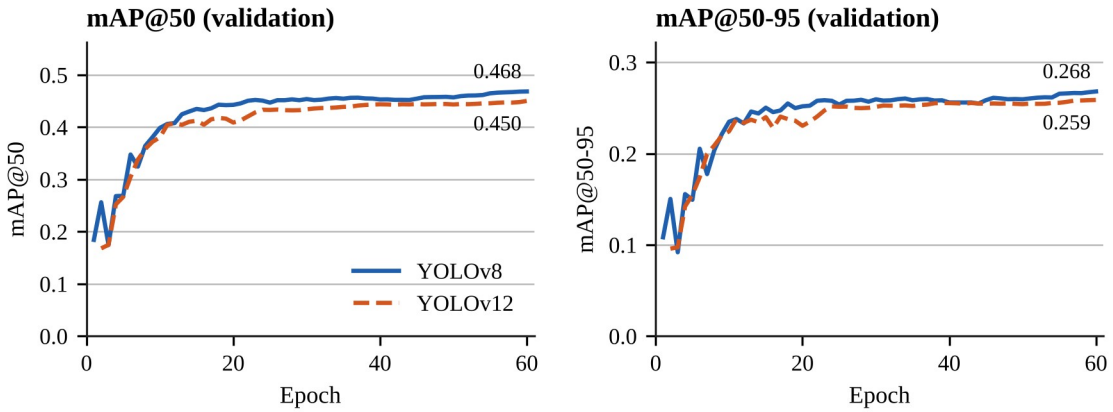


Figure 28 mAP@50 and mAP@50-95 per Training Epoch for YOLOv8 and YOLOv12.

Figure 28 shows how the scores grew during training. Both models improved quickly in the first ten rounds (epochs) and then levelled off. This means that training longer would not help much; better results need more and better training photographs.

4.1.3 How the Models Were Scored

The main scores are given per photograph. For each photograph and each condition, two things are compared: does the photograph show the condition, and did the model report it? This gives four possible outcomes:

- **True positive (TP):** the photograph shows the condition and the model reported it (a correct find).

- **False positive (FP):** the photograph does not show the condition, but the model reported it (a false alarm).
- **False negative (FN):** the photograph shows the condition, but the model did not report it (a missed finding).
- **True negative (TN):** the photograph does not show the condition and the model did not report it (correctly left unmarked).

Six scores are calculated from these four counts:

- **Accuracy:** how many photographs were judged correctly.
- **Precision:** when the model says a condition is there, how often it is right.
- **Recall:** of the photographs that show the condition, how many the model found.
- **Specificity:** of the photographs without the condition, how many the model correctly left unmarked.
- **F1-score:** one number that combines precision and recall.
- **AUROC:** how well the model's confidence separates photographs with the condition from photographs without it, without depending on a cut-off.

Any finding asks whether the photograph shows at least one of the four conditions (cavity, gum inflammation, plaque, or tartar). This is the main question a screening tool answers. A model reports a condition when its confidence reaches a cut-off. For each condition, the cut-off that gave the best F1-score on the validation set was chosen and then used unchanged on the test set. The live system currently uses fixed cut-offs (0.25, and 0.50 for cavity), and results at those cut-offs are also shown.

The standard YOLO score, mAP@50, is stricter: it counts a finding only when the model's box overlaps the hand-drawn box by at least half. It is given as a supporting result, because the job of Toothwise is to tell the dentist whether a condition is present, not to draw its exact outline.

4.1.4 Results for Any Finding

Table 19 gives the main result of this study.

Table 19 Image-Level Performance for Any Finding (Clean Test Set, Validation-Tuned Cut-offs)

Model	Accuracy	Precision	Recall	Specificity	F1-score	AUROC
YOLOv8	84.8%	87.9%	85.5%	83.9%	86.7%	90.0%
YOLOv12	84.3%	87.5%	85.0%	83.4%	86.3%	89.8%
Fused ensemble	85.3%	86.7%	88.0%	81.5%	87.4%	90.7%

N = 2,081 photographs, of which 1,204 show at least one condition. Cut-offs: 0.20 (YOLOv8), 0.20 (YOLOv12), 0.30 (ensemble), chosen on the validation set. Source: evaluate_image_level.py.

All three scored above 80% on all six measures. In simple terms, of every 100 photographs that showed a dental problem, YOLOv8 found about 86, YOLOv12 about 85, and the combined models 88. When a model said that a problem was present, it was right about 87 to 88 times out of 100. YOLOv8 was slightly better than YOLOv12, by less than one percentage point on every measure. The combined models found the most problems, with slightly more false alarms.

Table 20 shows the results at the cut-offs that the live system uses today.

Table 20 Image-Level Performance at the Deployed Cut-offs (Clean Test Set)

Model	Condition	Accuracy	Precision	Recall	Specificity	F1-score
YOLOv8	Any finding	83.6%	92.4%	78.0%	91.2%	84.6%
YOLOv12	Any finding	82.4%	91.9%	76.2%	90.8%	83.3%
Fused ensemble	Any finding	82.8%	92.5%	76.5%	91.4%	83.7%
YOLOv8	Cavity	87.7%	90.7%	77.9%	94.5%	83.8%
YOLOv12	Cavity	87.0%	89.8%	76.9%	94.0%	82.9%
Fused ensemble	Cavity	86.4%	93.1%	71.9%	96.3%	81.1%

Deployed cut-offs: 0.25, and 0.50 for cavity, applied to each model before the two are combined; as in the live system, the combined result keeps a cavity only when both models find it. AUROC does not depend on the cut-off and is given in Tables 19 and 21.

At these cut-offs the models are more careful: when they report a problem, they are right about 90% to 93% of the time, but they miss more (recall 71.9% to 78.0%). The live system also keeps a cavity only when both models find it, which makes its cavity

findings the most precise (93.1%) but lowers its cavity recall to 71.9%. Because a missed finding matters more than a false alarm in screening, the cut-offs chosen on the validation set (Table 19), which keep recall above 85%, are recommended for the live system. The braces filter of the live system (Section 4.2.1) is not included in these results.

4.1.5 Results per Condition

Table 21 shows the results for each condition with the two models combined, and Table 22 compares the F1-score of YOLOv8, YOLOv12, and the combined models. The full tables for each model are in Appendix B.

Table 21 Image-Level Performance per Condition, Fused Ensemble (Clean Test Set)

Condition	Photos	Accuracy	Precision	Recall	Specificity	F1-score	AUROC
Healthy tooth	34	98.8%	60.9%	82.4%	99.1%	70.0%	98.2%
Cavity	850	89.0%	87.6%	85.1%	91.7%	86.3%	93.8%
Gum inflammation	84	96.1%	51.7%	53.6%	97.9%	52.6%	93.5%
Plaque	98	95.2%	48.8%	20.4%	98.9%	28.8%	61.9%
Tartar	345	92.0%	75.8%	76.2%	95.2%	76.0%	90.4%
Any finding	1,204	85.3%	86.7%	88.0%	81.5%	87.4%	90.7%

Photos: test photographs that show the condition (out of 2,081). Cut-offs chosen on the validation set.

Table 22 Image-Level F1-Score per Condition for Each Model (Clean Test Set)

Condition	YOLOv8	YOLOv12	Fused ensemble
Healthy tooth	69.3%	70.1%	70.0%
Cavity	86.4%	85.5%	86.3%
Gum inflammation	59.3%	41.0%	52.6%
Plaque	31.1%	22.6%	28.8%
Tartar	75.2%	73.7%	76.0%
Any finding	86.7%	86.3%	87.4%

Cavity, the most serious condition, scored above 80% on all six measures for all three (see also Appendix B). **Tartar** came close, with an F1-score of about 74% to 76%. **Healthy tooth** scored well, but there are only 34 such test photographs, so that result is

less certain. **Gum inflammation** (F1-score 41% to 59%) and **plaque** (23% to 31%) were weak; the models missed most plaque photographs.

Plaque is the main weakness. Plaque is a thin, pale film that is hard to tell apart from the tooth and from tartar in an ordinary photograph; only one of the four photo collections marked plaque at all; and there are far fewer plaque examples than cavity examples. These are the same reasons why plaque is hard to judge by eye. Gum inflammation has only 84 clean test photographs, so its results are less certain.

Accuracy should be read together with the other scores. Most photographs do not show a given rare condition, so a model that rarely reports it still scores a high accuracy: plaque has an accuracy above 95% even though most plaque photographs are missed. This is why precision, recall, and F1-score are shown for every condition.

Table 23 gives the four counts behind these scores, so that the number of correct finds, false alarms, and missed findings can be seen directly.

Table 23 Image-Level True and False Positives and Negatives (Clean Test Set)

Model	Condition	True positives (found)	False positives (false alarm)	False negatives (missed)	True negatives (correctly clear)
YOLOv8	Any finding	1,029	141	175	736
YOLOv12	Any finding	1,024	146	180	731
Fused ensemble	Any finding	1,060	162	144	715
YOLOv8	Cavity	736	118	114	1,113
YOLOv12	Cavity	736	135	114	1,096
Fused ensemble	Cavity	723	102	127	1,129
Fused ensemble	Healthy tooth	28	18	6	2,029
Fused ensemble	Gum inflammation	45	42	39	1,955
Fused ensemble	Plaque	20	21	78	1,962
Fused ensemble	Tartar	263	84	82	1,652

N = 2,081 clean test photographs; each row adds up to 2,081. Cut-offs chosen on the validation set, as in Tables 19 and 21.

For any finding, the combined models correctly found a problem in 1,060 of the 1,204 photographs that showed one and missed 144, while raising 162 false alarms among the 877 photographs without a problem. For cavity, they found 723 of 850 cavity photographs, missed 127, and raised 102 false alarms. The counts also show the weakness in plaque plainly: of 98 plaque photographs, only 20 were found and 78 were missed. For screening, a false alarm is less serious than a miss, because the dentist checks every finding and can dismiss a false alarm, while a missed finding might not be looked at.

4.1.6 Box-Level (Standard YOLO) Performance (Supporting Results)

Table 24 gives the standard YOLO box-level scores at the end of training (validation set) and on the clean test set.

Table 24 Box-Level (Standard YOLO) Performance of YOLOv8 and YOLOv12

Metric	YOLOv8 (validation)	YOLOv12 (validation)	YOLOv8 (clean test)	YOLOv12 (clean test)
Precision	54.0%	49.4%	51.1%	53.3%
Recall	49.7%	49.1%	50.9%	48.1%
F1-score	51.8%	49.2%	51.0%	50.6%
mAP@50	46.8%	45.0%	50.0%	49.8%
mAP@50-95	26.8%	25.9%	29.1%	28.7%

Validation: scores saved in the models at the end of training (epoch 60). Clean test: evaluate_models.py on the 2,081 clean test photographs. Before the duplicates were removed, mAP@50 on the test set was 51.9% (YOLOv8) and 48.1% (YOLOv12).

Both models scored about 50% mAP@50 on the clean test set: YOLOv8 50.0% and YOLOv12 49.8%. In simple terms, about half of the hand-drawn boxes were matched by a box from the model. This is much lower than the per-photograph scores because the box must be placed almost exactly, and for conditions with soft edges, such as plaque and gum inflammation, even two people marking the same photograph would draw different boxes. Removing the duplicates changed mAP@50 by less than two percentage points. YOLOv12 reached its score with about 24% fewer parameters than YOLOv8 (19.66 million against 25.86 million), which makes it lighter to run.

4.1.7 Inference Speed

Speed was measured without a graphics card in two ways (Table 25): with the original PyTorch models on an ordinary laptop, like the computers in a clinic, and with the lighter ONNX models that the live system runs, limited to two processor threads as on the live server.

Table 25 Inference Time per Photograph on CPU

Configuration	Mean time per photograph	Standard deviation	Model forward pass only
YOLOv8, PyTorch (laptop)	215.9 ms	8.9 ms	203.7 ms
YOLOv12, PyTorch (laptop)	206.5 ms	7.6 ms	198.9 ms
Both models in sequence, PyTorch (laptop)	506.8 ms	138.7 ms	—
YOLOv8, ONNX (as in the live system)	657.8 ms	81.3 ms	650.0 ms
YOLOv12, ONNX (as in the live system)	720.7 ms	90.7 ms	712.9 ms
Both models in sequence, ONNX (as in the live system)	1,399.0 ms	166.3 ms	—

Laptop: Intel Core i7-13620H, 16 GB RAM, CPU only (PyTorch 2.14 CPU, Ultralytics 8.4.116); n = 20 timed predictions per configuration. ONNX: the optimized ONNX models and the system's own ONNX Runtime 1.31 code with 2 threads, as on the live server, on a 4-core Intel Xeon 2.8 GHz cloud processor, measured on 3 November 2026; n = 40 timed predictions per configuration (4 photographs, 5 runs each after one warm-up, in two separate runs).

On the laptop, each model analyzed a photograph in about one fifth of a second, and running both took about half a second. The ONNX models, limited to two threads, took about 0.7 seconds per model and 1.4 seconds for both. Either is fast enough for a dentist to upload a photograph and see the results right away.

The live system runs on a smaller, free cloud server with only 512 MB of memory. To fit, it uses a lighter copy of the same two models (ONNX format) and loads each model only while it is working. On that server, a full screening with both models takes about 3 to 4 seconds, and the system uses about 210 MB of memory at rest and about 400 MB at its busiest. The lighter models gave exactly the same results as the originals on 22 test photographs, so the accuracy results in this chapter also apply to the live system. A

wait of a few seconds is acceptable because a photograph is uploaded only once per screening.

4.1.8 Combining the Two Models

In the live system, both models look at every photograph and their findings are combined. When both models mark the same condition in about the same place, the two marks become one, and the finding gets a slightly higher confidence (plus 0.08, up to 0.98), because two models agreeing is stronger evidence than one.

The combined models found the most problems: for any finding, their recall was 88.0% against 85.5% for YOLOv8, with the best F1-score (87.4%) and AUROC (90.7%). They also had the highest AUROC of the three for every condition (Table 21 and Appendix B). The cost is slightly more false alarms (specificity 81.5% against 83.9% for YOLOv8). Since a missed finding matters more in screening than a false alarm that the dentist can dismiss, this supports running both models.

4.1.9 Summary of the Model Comparison

To answer the first problem: on 2,081 clean test photographs, YOLOv8, YOLOv12, and the combined models all scored above 80% on all six per-photograph measures for finding any dental problem and for finding cavities. Tartar came close, while gum inflammation and plaque were weak because of too few and inconsistently marked training photographs. YOLOv8 was slightly better than YOLOv12, YOLOv12 was smaller and slightly faster, and combining them found the most problems. On the stricter standard YOLO score, both reached about 50% mAP@50. Toothwise is therefore a screening aid, not a replacement for the dentist: its findings are suggestions that the dentist confirms or corrects.

4.2 Specialization and Technical Results

This section describes how the finished system works and how well it performed: the screening steps, the risk score, the prevention plan, the patient records, user roles and security, and the chatbot.

4.2.1 AI Detection Pipeline

When staff upload a photograph of a patient's teeth, the system goes through six steps: (1) it checks that the file is a HEIC/HEIF, JPEG, PNG, or WebP image of at most 15 MB, and converts HEIC photographs, the default camera format of iPhones, to a full-resolution JPEG; (2) it runs YOLOv8 and YOLOv12; (3) it combines their findings (Section 4.1.8) and removes braces brackets that the models mistake for findings; (4) it computes a risk score and risk level; (5) it writes a prevention plan; and (6) it saves the photograph, every finding with its box and confidence, the risk result, and the plan in the patient's record. The live system also keeps a copy of each photograph in the database, because the free server erases its own files whenever it restarts; the copy is used to restore the photograph for viewing and for the PDF report.

Neither model learned what braces look like, so in early testing the metal brackets on teeth with braces were reported as urgent cavities. The system now ignores any cavity, plaque, or tartar box in which 8% or more of the pixels are metal-grey, reports a cavity only when its confidence is at least 0.50, and keeps a cavity only when both models find it. When two or more brackets are seen, the prevention plan adds a note that braces were detected and that a dentist or orthodontist should check those teeth in person. The results in Section 4.1 do not include the braces filter; the cavity rules are included only in Table 20.

4.2.2 Risk Scoring Engine

The risk score follows a fixed rule, so that every score can be explained to the patient and checked by the dentist. Each condition has a severity weight, cavity being the most serious (Table 26). For each finding, the weight is multiplied by the model's

confidence; the score is the highest of these values, plus 6 points for each additional finding (up to five more), with a maximum of 100. The score then decides the risk level: LOW, MEDIUM, or HIGH.

Table 26 Risk Scoring Weights and Tier Thresholds

Class	Severity weight	Highest possible score from one finding
Healthy tooth	0	0 (no risk contribution)
Plaque	25	25 (LOW)
Tartar	35	35 (MEDIUM)
Gum inflammation	45	45 (MEDIUM)
Cavity	70	70 (HIGH)

Tiers: LOW 0–34, MEDIUM 35–69, HIGH 70–100. Source: backend/app/ml/labels.py (SEVERITY_WEIGHT) and backend/app/ml/risk_engine.py.

4.2.3 KNN-Based Recommendation Feature

The prevention plan uses a K-Nearest Neighbors (KNN) classifier. In simple terms, KNN looks at the five past cases that are most similar to the current screening and uses them to decide its risk level; the plan is then written from that level and the specific conditions found. Because there were no dentist-checked cases yet, the classifier first learned from made-up (synthetic) cases created with the risk rule. Each time a dentist confirms or corrects a risk level in the system, that case is saved as a real example for future retraining.

Table 27 shows how the classifier did on 2,000 new made-up cases that it had not seen before. Since these cases follow the same rule it learned from, this test shows how well it copies the rule, not how accurate it is on real patients.

Table 27 KNN Tier Classification on Unseen Synthetic Cases

Risk tier	Precision	Recall	F1-score	Cases
LOW	98.3%	96.3%	97.3%	864
MEDIUM	90.2%	95.5%	92.8%	971
HIGH	76.2%	58.2%	66.0%	165
Overall (weighted)	—	—	92.5%	2,000

Overall accuracy: 92.8% (1,855 of 2,000 cases). Source: python -m app.ml.evaluate_knn, re-run on the current code on 3 November 2026 with the same result.

The classifier agreed with the rule on 92.8% of the cases. Most mistakes happened near the borders between levels. The weakest level was HIGH: 69 of 165 HIGH cases were put in MEDIUM (recall 58.2%). It never mistook LOW for HIGH or HIGH for LOW. Because missing a HIGH case is the most serious mistake, the dentist still checks every risk level. No dentist-checked cases were recorded during the study, because the clinic's patients were not screened, so the classifier's accuracy on real cases could not yet be measured; this is recommended in Chapter V.

4.2.4 Results Interface and Patient Records

To answer the second problem: every screening is now saved as a standard visual record. It contains the original photograph, colored boxes showing each finding with its name and confidence, a plain-language explanation of each condition, the risk score and level, the prevention plan, and the dentist's notes. The patient record lists all screenings by date with a risk trend chart, so visits can be compared, and each screening can be downloaded as a PDF report (Figure 27). Below every result, a table lists the conditions detected in the screening with their image-level evaluation (Figure 22): for each detected condition, its highest confidence and number of marks in this photograph, next to the measured per-photograph scores of the model being viewed from Section 4.1, at the cut-offs the system uses. A condition with an F1-score below 60%, such as gum inflammation or plaque, is flagged to be confirmed by examination; when nothing is detected, the table gives the model's scores for detecting any condition and reminds the user that the models miss about a quarter of problem photographs. Because the dentist's own findings are not recorded per condition, the table shows the expected reliability of each result, not whether that particular result is correct. Figures 18, 21, and 25 to 27 in Section 3.11 show these screens. Patients can also create their own portal account, where they see their latest risk level and how it changed since their last visit, their latest findings explained simply with

the top next steps, a risk-over-time chart, the date of their next check-up (six months after the last screening), and their full screening history, but never another patient's data.

4.2.5 Role-Based Web Application

The staff side has three roles. The dentist is the administrator: the dentist can manage staff accounts and view the audit log. The receptionist registers and manages patients, and the dental assistant performs screenings. Patients have a separate login. The server checks each user's role on every request, not just by hiding menu buttons. Passwords are stored in scrambled (hashed) form, sign-ins use secure tokens, repeated login and chatbot requests are limited, and important staff actions are recorded in an audit log to which records are only added, never changed.

Staff and patients sign in with their email and password. Ticking “Remember me on this device” keeps them signed in for 30 days; otherwise they are signed out when the browser closes. A patient who forgets the password can reset it without email by confirming their email, birthdate, and mobile number, and then choosing a new password on the same page. A wrong answer always shows the same message, so nobody can use the form to find out which accounts exist; attempts are limited to five per minute, and every attempt is recorded in the audit log. Staff passwords are reset only by the dentist, from the Staff page, so a staff account cannot be taken over through the reset form. An earlier version emailed a reset link, but the free server could not send email reliably, so that step was removed. Before anyone can use the system, a Data Privacy Notice explains what data are collected and why, how they are protected, that AI results are not a diagnosis, and the user's rights under the Data Privacy Act of 2012; the user must tick the agreement to continue, once per browser session, so each person on a shared clinic computer sees it (Figure 9). Signed-in users can also change their password from the sidebar or the portal header.

4.2.6 AI Chatbot

The public website has an AI assistant, powered by Google Gemini, that answers general questions about dental care, the clinic, and Toothwise. Before a question is sent to Gemini, email addresses, phone numbers, and record numbers are removed, and any request for personal records or screening results is pointed to the secure patient portal instead. Each user can send up to 12 questions per minute. When the free Gemini quota runs out, questions about the clinic (hours, location, services, and prices) are still answered from the clinic's verified information, and other questions get a “temporarily unavailable” message.

4.3 Software Quality Evaluation Based on ISO/IEC 25010

This section answers the fourth problem: how acceptable is Toothwise? Two groups evaluated the system. Two IT experts rated its software quality on the six ISO/IEC 25010 characteristics using the IT Expert Evaluation Form, and 103 user respondents rated its main features using the user survey from October 19 to 28, 2026. Both forms were answered online (Google Forms) after the respondents agreed to take part. None of the respondents was a patient of CodiSmile Dental Clinic, and no clinic records were used.

4.3.1 Respondents

The two IT experts both specialize in software and web development; one has 1 to 3 years and the other 4 to 6 years of experience in IT, and they work as a senior AI engineer and an IT practitioner (Table 28). The 103 user respondents were 98 patient respondents, who tried the system from a patient's point of view, and 5 IT practitioners (Table 29). Not every part of the user survey was meant for every respondent, so the parts were answered by different numbers of people (Table 30).

Table 28 Profile of the IT Expert Evaluators

Profile item	Category	Frequency
Field of specialization	Software / Web Development	2

Profile item	Category	Frequency
Years of experience in IT	1–3 years	1
Years of experience in IT	4–6 years	1
Current position	Senior AI Engineer	1
Current position	IT	1

N = 2.

Table 29 Distribution of User Respondents

Respondent group	Frequency	Percentage
Patient respondents	98	95.1%
IT practitioners	5	4.9%
Total	103	100.0%

All 103 user respondents rated the three main features. The user survey also had a software quality part (Part B3) meant for dentists, clinic staff, and IT practitioners. No dentist or clinic staff member took part, so it was answered by 12 respondents (5 IT practitioners and 7 patient respondents), with an overall mean of 3.68 (Strongly Agree). Because these respondents were not all IT experts, the software quality results in this section come from the IT Expert Evaluation Form.

Table 30 User Respondents per Survey Part

Survey part	Patient respondents	IT practitioners	Total
AI Detection and Risk Classification	98	5	103
Results Interface and Patient Records	98	5	103
KNN-Based Personalized Prevention Plan	98	5	103
Current Manual Screening Process	2	0	2
Screening Consistency and Time	2	0	2
Security (Part B3)	7	5	12
Privacy (Part B3)	7	5	12
Usability (Part B3)	7	5	12
Reliability (Part B3)	7	5	12
Performance Efficiency (Part B3)	7	5	12
Data Protection (Part B3)	7	5	12

4.3.2 Evaluation Results per Quality Characteristic

Tables 31 to 36 show how the 2 IT experts rated the six quality characteristics. In each table, SA and A are the number of experts who chose Strongly Agree (4) and Agree (3), WM is the average rating, and SD shows how much the two ratings differed (0 means both experts gave the same rating).

Table 31 Evaluation Results for Security

Code	Statement	SA	A	WM	SD	Verbal interpretation
SE1	Only authorized users can log in and access the system's functions.	1	1	3.50	0.50	Strongly Agree
SE2	Each role (Dentist, Receptionist, Dental Assistant, Patient) can only access the features permitted for that role.	1	1	3.50	0.50	Strongly Agree
SE3	Passwords and login sessions are handled securely.	2	0	4.00	0.00	Strongly Agree
SE4	Sensitive staff actions are recorded in the audit log for administrative review.	2	0	4.00	0.00	Strongly Agree
SE5	Login, registration, and password-reset requests are limited against repeated attempts, and a failed password reset does not reveal whether an account exists.	2	0	4.00	0.00	Strongly Agree
	Security — mean			3.80	0.40	Strongly Agree

N = 2 IT experts. SA = Strongly Agree; A = Agree; WM = weighted mean; SD = standard deviation. No statement was rated Disagree or Strongly Disagree.

Security received an average rating of 3.80 out of 4 (SD 0.40), which means Strongly Agree. One expert chose Agree for SE1 and SE2, on who can sign in and what each role may open, so these two statements averaged 3.50, still Strongly Agree; every other statement was rated Strongly Agree by both experts.

Table 32 Evaluation Results for Privacy

Code	Statement	SA	A	WM	SD	Verbal interpretation
PR1	Screening images are not publicly exposed and can only be viewed by authorized users.	2	0	4.00	0.00	Strongly Agree

Code	Statement	SA	A	WM	SD	Verbal interpretation
PR2	Patients can only view their own records through the patient portal.	2	0	4.00	0.00	Strongly Agree
PR3	The system handles personal and health information in line with the Data Privacy Act of 2012.	2	0	4.00	0.00	Strongly Agree
PR4	Users must read and agree to the Data Privacy Notice before they can use the system.	2	0	4.00	0.00	Strongly Agree
PR5	The AI assistant answers general dental and clinic questions without reading or revealing patient records.	2	0	4.00	0.00	Strongly Agree
	Privacy — mean			4.00	0.00	Strongly Agree

N = 2 IT experts. SA = Strongly Agree; A = Agree; WM = weighted mean; SD = standard deviation. No statement was rated Disagree or Strongly Disagree.

Privacy received an average rating of 4.00 out of 4 (SD 0.00), which means Strongly Agree. Both IT experts chose Strongly Agree for all 5 statements.

Table 33 Evaluation Results for Usability

Code	Statement	SA	A	WM	SD	Verbal interpretation
US1	The system is easy to learn and use, even for first-time users.	2	0	4.00	0.00	Strongly Agree
US2	Navigation between the dashboard, patient records, and screening pages is clear and logical.	2	0	4.00	0.00	Strongly Agree
US3	Labels, buttons, and messages are easy to understand.	2	0	4.00	0.00	Strongly Agree
US4	The layout and design of the interface are visually pleasing and organized.	2	0	4.00	0.00	Strongly Agree
US5	The system works well in both desktop and mobile web browsers.	2	0	4.00	0.00	Strongly Agree
	Usability — mean			4.00	0.00	Strongly Agree

N = 2 IT experts. SA = Strongly Agree; A = Agree; WM = weighted mean; SD = standard deviation. No statement was rated Disagree or Strongly Disagree.

Usability received an average rating of 4.00 out of 4 (SD 0.00), which means Strongly Agree. Both IT experts chose Strongly Agree for all 5 statements. The experts suggested fewer clicks and less scrolling, less white space, and a layout for all screen sizes (Appendices C and D).

Table 34 Evaluation Results for Reliability

Code	Statement	SA	A	WM	SD	Verbal interpretation
RE1	The system runs consistently without unexpected crashes or failures.	2	0	4.00	0.00	Strongly Agree
RE2	When an error occurs, the system shows a clear message and recovers without losing data.	2	0	4.00	0.00	Strongly Agree
RE3	The system produces consistent results when the same image is screened again.	2	0	4.00	0.00	Strongly Agree
RE4	Screening photos and records remain available after the server restarts.	2	0	4.00	0.00	Strongly Agree
RE5	After the server wakes from an idle period, the system works normally without errors.	2	0	4.00	0.00	Strongly Agree
	Reliability — mean			4.00	0.00	Strongly Agree

N = 2 IT experts. SA = Strongly Agree; A = Agree; WM = weighted mean; SD = standard deviation. No statement was rated Disagree or Strongly Disagree.

Reliability received an average rating of 4.00 out of 4 (SD 0.00), which means Strongly Agree. Both IT experts chose Strongly Agree for all 5 statements. This includes RE5, on the system working normally after the free server wakes from an idle period, a limitation noted in Chapter I.

Table 35 Evaluation Results for Performance Efficiency

Code	Statement	SA	A	WM	SD	Verbal interpretation
PE1	The system analyzes an uploaded image and displays the results within an acceptable time.	2	0	4.00	0.00	Strongly Agree
PE2	Pages and patient records load quickly.	2	0	4.00	0.00	Strongly Agree
PE3	The system remains responsive while performing AI detection.	2	0	4.00	0.00	Strongly Agree
PE4	Photos in common formats (JPEG, PNG, WebP, HEIC) of up to 15 MB are uploaded and processed without errors.	2	0	4.00	0.00	Strongly Agree
PE5	The model-comparison view clearly reports the processing time of YOLOv8 and YOLOv12 for each screening.	2	0	4.00	0.00	Strongly Agree
	Performance Efficiency — mean			4.00	0.00	Strongly Agree

N = 2 IT experts. SA = Strongly Agree; A = Agree; WM = weighted mean; SD = standard deviation. No statement was rated Disagree or Strongly Disagree.

Performance Efficiency received an average rating of 4.00 out of 4 (SD 0.00), which means Strongly Agree. Both IT experts chose Strongly Agree for all 5 statements.

Table 36 Evaluation Results for Data Protection

Code	Statement	SA	A	WM	SD	Verbal interpretation
DP1	Patient, screening, and detection records are stored accurately and completely.	2	0	4.00	0.00	Strongly Agree
DP2	Related records (patients, screenings, findings, recommendations, notes) remain consistent with one another.	2	0	4.00	0.00	Strongly Agree
DP3	The system prevents invalid or incomplete data from being saved.	2	0	4.00	0.00	Strongly Agree
DP4	The generated PDF screening report matches the stored screening record.	2	0	4.00	0.00	Strongly Agree
DP5	The system protects data against loss through database backups and stored copies of screening photos.	2	0	4.00	0.00	Strongly Agree
	Data Protection — mean			4.00	0.00	Strongly Agree

N = 2 IT experts. SA = Strongly Agree; A = Agree; WM = weighted mean; SD = standard deviation. No statement was rated Disagree or Strongly Disagree.

Data Protection received an average rating of 4.00 out of 4 (SD 0.00), which means Strongly Agree. Both IT experts chose Strongly Agree for all 5 statements.

4.3.3 Evaluation Results for the System Features

Tables 37 to 39 show the results for the system's three main features, rated by all 103 user respondents.

Table 37 Evaluation Results for AI Detection and Risk Classification

Code	Statement	WM	SD	Verbal interpretation
AD1	Toothwise accurately detects tartar, plaque, cavity, gum inflammation, and healthy tooth from a still intraoral photograph.	3.48	0.54	Strongly Agree
AD2	Combining the YOLOv8 and YOLOv12 outputs produces dependable detection results.	3.43	0.53	Strongly Agree
AD3	The computed risk score and risk classification (Low, Medium, High) reflect the detected findings appropriately.	3.50	0.50	Strongly Agree

Code	Statement	WM	SD	Verbal interpretation
AD4	The system works well with photos taken using a standard device camera such as a smartphone or webcam.	3.47	0.52	Strongly Agree
AD5	The detection results are a helpful reference that supports, but does not replace, the dentist's judgment.	3.40	0.56	Strongly Agree
	AI Detection and Risk Classification — mean	3.45	0.53	Strongly Agree

N = 103 (98 patient respondents, 5 IT practitioners). WM = weighted mean; SD = standard deviation.

AI Detection and Risk Classification received an average rating of 3.45 out of 4 (SD 0.53), which means Strongly Agree. The best-rated statement was AD3, “The computed risk score and risk classification (Low, Medium, High) reflect the detected findings appropriately” (3.50), and the lowest was AD5, “The detection results are a helpful reference that supports, but does not replace, the dentist's judgment” (3.40). Only 9 of the 515 answers in this section were Disagree or Strongly Disagree. Note that this rating shows how accurate the detection seemed to the respondents, not how accurate it measured: the measured accuracy in Section 4.1 is moderate, and plaque is often missed. One respondent wrote that results may be wrong when the photograph is blurry and should still be checked by a dentist, which is how the system is meant to be used.

Table 38 Evaluation Results for Results Interface and Patient Records

Code	Statement	WM	SD	Verbal interpretation
RI1	The bounding boxes and condition labels clearly show where each finding is located on the image.	3.49	0.50	Strongly Agree
RI2	The plain-language descriptions make the findings easy to understand.	3.45	0.52	Strongly Agree
RI3	The per-patient record makes it easy to review and compare findings across visits.	3.48	0.54	Strongly Agree
RI4	The downloadable PDF screening report is complete and useful for clinical recordkeeping.	3.47	0.55	Strongly Agree
RI5	The self-service patient portal lets patients conveniently view their own screening history and past findings.	3.48	0.59	Strongly Agree
	Results Interface and Patient Records — mean	3.47	0.54	Strongly Agree

N = 103 (98 patient respondents, 5 IT practitioners). WM = weighted mean; SD = standard deviation.

Results Interface and Patient Records received an average rating of 3.47 out of 4 (SD 0.54), which means Strongly Agree. The best-rated statement was RI1, “The bounding boxes and condition labels clearly show where each finding is located on the image” (3.49), and the lowest was RI2, “The plain-language descriptions make the findings easy to understand” (3.45). Only 7 of the 515 answers in this section were Disagree or Strongly Disagree. In their written comments, respondents especially liked the plain-language explanations and how clearly the results are shown.

Table 39 Evaluation Results for KNN-Based Personalized Prevention Plan

Code	Statement	WM	SD	Verbal interpretation
PP1	The prevention plan is relevant to the patient's detected findings and risk classification.	3.54	0.50	Strongly Agree
PP2	The recommendations are clear and easy for patients to follow.	3.47	0.52	Strongly Agree
PP3	The prevention plan encourages earlier preventive action and better oral health awareness.	3.51	0.52	Strongly Agree
PP4	The prevention plan appropriately serves as guidance and not as a formal medical prescription.	3.49	0.50	Strongly Agree
	KNN-Based Personalized Prevention Plan — mean	3.50	0.51	Strongly Agree

N = 103 (98 patient respondents, 5 IT practitioners). WM = weighted mean; SD = standard deviation.

KNN-Based Personalized Prevention Plan received an average rating of 3.50 out of 4 (SD 0.51), which means Strongly Agree. The best-rated statement was PP1, “The prevention plan is relevant to the patient's detected findings and risk classification” (3.54), and the lowest was PP2, “The recommendations are clear and easy for patients to follow” (3.47). Only 2 of the 412 answers in this section were Disagree or Strongly Disagree. This was the best-rated of the three features. Respondents wrote that the tips were clear and reminded them to care for their teeth, and several asked for more dental tips, check-up reminders, and short videos.

4.3.4 Summary of the Evaluation

Table 40 Summary of Evaluation Results

Section	Rated by	WM	SD	Interpretation
Security	2 IT experts	3.80	0.40	Highly Acceptable
Privacy	2 IT experts	4.00	0.00	Highly Acceptable
Usability	2 IT experts	4.00	0.00	Highly Acceptable
Reliability	2 IT experts	4.00	0.00	Highly Acceptable
Performance Efficiency	2 IT experts	4.00	0.00	Highly Acceptable
Data Protection	2 IT experts	4.00	0.00	Highly Acceptable
Quality characteristics (mean of six)	2 IT experts	3.97	0.18	Highly Acceptable
AI Detection and Risk Classification	103 user respondents	3.45	0.53	Highly Acceptable
Results Interface and Patient Records	103 user respondents	3.47	0.54	Highly Acceptable
KNN-Based Personalized Prevention Plan	103 user respondents	3.50	0.51	Highly Acceptable
System features (mean of three)	103 user respondents	3.48	0.53	Highly Acceptable

All nine sections were rated Strongly Agree, so Toothwise is Highly Acceptable to its evaluators. The IT experts rated the six quality characteristics 3.97 out of 4: 58 of their 60 ratings were Strongly Agree and the other 2 were Agree, both on Security. The user respondents rated the three features 3.48, and only 18 of their 1,442 answers were Disagree or Strongly Disagree.

Two cautions apply. First, the software quality results come from 2 IT experts, so they describe the judgment of a small group. Second, the ratings show what the evaluators thought of the system; they should be read together with the measured results in Sections 4.1 and 4.2.

In their written comments, the IT experts described the system as innovative and praised that it keeps the clinic's records in one place, shows each user only what they need, and is simple enough to use without much training. Neither reported a problem, bug, or security concern. They recommended fewer clicks and less scrolling, and noted that the web-based design is a good base for features such as billing and inventory.

The user respondents most often praised how easy the system is to use, its clear design, and its simple explanations. The problems they mentioned were slow loading the first time, pages with too much information at once, some features being hard to find, and wrong results from blurry photographs. Their suggestions were reminders for brushing and check-ups, more dental tips and short videos, an easier search, more conditions to detect with better accuracy, and an extra security step for staff accounts.

4.4 Speed and Reliability of the System

This section answers the fifth problem: is Toothwise fast and reliable enough for daily use in a clinic? It uses three kinds of measured evidence: the time the system takes to complete a screening, the results of the automated tests of the system's code, and an audit of the website's speed and accessibility. The respondents' ratings are added at the end as supporting information.

4.4.1 Screening Speed

On an ordinary computer without a graphics card, each model analyzed a photograph in about one fifth of a second, and both models together in about half a second; with the ONNX models the live system uses, limited to two threads, both models took about 1.4 seconds (Table 25). On the free cloud server used for the study, a full screening with both models takes about 3 to 4 seconds and uses at most about 400 MB of memory (Section 4.1.7). This is fast enough to show the results right after the photograph is uploaded. The only noticeable delay is the first request after the free server has been idle for 15 minutes, which can take about a minute while the service starts again (Section 1.4); a paid hosting plan that does not sleep removes this delay.

4.4.2 Reliability

The system's server code is checked by 118 automated tests. They cover signing in, password recovery, the braces filter, the chatbot, staff accounts, the condition list, the detection models, patient records, the prevention plan, the risk score, user roles, and screenings, including restoring photographs from the database. All 118 tests passed (Table 41).

Table 41 Automated Backend Test Results

Test module	Area tested	Test cases	Result
test_auth.py	Staff and patient login, tokens	11	Passed
test_braces_filter.py	Removal of orthodontic brackets from detections	4	Passed
test_chatbot.py	AI assistant retries, fallbacks, and verified clinic facts	6	Passed
test_dentists.py	Staff account management	14	Passed
test_labels.py	Class list consistency	11	Passed
test_onnx_yolo.py	ONNX Runtime letterboxing, NMS, and model loading	4	Passed
test_password_recovery.py	Remember me, patient forgot-password verification, and change password	14	Passed
test_patients.py	Patient records	10	Passed
test_recommender.py	KNN tier and prevention plan	12	Passed
test_risk_engine.py	Risk score and tier	9	Passed
test_roles.py	Three-role access control	11	Passed
test_screenings.py	Screening workflow, storage, and photo restore	12	Passed
Total		118	118 passed, 0 failed

Source: python -m pytest in backend/ (run with YOLO_BACKEND=onnx on 3 November 2026).

The lighter (ONNX) copies of the models that the live system uses gave exactly the same results as the original models on 22 test photographs (Section 4.1.7), so the live system detects the same findings that were evaluated in Section 4.1.

4.4.3 Website Speed and Accessibility

The website's public pages were tested with Google Lighthouse, a free tool that scores speed, accessibility, and good practice from 0 to 100, on both a phone-sized screen (with a slow connection) and a desktop screen. Because the home page was redesigned after the first test, all six tests were run again on October 20, 2026. Table 42 shows the scores.

Table 42 Lighthouse Audit Results

Page	Viewport	Performance	Accessibility	Best practices	Largest Contentful Paint
Landing	Desktop	100	100	100	0.7 s
Landing	Mobile	86	100	100	2.5 s
Staff login	Desktop	99	100	100	0.9 s
Staff login	Mobile	98	100	100	2.2 s
Patient portal	Desktop	99	100	100	0.9 s
Patient portal	Mobile	97	100	100	2.4 s

Landing (mobile) is the median of three runs (67, 86, 92); other rows are single runs. Cumulative Layout Shift was at most 0.014 on every page.

Accessibility and best practices scored a perfect 100 on every page. Speed scored 99 to 100 on desktop and 86 to 98 on phones. The lowest, the home page on a phone, is caused by the amount of code and pictures the redesigned page loads; its score also changed the most between runs. Splitting the website's code into smaller parts cut the main download from 231 KB to 95 KB. The staff pages, which Lighthouse cannot open because they need a login, were checked separately with axe-core against the WCAG 2.0 and 2.1 accessibility rules (Levels A and AA) on all eight staff pages and five pop-up windows, in both light and dark mode; no problems were found.

4.4.4 Respondents' Views (Supporting Information)

The respondents' ratings agree with the measured results. The 2 IT experts rated Performance Efficiency 4.00 and Reliability 4.00, both Highly Acceptable (Section 4.3). Two respondents also answered the sections on the current manual process and on

screening consistency and time (Tables 43 and 44). A timed comparison of the same screenings with and without Toothwise was not possible because the study did not involve the clinic's patients, so these two tables are shown only as supporting information.

Table 43 Respondents' Assessment of the Current Manual Screening Process

Code	Statement	WM	SD	Verbal interpretation
MS1	Early-stage tartar, plaque, cavity, and gum inflammation are difficult to detect consistently through visual inspection alone.	3.00	0.00	Agree
MS2	Findings from visual-only screening can differ between examinations of the same patient.	3.50	0.50	Strongly Agree
MS3	Lighting, viewing angle, and condition stage affect how reliably dental conditions are identified by eye.	3.50	0.50	Strongly Agree
MS4	The clinic lacks a standardized visual record of screening findings that can be compared across visits.	3.00	0.00	Agree
MS5	Patients currently have no convenient way to review their own screening history.	3.50	0.50	Strongly Agree
MS6	Patients rarely receive preventive guidance that is specific to their own screening findings.	2.50	0.50	Disagree
MS7	A tool that supports consistent early screening would be useful to the clinic.	3.50	0.50	Strongly Agree
	Current Manual Screening Process — mean	3.21	0.56	Agree

N = 2. These two sections were answered by only 2 respondents, so the results are shown for completeness and should not be generalized.

Table 44 Respondents' Assessment of Screening Consistency and Time with Toothwise

Code	Statement	WM	SD	Verbal interpretation
SC1	Toothwise makes screening results more consistent than the manual visual-only process.	3.50	0.50	Strongly Agree
SC2	Toothwise reduces the time needed to complete and document a screening.	3.50	0.50	Strongly Agree
SC3	Toothwise's findings agree well with the dentist's own findings.	3.50	0.50	Strongly Agree
SC4	Toothwise fits smoothly into the clinic's existing screening workflow.	3.00	0.00	Agree
SC5	I would recommend the continued use of Toothwise at CodiSmile Dental Clinic.	3.50	0.50	Strongly Agree

Code	Statement	WM	SD	Verbal interpretation
	Screening Consistency and Time — mean	3.40	0.49	Strongly Agree

N = 2. These two sections were answered by only 2 respondents, so the results are shown for completeness and should not be generalized.

The 2 respondents agreed that checking teeth by eye alone has the problems described in the first problem (average 3.21): results can differ from one examination to the next, lighting and angle affect what can be seen, and a tool for consistent early screening would be useful. The only statement they disagreed with was MS6, “Patients rarely receive preventive guidance that is specific to their own screening findings” (2.50). They also agreed that Toothwise makes results more consistent, saves time in completing and recording a screening, and agrees with the dentist's own findings (average 3.40), and they would recommend that the clinic keep using it.

Because only 2 respondents answered these sections and no screening was timed, these results show a likely direction, not a measured effect.

4.4.5 Summary

To answer the fifth problem: Toothwise completes a screening in about 3 to 4 seconds on a small, free cloud server, passed all 118 automated tests, and scored 100 for accessibility and best practices and 86 to 100 for speed on every audited page. The system is therefore fast and reliable enough for daily clinic use, provided it is moved to a hosting plan that does not sleep. Whether it also makes screening more consistent than the manual process should be measured in a future timed comparison with the dentist's findings.

CHAPTER V

SUMMARY OF FINDINGS, CONCLUSIONS, AND RECOMMENDATIONS

This chapter summarizes the findings of the study, states the conclusions drawn from them, and presents recommendations for CodiSmile Dental Clinic, for the further development of Toothwise, and for future researchers.

5.1 Summary of Findings

The findings are summarized according to the five problems stated in Chapter I.

1. **Detection of the five conditions (Problem 1).** YOLOv8 and YOLOv12 were trained on 40,271 images merged from four public datasets. A duplicate check found that 849 of the 2,930 test images also appeared, as copies or near-copies, in the training or validation split; these were removed, and the models were evaluated on the remaining 2,081 clean test images. In the primary, image-level evaluation with validation-tuned thresholds, YOLOv8, YOLOv12, and the fused ensemble each exceeded 80% in accuracy, precision, recall, specificity, F1-score, and AUROC in identifying images that show a dental condition (F1-scores of 86.7%, 86.3%, and 87.4%; AUROC 90.0%, 89.8%, and 90.7%) and images that show a cavity (F1-scores of 86.4%, 85.5%, and 86.3%). At the cut-offs the live system currently uses, precision rose to about 90% to 93% but recall fell to 71.9% to 78.0%, so the validation-tuned cut-offs are recommended. Tartar reached an F1-score of 73.7% to 76.0%, gum inflammation 41.0% to 59.3%, and plaque 22.6% to 31.1%. In the secondary, box-level evaluation on the clean split, YOLOv8 obtained an mAP@50 of 50.0% and YOLOv12 49.8%. YOLOv12 has about 24% fewer parameters and was slightly faster (about 206 ms against 216 ms per image on CPU), and the fused ensemble had the highest recall, F1-score, and AUROC. In the deployed system, a screening takes about 3 to 4 seconds, and metal brackets are filtered out so that they are not reported as cavities. The

respondents who evaluated the current manual process agreed that it has the limitations described in the first problem (mean 3.21).

2. **Recording and access to findings (Problem 2).** Every screening is saved as a standardized visual record with the photograph, labeled bounding boxes, the risk result, the prevention plan, and the dentist's notes; it can be compared across visits, exported as a PDF report, and viewed by the patient through the self-service portal, which now also shows the latest results, a risk-over-time chart, and the next check-up date. Respondents rated the results interface and patient records 3.47, Strongly Agree.
3. **Finding-specific preventive guidance (Problem 3).** The rule-based risk engine and the KNN classifier (K = 5) produce a risk tier and an ordered prevention plan for every screening. On 2,000 unseen synthetic cases, the classifier reproduced the clinical rule with 92.8% accuracy; its weakest tier was HIGH (recall 58.2%). Respondents rated the prevention plan 3.50, Strongly Agree, the highest of the three feature sections.
4. **Acceptability (Problem 4).** The 2 IT experts rated the six ISO/IEC 25010 quality characteristics with an overall mean of 3.97 (Highly Acceptable): every characteristic received 4.00 except Security (3.80), where one expert chose Agree for two statements. The three feature sections, rated by all 103 user respondents, obtained an overall mean of 3.48 (Highly Acceptable).
5. **Speed and reliability (Problem 5).** In the deployed system, a full screening with both models takes about 3 to 4 seconds, and a model analyzes a photograph in about 0.2 seconds on an ordinary computer. All 118 automated backend tests passed. The website scored 100 in Lighthouse accessibility and best practices on all audited pages, and 99 to 100 for speed on desktop and 86 to 98 on phones. The IT experts rated Performance Efficiency 4.00 and Reliability 4.00 (Highly

Acceptable). A timed comparison of screening consistency and time against the manual process was not conducted.

5.2 Conclusions

Based on the findings, the following conclusions are drawn.

1. YOLO-based detection on ordinary intraoral photographs reliably identifies whether a photograph shows a dental condition, and in particular a cavity, with all image-level measures above 80% on unseen, duplicate-free images. It is less reliable for tartar and gum inflammation, unreliable for plaque, and places its boxes exactly on only about half of the hand-marked findings (mAP@50 of about 50%). It is therefore suitable as a screening aid that directs the dentist's attention, not as an independent diagnosis. This confirms the delimitation set in Chapter I.
2. The difference between YOLOv8 and YOLOv12 is small compared with the differences between conditions, so the quality, balance, and consistency of the training data matter more than the choice of architecture. Combining both models found more conditions than either model alone, which supports the ensemble design. Checking the test set for duplicates of training images is necessary when public datasets are merged: about a quarter of the original test images were such duplicates.
3. Storing each screening as a standardized visual record, with patient portal access, addresses the absence of comparable records across visits identified in the second problem.
4. Finding-specific prevention plans can be generated automatically for every screening without adding to the dentist's work, and the respondents found them relevant and easy to follow.

5. Toothwise is highly acceptable to its evaluators on all six quality characteristics and on its main features (overall means of 3.97 and 3.48 on a four-point scale). Because the quality characteristics were rated by 2 IT experts, this conclusion should be confirmed with a larger group of IT experts, dentists, and clinic staff.
6. Toothwise is fast and reliable enough for daily clinic use: it completes a screening in seconds, passed all of its automated tests, and its website is fully accessible. The free hosting tier is the only obstacle, because its first request after a quiet period is slow. Whether Toothwise also makes screening more consistent than the manual process could not be established without a controlled comparison.

5.3 Recommendations

In light of the findings and conclusions, the following are recommended.

For CodiSmile Dental Clinic

1. Use Toothwise as a screening and documentation aid, always with the dentist's review of the detections and confirmation of the risk tier, and take clear, well-lit photographs, since blurry images reduce accuracy.
2. Confirm or correct the risk tier of every screening so that the KNN classifier is retrained on real, dentist-labeled cases.
3. Move the system from the free hosting tier to a plan that does not sleep, so that the first page load is not delayed, before using it in daily operations.

For the further development of Toothwise

1. Improve detection of plaque and other weak classes by adding more and better-annotated photographs, including consented photographs from the clinic, and add a braces class instead of filtering brackets. Before retraining, split the data by original photograph so that copies of one photograph never fall into different splits, and give each photograph one consistent set of labels.

2. Switch the live system to the validation-tuned confidence thresholds reported in Chapter IV, which find more conditions than the current fixed thresholds, after confirming them on clinic photographs.
3. Add the features requested by respondents: reminders for brushing and check-ups, more dental health tips and short videos, a faster and simpler search, and lighter pages for slow Internet connections.
4. Add a second authentication factor for staff accounts, as one respondent suggested, and keep the audit and privacy controls aligned with the Data Privacy Act of 2012.
5. Simplify pages that present much information at once, reduce the clicks and scrolling needed to finish a task, reduce unneeded white space, make the layout adapt to all screen sizes and aspect ratios, and make less-used features easier to find, as the IT experts and user respondents suggested. In the longer term, billing and inventory modules can be added, as one IT expert noted.

For future researchers

1. Conduct a controlled comparison of screening consistency and time: time the same screenings with and without Toothwise and measure agreement with the dentist's findings (for example, with Cohen's kappa).
2. Repeat the ISO/IEC 25010 evaluation with a larger and more balanced group of IT experts, dentists, and clinic staff, the groups for which the quality evaluation was designed.
3. Validate the models on photographs from other clinics and devices, and compare them with newer detection architectures and with segmentation-based approaches.
4. Train the detection models on more and larger datasets. The current models learned from four public datasets in which plaque appears in only one source, which explains much of the weak plaque result. More public datasets, consented

and dentist-annotated clinic photographs, and a more balanced number of examples per condition should raise accuracy.

5. Add an oral diagnosis form to the system: a digital examination form in which the dentist records the clinical findings for each tooth (for example, on a tooth chart), the final diagnosis, and the treatment plan, next to the AI detections. The recorded diagnoses could then be compared with the detections to measure accuracy in real use and supply labeled cases for retraining.
6. Extend the system to detect more classes. The source datasets already contain conditions that this study left out, such as missing teeth, malalignment, spacing, tooth stains and discoloration, oral ulcers, and hypodontia, and orthodontic braces could be added as their own class instead of being filtered out. Each new class needs enough annotated examples, a risk weight, and its own advice.
7. Use an intraoral camera to capture the images for detection. Smartphone photographs vary in lighting and angle and cannot clearly show the back teeth; an intraoral camera gives close, well-lit, and consistent views of each tooth. Future researchers could connect one to the screening page and retrain the models on its images.

REFERENCES

Almajed, O. S., Alayadi, H., & Sabbah, W. (2023). Inequalities in the oral health-related quality of life among children in Saudi Arabia. *Cureus, 15*(11), e49456. https://doi.org/10.7759/cureus.49456

António, J., Malheiro, R., & Jardim, S. (2024). Feedback collection and nearest-neighbor profiling for recommendation systems in healthcare scenarios. *Technologies, 12*(8), 127. https://doi.org/10.3390/technologies12080127

Askar, H., Krois, J., Rohrer, C., Mertens, S., Elhennawy, K., Ottolenghi, L., Mazur, M., Paris, S., & Schwendicke, F. (2021). Detecting white spot lesions on dental photography using deep learning: A pilot study. *Journal of Dentistry, 107*, 103615. https://doi.org/10.1016/j.jdent.2021.103615

Cover, T. M., & Hart, P. E. (1967). Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, *13*(1), 21–27. https://doi.org/10.1109/TIT.1967.1053964

de Brito, D. H. S., dos Santos, T. G. F. T., Costa, S. G. A., dos Santos, A. S., da Silva, I. L. B., da Silva, N. R. C., Fernandes, B. J. T., Niederman, R., Mota, C. C. B. O., da Silveira, M. M. F., Heimer, M. V., & Rosenblatt, A. (2025). Accuracy in diagnosing caries in young permanent molars using interproximal radiographic imaging and validation by artificial intelligence. *Journal of Clinical and Experimental Dentistry*, *17*(6), e648–e655. https://doi.org/10.4317/jced.62396

Etikan, I., Musa, S. A., & Alkassim, R. S. (2016). Comparison of convenience sampling and purposive sampling. *American Journal of Theoretical and Applied Statistics, 5*(1), 1–4. https://doi.org/10.11648/j.ajtas.20160501.11

Gutierrez, K. L. T., & Viacrusis, P. M. L. (2023). Bridging the gap or widening the divide: A call for capacity-building in artificial intelligence for healthcare in the

Philippines. *Journal of Medicine, University of Santo Tomas, 7*(2). https://doi.org/10.35460/2546-1621.2023-0081

Jocher, G., Chaurasia, A., & Qiu, J. (2023). *Ultralytics YOLOv8* (Version 8.0.0) [Computer software]. https://github.com/ultralytics/ultralytics

Kang, S., Shon, B., Park, E. Y., Jeong, S., & Kim, E.-K. (2024). Diagnostic accuracy of dental caries detection using ensemble techniques in deep learning with intraoral camera images. *PLOS ONE, 19*(9), e0310004. https://doi.org/10.1371/journal.pone.0310004

Karakuş, R., Öziç, M. Ü., & Tassoker, M. (2024). AI-assisted detection of interproximal, occlusal, and secondary caries on bite-wing radiographs: A single-shot deep learning approach. *Journal of Imaging Informatics in Medicine, 37*, 3146–3159. https://doi.org/10.1007/s10278-024-01113-x

Liu, Y., Cheng, Y., Song, Y., Cai, D., & Zhang, N. (2024). Oral screening of dental calculus, gingivitis and dental caries through segmentation on intraoral photographic images using deep learning. *BMC Oral Health, 24*(1), 1287. https://doi.org/10.1186/s12903-024-05072-1

Mendoza, M. F., Yacapin, C. P., Alfaro, A. C. A., Ulitin, A. R., Valverde, H. A., Medina, V. O., & Lam, H. Y. (2025). Cost-effectiveness analysis of oral health care package of services within a comprehensive PhilHealth benefit package. *Acta Medica Philippina, 59*(14), 49–59. https://doi.org/10.47895/amp.v59i14.9619

Mohammad-Rahimi, H., Motamedian, S. R., Rohban, M. H., Krois, J., Uribe, S. E., Mahmoudinia, E., Rokhshad, R., Nadimi, M., & Schwendicke, F. (2022). Deep learning for caries detection: A systematic review. *Journal of Dentistry, 122*, 104115. https://doi.org/10.1016/j.jdent.2022.104115

Moharrami, M., Farmer, J., Singhal, S., Watson, E., Glogauer, M., Johnson, A. E. W., Schwendicke, F., & Quiñonez, C. (2024). Detecting dental caries on oral

photographs using artificial intelligence: A systematic review. *Oral Diseases, 30*(4), 1765–1783. https://doi.org/10.1111/odi.14659

Moharrami, M., Vahab, E., Bagherianlemraski, M., Hemmati, G., Singhal, S., Quiñonez, C., Schwendicke, F., & Glogauer, M. (2025). Deep learning for detecting dental plaque and gingivitis from oral photographs: A systematic review. *Community Dentistry and Oral Epidemiology, 53*(6), 617–632. https://doi.org/10.1111/cdoe.70001

Oliveros-Villarico, M., Flores, K. K. R., de Guzman, J. C. Z., Manrique, A. L. V., Millo, G. A., & Estrera, K. R. P. (2023). Pediatrician knowledge, attitude and practices on children's oral health in a tertiary public hospital in the Philippines: A descriptive study. *Acta Medica Philippina, 57*(7), 56–63. https://doi.org/10.47895/amp.vi0.5323

Park, E. Y., Cho, H., Kang, S., Jeong, S., & Kim, E.-K. (2022). Caries detection with tooth surface segmentation on intraoral photographic images using deep learning. *BMC Oral Health, 22*(1), 573. https://doi.org/10.1186/s12903-022-02589-1

Ramírez-Pedraza, A., Salazar-Colores, S., Cardenas-Valle, C., Terven, J., González-Barbosa, J.-J., Ornelas-Rodriguez, F.-J., Hurtado-Ramos, J.-B., Ramírez-Pedraza, R., Córdova-Esparza, D.-M., & Romero-González, J.-A. (2025). Deep learning in oral hygiene: Automated dental plaque detection via YOLO frameworks and quantification using the O'Leary index. *Diagnostics, 15*(2), 231. https://doi.org/10.3390/diagnostics15020231

Sarmiento, R. F. R., Overgaard, S. M., Gai, C., Overgaard, J. D., & Ohde, J. W. (2025). Guiding responsible AI in healthcare in the Philippines. *npj Digital Medicine, 8*, 338. https://doi.org/10.1038/s41746-025-01755-3

Schwendicke, F., Golla, T., Dreher, M., & Krois, J. (2019). Convolutional neural networks for dental image diagnostics: A scoping review. *Journal of Dentistry, 91*, 103226. https://doi.org/10.1016/j.jdent.2019.103226

Tian, Y., Ye, Q., & Doermann, D. (2025). *YOLOv12: Attention-centric real-time object detectors* (arXiv:2502.12524). arXiv. https://doi.org/10.48550/arXiv.2502.12524

Wang, Z., Zhu, Z., Jiang, T., & Gu, Y. (2026). Artificial intelligence for dental caries diagnosis: Translating algorithms to clinical practice. *Frontiers in Medicine, 13*, 1879708. https://doi.org/10.3389/fmed.2026.1879708

World Health Organization. (2022). *Oral health Philippines 2022 country profile*. https://www.who.int/publications/m/item/oral-health-phl-2022-country-profile

World Health Organization. (2025, March 17). *Oral health* [Fact sheet]. https://www.who.int/news-room/fact-sheets/detail/oral-health

Wu, P.-Y., Lin, Y.-J., Chang, Y.-J., Wei, S.-T., Chen, C.-A., Li, K.-C., Tu, W.-C., & Abu, P. A. R. (2025). Deep learning-assisted diagnostic system: Apices and odontogenic sinus floor level analysis in dental panoramic radiographs. *Bioengineering*, *12*(2), 134. https://doi.org/10.3390/bioengineering12020134

Yue, W., Wang, Z., Zhang, J., & Liu, X. (2021). An overview of recommendation techniques and their applications in healthcare. *IEEE/CAA Journal of Automatica Sinica*, *8*(4), 701–717. https://doi.org/10.1109/JAS.2021.1003919

APPENDIX A

Letter of Approval from CodiSmile Dental Clinic

October 26, 2026

DRA. ANGELICA SHALA C. CODILLA
Head Dentist
CodiSmile Dental Clinic
NW Phase LRT Central Terminal, Arroceros St.
Brgy. 659-A, Ermita, Manila City

Dear Dra. Codilla:

Greetings of peace!

We, the undersigned students of the College of Computing Studies, Universidad de Manila, are currently conducting our capstone project entitled **"Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic."** The system analyzes still intraoral photographs to detect tartar, plaque, cavity, gum inflammation, and healthy tooth using the YOLOv8 and YOLOv12 object detection models, computes a risk classification, and generates a personalized prevention plan, while all findings remain subject to the review and final judgment of the clinic's licensed dentist.

In line with this, we respectfully request your approval to conduct our study at CodiSmile Dental Clinic, specifically to:

1. use CodiSmile Dental Clinic as the deployment site and clinical reference of the Toothwise system;
2. demonstrate and pilot the system with the clinic's dentist, dental assistants, and receptionist, using sample dental images only; and
3. consult the clinic's dentist on the system's detection results, risk classification, and prevention plans so that they reflect sound clinical practice.

The study does not involve the clinic's patients. The ISO/IEC 25010 software quality evaluation has already been answered by 50 respondents from outside the clinic, so no patient photograph, record, or personal information from CodiSmile Dental Clinic will be collected. Any information shared with us will be kept strictly confidential and used solely for academic purposes, in accordance with the Data Privacy Act of 2012 (Republic Act No. 10173), and all activities will be scheduled so as not to disrupt the clinic's daily operations.

Figure 29 Letter Requesting Approval to Conduct the Study (page 1 of 2).

Your approval will be a great help to the completion of our study. Thank you very much for your time and kind consideration.

Respectfully yours,

_____	_____
Brayan Axel L. Tubillas	**Robin C. Ponce**
Proponent	Proponent
_____	_____
Keith Ashly V. Sales	**Rizelle Carla N. Magpantay**
Proponent	Proponent

_____
Reizen S. Arreola
Proponent

Noted by:

_____
PROF. FRANCIS ARLANDO L. ATIENZA
Capstone Adviser, College of Computing Studies

Approved by:

_____
DRA. ANGELICA SHALA C. CODILLA
Head Dentist, CodiSmile Dental Clinic

Date: _____

☐ Approved
☐ Not approved

Figure 30 Letter Requesting Approval to Conduct the Study (page 2 of 2).

APPENDIX B

Detailed Image-Level Results per Model

The tables below give the per-condition results of each model on the 2,081 clean test photographs. The results of the two models combined are in Table 21.

Table 45 Image-Level Performance per Condition, YOLOv8 (Clean Test Set)

Condition	Photos	Thr.	Acc.	Prec.	Rec.	Spec.	F1-score	AUROC
Healthy tooth	34	0.60	98.9%	63.4%	76.5%	99.3%	69.3%	96.7%
Cavity	850	0.30	88.9%	86.2%	86.6%	90.4%	86.4%	92.8%
Gum inflammation	84	0.20	96.6%	58.0%	60.7%	98.1%	59.3%	91.7%
Plaque	98	0.05	95.1%	46.0%	23.5%	98.6%	31.1%	61.1%
Tartar	345	0.20	92.2%	79.6%	71.3%	96.4%	75.2%	89.7%
Any finding	1,204	0.20	84.8%	87.9%	85.5%	83.9%	86.7%	90.0%
Mean of the five classes	—	—	94.3%	66.6%	63.7%	96.6%	64.3%	86.4%

Photos: test photographs that show the condition (out of 2,081). Thr.: validation-tuned cut-off. Acc.: accuracy; Prec.: precision; Rec.: recall; Spec.: specificity.

Table 46 Image-Level Performance per Condition, YOLOv12 (Clean Test Set)

Condition	Photos	Thr.	Acc.	Prec.	Rec.	Spec.	F1-score	AUROC
Healthy tooth	34	0.60	98.9%	62.8%	79.4%	99.2%	70.1%	92.3%
Cavity	850	0.25	88.0%	84.5%	86.6%	89.0%	85.5%	93.0%
Gum inflammation	84	0.25	95.6%	44.4%	38.1%	98.0%	41.0%	90.4%
Plaque	98	0.15	95.4%	53.8%	14.3%	99.4%	22.6%	57.2%
Tartar	345	0.15	91.4%	74.3%	73.0%	95.0%	73.7%	88.6%
Any finding	1,204	0.20	84.3%	87.5%	85.0%	83.4%	86.3%	89.8%
Mean of the five classes	—	—	93.8%	64.0%	58.3%	96.1%	58.6%	84.3%

Photos: test photographs that show the condition (out of 2,081). Thr.: validation-tuned cut-off. Acc.: accuracy; Prec.: precision; Rec.: recall; Spec.: specificity.

APPENDIX C

Evaluator Feedback and Panel Consultation

The IT experts received a video demonstration of Toothwise and the evaluation form through Facebook Messenger and gave their feedback in the same conversations; their names and photos are blurred. Messages in Filipino are summarized in English in the captions.

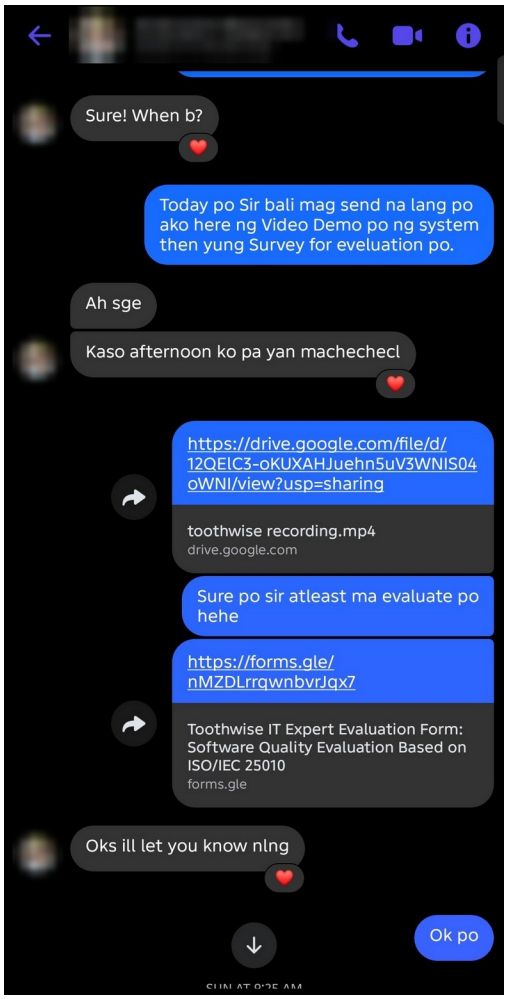


Figure 31 IT Expert 1 — Demonstration and Form Sent. *The researchers sent the demonstration video and the IT Expert Evaluation Form link.*

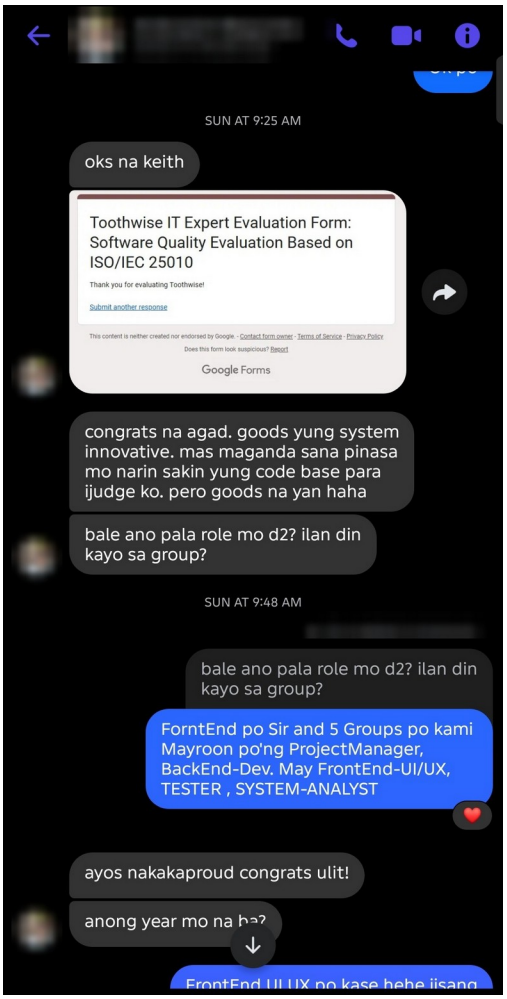


Figure 32 IT Expert 1 — Feedback. *After submitting the form, the expert found the system good and innovative and noted that the code base could also have been shared for review.*

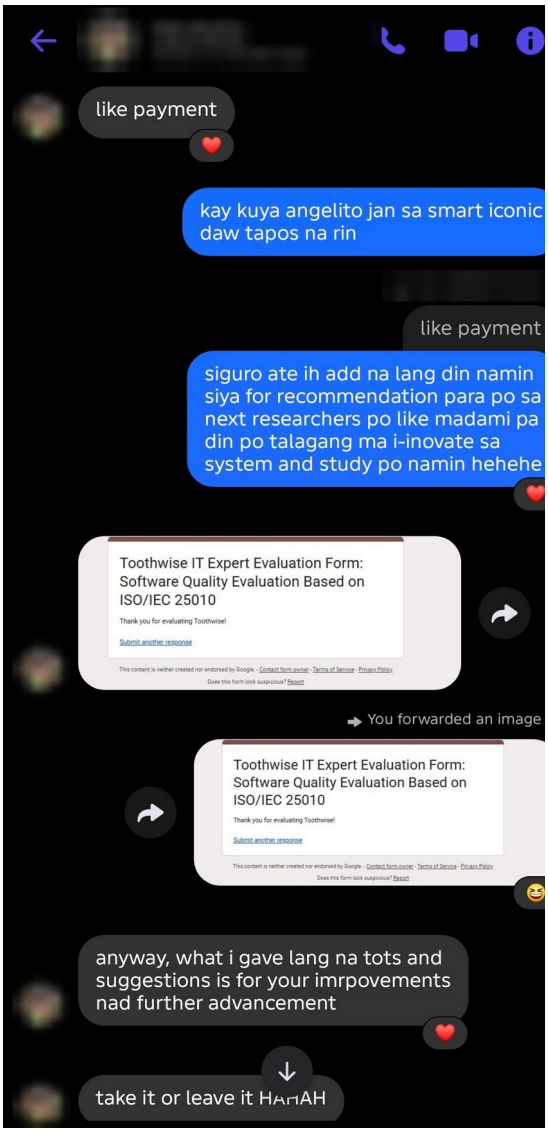


Figure 33 IT Expert 2 — Form Submitted. *The expert submitted the form and offered the comments as suggestions for improvement and further advancement.*

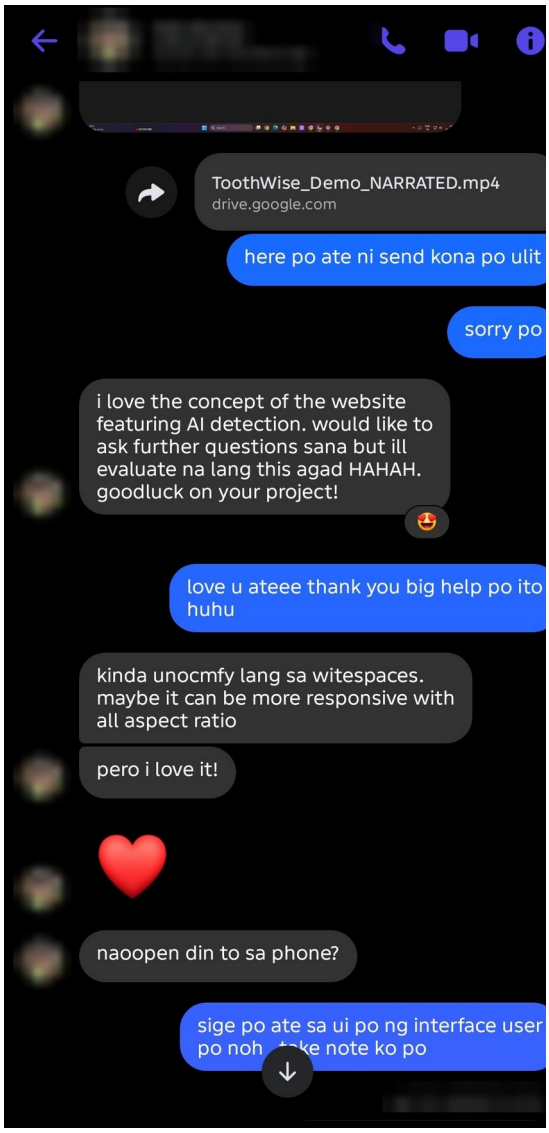


Figure 34 IT Expert 2 — Feedback. *The expert liked the concept of a website with AI detection, found some white space uncomfortable, suggested a layout responsive to all aspect ratios, and asked if it opens on a phone.*

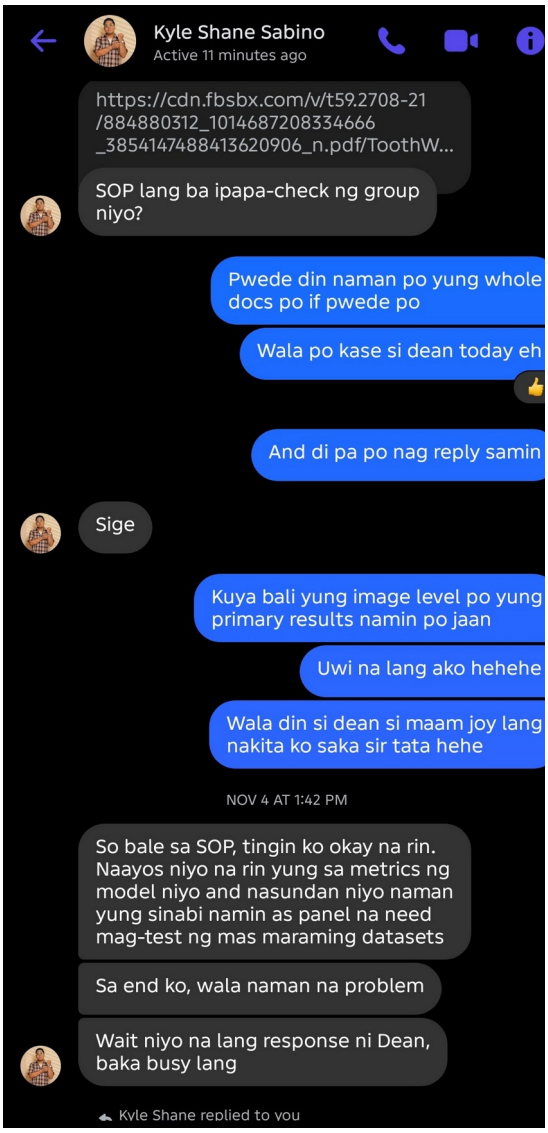


Figure 35 Panel Consultation — Request for Review. *The researchers sent the revised paper to a panel member for checking of the statement of the problem.*

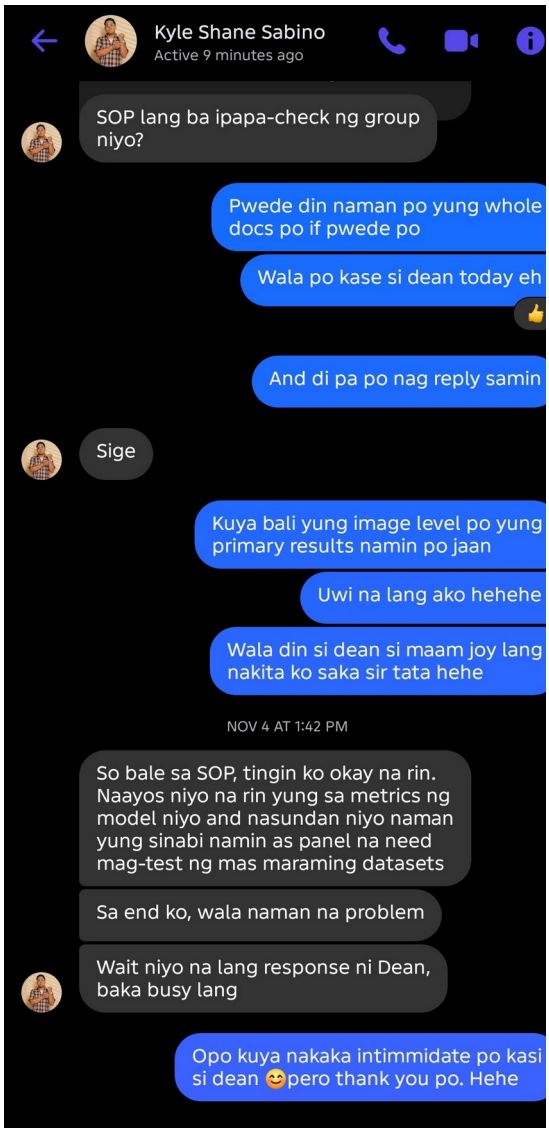


Figure 36 Panel Consultation — Approval. *The panel member found the statement of the problem acceptable, noted that the model metrics were corrected and that the panel's advice to test on more datasets was followed, and saw no problem.*

APPENDIX D

IT Expert Evaluation Results

The table gives, for every statement of the IT Expert Evaluation Form, how many of the IT experts chose each rating and the weighted mean (WM); the statements are listed in Tables 31 to 36. The figures show the response summary from Google Forms, with the evaluators' names blurred, followed by the experts' written comments.

Table 47 Frequency Distribution of the IT Experts' Ratings

Code	Strongly Agree (4)	Agree (3)	Disagree (2)	Strongly Disagree (1)	N	WM
SE1	1	1	0	0	2	3.50
SE2	1	1	0	0	2	3.50
SE3	2	0	0	0	2	4.00
SE4	2	0	0	0	2	4.00
SE5	2	0	0	0	2	4.00
PR1	2	0	0	0	2	4.00
PR2	2	0	0	0	2	4.00
PR3	2	0	0	0	2	4.00
PR4	2	0	0	0	2	4.00
PR5	2	0	0	0	2	4.00
US1	2	0	0	0	2	4.00
US2	2	0	0	0	2	4.00
US3	2	0	0	0	2	4.00
US4	2	0	0	0	2	4.00
US5	2	0	0	0	2	4.00
RE1	2	0	0	0	2	4.00
RE2	2	0	0	0	2	4.00
RE3	2	0	0	0	2	4.00
RE4	2	0	0	0	2	4.00
RE5	2	0	0	0	2	4.00
PE1	2	0	0	0	2	4.00
PE2	2	0	0	0	2	4.00
PE3	2	0	0	0	2	4.00
PE4	2	0	0	0	2	4.00
PE5	2	0	0	0	2	4.00

Code	Strongly Agree (4)	Agree (3)	Disagree (2)	Strongly Disagree (1)	N	WM
DP1	2	0	0	0	2	4.00
DP2	2	0	0	0	2	4.00
DP3	2	0	0	0	2	4.00
DP4	2	0	0	0	2	4.00
DP5	2	0	0	0	2	4.00

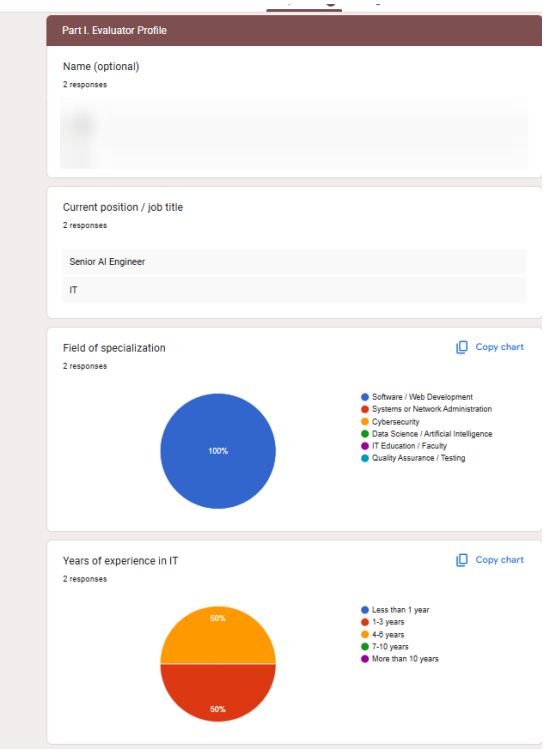


Figure 37 IT Expert Form Results — Evaluator Profile.



Figure 38 IT Expert Form Results — Security, Privacy, and Usability.



Figure 39 IT Expert Form Results — Reliability, Performance Efficiency, and Data Protection.

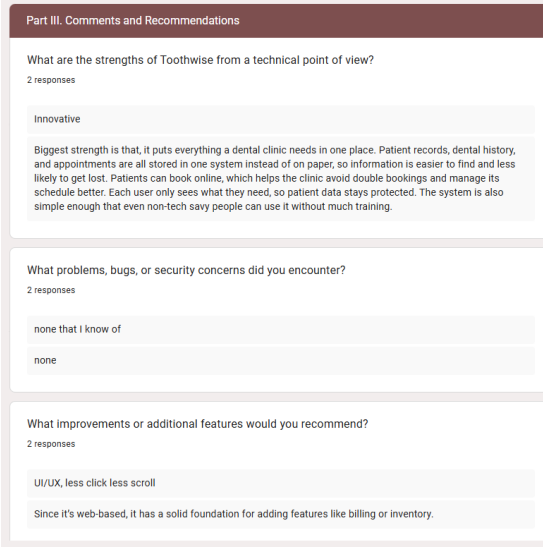


Figure 40 IT Expert Form Results — Comments and Recommendations.

Comments and recommendations, as written

What are the strengths of Toothwise from a technical point of view?

- IT expert 1: “Innovative”
- IT expert 2: “Biggest strength is that, it puts everything a dental clinic needs in one place. Patient records, dental history, and appointments are all stored in one system instead of on paper, so information is easier to find and less likely to get lost. Patients can book online, which helps the clinic avoid double bookings and manage its schedule better. Each user only sees what they need, so patient data stays protected. The system is also simple enough that even non-tech savy people can use it without much training.”

What problems, bugs, or security concerns did you encounter?

- IT expert 1: “none that I know of”

- IT expert 2: “none”

What improvements or additional features would you recommend?

- IT expert 1: “UI/UX, less click less scroll”
- IT expert 2: “Since it’s web-based, it has a solid foundation for adding features like billing or inventory.”