

Universidad De Manila



Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic

A Capstone Project Presented to the College of
Computing Studies
Universidad de Manila

In Partial Fulfillment
of the Requirements for the Course
CAPSTONE 2

Brayan Axel L. Tubillas
Keith Ashly V. Sales
Reizen S. Arreola
Robin C. Ponce
Rizelle Carla N. Magpantay

Prof. Francis Arlando L. Atienza
Professor

September 2026

TABLE OF CONTENTS

CHAPTER I: INTRODUCTION

1.1 Background of the Study 4

1.2 Statement of the Problem.....7

1.3 Objectives of the Study.....8

1.3.1 General Objective 8

1.3.2 Specific Objectives. 8

1.4 Scope, Delimitation and Limitation of the Study 9

1.5 Significance of the Study.....11

1.6 Definition of Terms. 12

CHAPTER II: REVIEW OF RELATED LITERATURE AND STUDIES

2.1 Related Literature. 13

2.2 Related Studies. 18

2.3 Synthesis. 22

CHAPTER III: METHODOLOGY

3.1 Developmental Methodology 24

3.2 Project Planning.....26

3.3 Roles and Responsibilities of the Five-Member Development Team. 28

3.4 Requirement Analysis.....30

3.5 System Design 32

3.5.1 Use Case Diagram 32

3.5.2 Entity Relationship Diagram (ERD) 34

3.5.3 Data Flow Diagram Level 0 37

3.5.4 Data Flow Diagram Level 1 38

3.5.5 Gantt Chart 40

3.6 Software Specifications.....42

3.7 Hardware Specifications.....45

3.8 System Development.....48

3.9 System Testing 50

3.10 Full Current System Screenshots 51

References 57f

CHAPTER I
INTRODUCTION

1.1 Background of the Study

CodiSmile Dental Clinic currently relies on manual visual examination and professional judgment to identify dental caries and other visible dental conditions from patient examinations and intraoral images. Although this approach is essential to dental practice, manually reviewing images may be time-consuming and can result in variations in assessment and documentation. To support the existing process, the proponents propose Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic, an AI-assisted system that analyzes intraoral images and detects possible dental conditions — tartar, plaque, cavity, gum inflammation, and healthy tooth — using the YOLOv12 and YOLOv8 object detection models. The system provides visual detection results, condition labels, descriptions, a computed risk classification, and preventive recommendations to assist dentists and improve patient understanding. Toothwise is intended to serve as a supplementary screening and documentation tool, while final interpretation and clinical decisions remain the responsibility of licensed dental professionals.

Dental caries, commonly referred to as tooth decay or cavities, remains one of the most prevalent non-communicable diseases in the world. The World Health Organization (2024) estimates that oral diseases affect nearly 3.7 billion people globally, with untreated dental caries in permanent teeth identified as the single most widespread health condition based on Global Burden of Disease 2021 data. Despite its prevalence, a significant portion of the population, particularly those in geographically isolated or economically disadvantaged communities, continues to face substantial barriers in accessing timely and adequate dental care. The financial cost of consultations, the distance to dental facilities, and low oral health literacy frequently result in delayed care-seeking, by which point

cavities have often advanced beyond their most treatable stages. According to Qari et al. (2023), untreated dental caries does not just cause physical pain, it also affects a person's confidence and overall quality of life.

A critical dimension of the current dental care challenge lies in the inconsistency of manual cavity screening practices. Wang et al. (2026), in a review on the clinical translation of AI-based caries diagnosis, noted that traditional diagnostic methods suffer from limited sensitivity and high variability between examiners, particularly for early-stage lesions. This inter-examiner variability is particularly problematic in high-volume clinic settings where practitioner fatigue and time pressure can further compromise screening reliability. Research has shown that deep learning models are capable of detecting dental cavities with accuracy that is on par with, or sometimes even better than, trained clinicians, as confirmed by systematic reviews from Schwendicke et al. (2021) and Mohammad Rahimi et al. (2022). This makes AI a genuinely useful tool for supplementing human judgment during screenings.

Standard digital cameras and smartphone cameras have also changed the way dental screening can be supported outside specialized clinical equipment. Instead of relying purely on what a dentist can see with a mirror and a light, or requiring dedicated intraoral camera hardware, clear intraoral photographs can now be captured using widely available consumer devices. Studies such as those by Moharrami et al. (2025) and Askar et al. (2021) show that RGB photographs captured under ordinary imaging conditions are detailed enough to support reliable artificial intelligence-based detection of dental conditions. It is this still-image, standard-camera modality that serves as the primary data source for Toothwise during patient screenings at CodiSmile Dental Clinic, since the system is designed to work from uploaded or device-captured photographs rather than requiring a dedicated intraoral camera.

The AI model powering the primary detection module of this system is YOLOv12, which was introduced by Tian et al. (2025) as an attention-centric object detection framework. Compared to previous versions in the YOLO family, YOLOv12 uses area attention mechanisms that help it focus selectively on spatially relevant regions, something that happens regularly in dental images where carious lesions appear as small, low-contrast areas across multiple tooth surfaces simultaneously. Its combination of speed and accuracy makes it well-suited for efficient use during clinical consultations, where screening results need to be generated quickly enough to fit within a standard patient visit.

To evaluate whether YOLOv12's architectural improvements translate into meaningful gains for this specific screening task, this study also trains and evaluates YOLOv8, an earlier and widely adopted model in the YOLO family (Ultralytics, 2023). YOLOv8 has been applied in several prior dental imaging studies with strong results; de Brito et al. (2025) used YOLOv8 together with EfficientNet-B0 to diagnose caries in interproximal radiographic images of young permanent molars, and other work has applied YOLOv8's anchor-free detection mechanism to interproximal caries detection with improved accuracy and speed over earlier YOLO versions. This makes it a stable, well-benchmarked baseline for object detection tasks. While YOLOv12 and YOLOv8 are formally compared against each other using precision, recall, mAP, and inference speed on a held-out labeled test set, both models are also deployed together in the operational system, where their independent detection outputs are fused into a single clinical detection set. This dual approach allows the study to identify which architecture performs better on this specific task while still taking advantage of both models' detections in the deployed screening workflow.

Beyond detection, Toothwise also generates a personalized prevention plan for each patient based on their screening results. This is done using the K-Nearest Neighbors (KNN) algorithm, a supervised learning method originally introduced by Cover and Hart (1967), which classifies a new case by comparing it against the most similar cases in a

reference dataset. In this study, KNN classifies each screening into a risk tier (low, medium, or high) based on the patient's detected findings — condition type, number of affected areas, and severity indicators derived from the detection output — and this risk classification, together with condition-specific rules, is used to compose an ordered, patient-friendly prevention plan. The reference cases used by the KNN model are validated by a licensed dentist prior to deployment, ensuring that the resulting classifications and recommendations reflect clinically sound guidance rather than unsupervised pattern matching alone.

While AI-assisted dental screening has grown internationally, its application in the Philippine setting remains an emerging area of research. Wu et al. (2025), in collaboration with Ateneo de Manila University's Laboratory for Intelligent Visual Environments, applied a deep learning model to identify tooth and sinus structures in dental panoramic radiographs, illustrating the early but growing interest in AI-assisted dental diagnostics locally. This study seeks to build on that emerging interest by applying AI-assisted screening directly within a clinic setting through Toothwise.

This project was created in partnership with CodiSmile Dental Clinic, which serves as the deployment site and clinical reference point throughout the study. The goal is not to take the place of the dentist but to give them a more efficient and objective tool to work with. With Toothwise integrated into the clinic's workflow, screenings can be documented more consistently, and patients can walk away with a clearer understanding of their oral health through the system's risk classification and personalized prevention plan features.

1.2 Statement of the Problem

Dental screening at CodiSmile Dental Clinic relies on visual inspection. Early-stage tartar, plaque, cavity, and gum inflammation are small, look similar to one another, and vary with lighting, angle, and stage, so visual assessment is subjective and findings can differ between examinations. The clinic has no standardized visual record of findings

to compare across visits, and patients are rarely given finding-specific preventive guidance. Undetected conditions can progress to pain, tooth loss, and costlier treatment, which makes consistent early screening a challenge of the method itself.

This study seeks to answer the following:

1. What limitations of visual-only screening make early-stage tartar, plaque, cavity, and gum inflammation and healthy tooth, as well as their relative severity, difficult to detect and distinguish consistently at CodiSmile Dental Clinic?
2. What gaps exist in how screening findings are recorded and made accessible to dentists and patients, given the absence of a standardized visual record and a way for patients to review their own history?
3. What barriers prevent patients from receiving timely, finding-specific preventive guidance after screening?
4. What is the level of acceptability of the Toothwise system in terms of the six quality characteristics of ISO/IEC 25010 (security, privacy, usability, reliability, performance efficiency, and data protection), as evaluated by IT experts and clinical end users?
5. What is the difference on the use of Toothwise to improve screening consistency and reduce screening time compared with the current manual process?

1.3 Objectives of the Study

1.3.1 General Objective:

To design, develop, and evaluate Toothwise, a dental screening support system that uses YOLO-based detection on still intraoral photographs, a KNN-based risk classification and recommendation feature, to provide consistent and accessible screening support for licensed dentists and patients at CodiSmile Dental Clinic.

1.3.2 Specific Objectives:

1. To train YOLOv12 and YOLOv8 on annotated still intraoral images to detect tartar, plaque, cavity, gum inflammation, and healthy tooth and compute an associated risk score and risk classification (low, medium, high) from the detected findings, and compare the two models by precision, recall, mAP, and inference speed.

2. To develop a results interface showing bounding boxes, condition labels, plain-language descriptions, and risk classification, with a per-patient record of findings across visits accessible to dentists and, through a self-service patient portal, to patients themselves.
3. To develop a KNN-based feature that generates a personalized prevention plan from the detected findings and computed risk classification.
4. To evaluate Toothwise using the six ISO/IEC 25010 characteristics with IT experts and clinical end users.
5. To measure changes in screening consistency and time with Toothwise versus the manual process, using the dentist's findings as the reference.

1.4 Scope, Delimitation and Limitation of the Study

Scope

This study covers the design, development, training, and evaluation of Toothwise, an AI-assisted dental screening support system for CodiSmile Dental Clinic. Using still images captured through a standard device camera, such as a smartphone or webcam, and uploaded to the system, Toothwise analyzes the patient's teeth and gums to detect five classes: tartar, plaque, cavity, gum inflammation, and healthy tooth. YOLOv12 and YOLOv8 models are trained on publicly available annotated dental image datasets and compared in terms of precision, recall, mean average precision (mAP), and inference speed; both models are also used together in the deployed system, with their outputs fused into a single clinical detection set. After training, the system is tested at CodiSmile Dental Clinic to assess how it performs in an actual clinical environment. Detection results are displayed through a results interface that shows bounding boxes, condition labels, plain-language descriptions, and a computed risk classification for licensed dentists, and findings are saved in a per-patient record that can be reviewed across visits. Based on the detected findings, a recommendation feature that uses the K-Nearest Neighbors (KNN) algorithm classifies the screening into a risk tier and generates a personalized prevention plan for each patient. The system also generates a downloadable PDF screening report for clinical

recordkeeping and maintains an audit log of sensitive staff actions for administrative review. The intended users are licensed dentists and clinic personnel at CodiSmile Dental Clinic; in addition, patients may register for a self-service portal to view their own screening history and past findings. The system is evaluated by IT experts and clinical end users using the six quality characteristics of ISO/IEC 25010, and its effect on screening consistency and screening time is compared with the current manual process, using the dentist's findings as the reference.

Delimitations

Toothwise is a support tool and not a replacement for a licensed dentist. The dentist at CodiSmile Dental Clinic always has the final say on any diagnosis or treatment decision. The prevention plan is meant to guide the patient and is not a formal medical prescription, so it does not include medication, treatment procedures, or dosages. The system works only with still photographic images captured through a standard device camera or uploaded from an existing file, and does not process X-rays or any other radiographic input, nor does it require specialized intraoral camera hardware. It is limited to the five classes named above and does not detect oral cancer, bone loss, impacted teeth, orthodontic issues, or other conditions that surface-level visual detection cannot reliably identify. The comparison of detection models is limited to YOLOv12 and YOLOv8 and does not include other detection architectures. The system is developed and tested only at CodiSmile Dental Clinic and does not integrate with billing or appointment scheduling systems. The evaluation covers software performance, quality, and acceptability, and does not measure long-term clinical outcomes such as changes in caries rates over time.

Limitations

The accuracy of the detection models depends on image quality, lighting, camera positioning, and the patient's cooperation during capture, and may vary more than it would with specialized intraoral camera hardware since a standard device camera is used instead. Because the models are trained on public datasets, performance in the clinic may differ

from training results if the clinic's captured images look different from the dataset images. The quality of the KNN-generated risk classification and recommendations depends on the validated cases used as its reference data. Findings may not generalize to other clinics, camera devices, or patient populations.

1.5 Significance of the Study

This study is expected to benefit the following groups:

Patients. Patients of CodiSmile Dental Clinic may benefit from a more consistent screening experience, in which findings such as tartar, plaque, cavities, gum inflammation, and healthy tooth are shown visually, explained in plain language, and accompanied by a risk classification. The personalized prevention plan gives them specific guidance based on their own findings, which can support earlier preventive action and better awareness of their oral health, and the patient portal allows them to review their own screening history at any time.

Dentists and CodiSmile Dental Clinic. Toothwise serves as a screening support tool that gives licensed dentists a consistent, image-based reference during consultations, while diagnosis and treatment decisions remain with them. The per-patient record of images, findings, and risk classifications gives the clinic a standardized documentation of screenings that can be reviewed and compared across visits, which the clinic currently lacks, and the PDF reporting and audit logging features support clinical recordkeeping and administrative oversight.

IT Practitioners and Software Developers. The comparison of YOLOv12 and YOLOv8, the use of KNN for risk classification and recommendations, and the ISO/IEC 25010 evaluation provide a documented example of building and assessing an AI-based health support system in a real clinic setting.

Future Researchers. Students and researchers interested in AI in dentistry can use the methodology, dataset preparation, model comparison, and system design of this study as a basis for further work, such as expanding the detectable conditions, testing other clinics, or studying long-term clinical outcomes.

The Research Community. There are relatively few local studies that apply AI-based detection to dental screening in the Philippine setting. This project contributes to that growing body of work and may serve as a reference for similar initiatives at other clinics across the country.

1.6 Definition of Terms

For clarity and common understanding, the following terms are defined based on how they are used in this study.

Calculus (Tartar). A hardened deposit of mineralized plaque that forms on the tooth surface and along the gumline. In this study, it refers to one of the five conditions detected by the Toothwise system.

Dental Caries (Cavity). The decay and structural breakdown of tooth enamel and underlying tissue caused by bacterial activity. In this study, it refers to one of the five conditions detected by the Toothwise system.

Gum Inflammation (Gingival Inflammation). A visible reddening, swelling, or irritation of the gum tissue, often an early indicator of gingivitis. In this study, it refers to one of the five conditions detected by the Toothwise system.

Inference Speed. The amount of time a trained detection model takes to process an input image and produce a detection result. In this study, it is one of the metrics used to compare the performance of YOLOv12 and YOLOv8.

K-Nearest Neighbors (KNN). A supervised machine learning algorithm introduced by Cover and Hart (1967), which classifies a new case by identifying and comparing it to the most similar cases in a reference dataset. In this study, KNN is used to classify each screening into a risk tier and generate a personalized prevention plan based on a patient's detected findings.

Mean Average Precision (mAP). A standard evaluation metric for object detection models that summarizes a model's precision and recall performance across all detected classes. In this study, it is one of the metrics used to compare YOLOv12 and YOLOv8.

Plaque. A soft, sticky film of bacteria that forms on the surface of teeth. In this study, it refers to one of the five conditions detected by the Toothwise system.

Precision. A performance metric that measures the proportion of correctly identified detections out of all detections made by the model. In this study, it is used to evaluate and compare the performance of YOLOv12 and YOLOv8.

Recall. A performance metric that measures the proportion of actual conditions in an image that were correctly identified by the model. In this study, it is used to evaluate and compare the performance of YOLOv12 and YOLOv8.

Risk Score. A numerical value from 0 to 100, classified as low, medium, or high, generated by the system from a patient's detected findings to indicate the overall severity of a screening result. In this study, the risk score is produced using the KNN algorithm.

Toothwise. The AI-assisted dental screening support system developed in this study, which uses YOLOv12 and YOLOv8 for detection and KNN for risk classification and generating personalized prevention plans.

YOLOv8 (You Only Look Once, version 8). An object detection model architecture in the YOLO family, used in this study as a baseline for comparison against YOLOv12 in detecting dental conditions.

YOLOv12 (You Only Look Once, version 12). An attention-centric object detection model architecture introduced by Tian et al. (2025), used in this study as the primary detection model for identifying dental conditions from intraoral images.

CHAPTER II

REVIEW OF RELATED LITERATURE AND STUDIES

This chapter presents the related literature and studies that have been gathered by the proponents. It introduces many ideas linked to the topic, which includes the research's principal focus. The study's synthesis is offered to analyze the previous research that the researchers have gathered.

2.1 Related Literature

Convolutional Neural Networks for Dental Image Diagnostics. Schwendicke et al. (2021) put together a systematic review of how convolutional neural networks had been applied in dental imaging, covering everything from caries detection to tooth segmentation. What stood out from their findings was that CNN models were not just competitive with experienced dental clinicians but actually outperformed them in some cases. The authors were careful to point out, though, that the quality of training data matters enormously. Models trained on poorly labeled or inconsistent datasets tend to underperform in real clinical scenarios. They also made it clear that these tools are meant to assist dentists, not replace them. That design principle carries directly into Toothwise, where all detected findings at CodiSmile Dental Clinic are reviewed by a licensed dentist before being communicated to any patient. The review also noted the absence of systems that go beyond detection and actually provide clinical guidance afterward, which is precisely the gap that the personalized prevention plan feature of Toothwise was built to address.

Deep Learning for Caries Detection: A Systematic Review. Mohammad Rahimi et al. (2022) reviewed a large number of studies that used machine learning and deep learning to detect caries across different image types, including intraoral photographs. Their main takeaway was that deep learning, particularly convolutional neural networks,

worked reliably for identifying both early and more advanced carious lesions. They also raised a concern that is very relevant to this project: there are not many large, standardized, and publicly available dental datasets out there, which makes it harder to train models that perform well across different clinical settings. For Toothwise, this limitation was addressed by combining publicly available annotated datasets with data augmentation techniques during model training. The review also found that no existing study had combined automated caries detection with a recommendation output, which reinforces why the prevention plan feature of Toothwise fills a genuine gap in the existing body of work.

Deep Learning for Detecting Dental Plaque and Gingivitis from Oral Photographs. Moharrami et al. (2025) conducted a systematic review evaluating deep learning models specifically for detecting dental plaque and gingivitis from RGB intraoral photographs, including images captured through standard consumer devices rather than specialized dental equipment. Across the studies they reviewed, models showed stronger and more consistent performance for plaque detection than for gingivitis, where the certainty of evidence was rated lower due to the more subtle and gradual visual presentation of early gingival inflammation. This finding is directly relevant to Toothwise, since gum inflammation is one of the five conditions the system is trained to detect from standard-camera photographs; it signals that this class may require more careful dataset curation and potentially larger sample sizes than the other conditions to reach comparable performance during training.

Oral Screening of Dental Calculus, Gingivitis, and Dental Caries through Segmentation on Intraoral Photographic Images. Zhang (2024) developed a deep learning segmentation system, Oral-Mamba, capable of simultaneously detecting three conditions — dental caries, dental calculus, and gingivitis — from a single set of intraoral photographic images, achieving accuracy rates above 0.81 for each condition. This study is one of the few in the literature that approaches multi-condition detection from a single intraoral image source rather than treating each condition as a separate, isolated detection task. It directly supports the feasibility of the multi-class detection approach used in

Toothwise, which similarly trains a single model to identify several distinct conditions — tartar, plaque, cavity, and gum inflammation — from one still-image source.

Oral Health Philippines Country Profile. The World Health Organization (2022) released a country-specific oral health profile for the Philippines as part of its Global Oral Health Status Report. The numbers are difficult to ignore. Dental caries was listed as one of the most prevalent health concerns in the country, with high rates of untreated cases across all age groups. The report also pointed out something that many Filipinos already know from experience: access to dental care is not evenly distributed. People in rural and community settings have significantly fewer options, and the limited dental workforce makes it harder to reach those who need care the most. This profile is the most authoritative reference for understanding the national oral health landscape that Toothwise seeks to contribute to. Deploying an artificial intelligence-assisted screening system at CodiSmile Dental Clinic is a concrete step toward addressing the access and consistency gaps that the WHO profile describes.

Pediatrician Knowledge, Attitude and Practices on Children's Oral Health in the Philippines. Oliveros Villarico et al. (2023) looked at something that often goes undiscussed, which is the role of pediatricians in preventing dental caries among Filipino children. Their study, conducted at a tertiary public hospital in Metro Manila, found that a significant number of pediatricians had limited awareness of evidence-based caries prevention approaches, even though the 2018 DOH National Survey on Oral Health already documented that 85.2% of Filipino five-year-olds had dental caries. The gap between what the data shows and what frontline healthcare workers actually know about it is striking. For this project, this literature helps explain why a tool like Toothwise matters at the clinic level. If even pediatricians are not fully equipped to address oral caries early, having an AI-assisted screening system that flags dental conditions during a regular consultation at CodiSmile Dental Clinic can make a real difference in catching problems before they get worse.

Guiding Responsible Artificial Intelligence in Healthcare in the Philippines. Alibudbud et al. (2025) wrote a narrative review and policy analysis on responsible AI

adoption in Philippine healthcare. They went through the current landscape of AI applications in the country, examined what regulatory frameworks exist, and laid out recommendations for making sure AI tools in clinical settings are used ethically and safely. Some of the key concerns they raised were limited digital infrastructure, the lack of local clinical validation standards, and the need for healthcare workers to understand and critically evaluate AI outputs. They also emphasized that AI systems used in Philippine healthcare should be transparent, subject to human oversight, and compliant with the Data Privacy Act of 2012. Reading this study shaped several decisions in how Toothwise was designed, particularly the requirement that all AI detection results at CodiSmile Dental Clinic must be reviewed and confirmed by a licensed dentist before they are communicated to patients, and that all patient data is stored securely with authenticated access and audit logging.

An Overview of Recommendation Techniques and Their Applications in Healthcare. Yue et al. (2021) reviewed the major recommendation approaches used in healthcare contexts, namely content-based, collaborative filtering, and hybrid methods, and mapped their application across five scenarios: dietary recommendation, lifestyle recommendation, training recommendation, decision-making support for patients and physicians, and disease-related prediction. Their review positions K-Nearest Neighbors-based approaches as a foundational technique underlying several of these healthcare recommendation scenarios, particularly where recommendations are generated by classifying a case according to its similarity to previously seen cases. This framework directly informs how the KNN-based risk classification and prevention plan feature of Toothwise was conceptualized, positioning it within the "decision-making support for patients" category identified in their review, an application area that, per their analysis, remains comparatively underexplored relative to dietary and lifestyle recommendation systems.

2.2 Related Studies

Detecting White Spot Lesions on Dental Photography Using Deep Learning. Askar et al. (2021) worked on detecting white spot lesions in intraoral dental photos captured with standard consumer cameras, using deep learning. White spot lesions are early stage demineralization, essentially the very beginning of a cavity before it becomes a visible hole. What made this study particularly useful is that it was done with a small dataset, which is a realistic constraint for many research projects, including this one. Despite that limitation, the model still managed to detect subtle surface level changes that would be easy to miss in a routine visual check. For Toothwise, this matters because it shows that even with limited annotated data and images from standard, non-specialized cameras, deep learning can pick up early, crucial changes from photographs, which is the same image modality the system uses at CodiSmile Dental Clinic. It also gave the team confidence that training a YOLOv12 model for early cavity detection is achievable even when the training dataset is not massive and specialized imaging hardware is not required.

Ultralytics YOLO (Version 8.0.0). Jocher et al. (2023) released YOLOv8 as an open source framework under the Ultralytics platform. Compared to older YOLO versions, YOLOv8 came with a cleaner training pipeline, an anchor free detection head, and support for multiple tasks like detection, segmentation, and classification all in one package. It became one of the most widely used detection frameworks in both academic and applied settings because it was relatively easy to work with while still delivering strong results. Understanding YOLOv8 matters for this project because YOLOv12 was built on the same foundation. The training tools, dataset formats, and evaluation methods used in YOLOv8 largely carry over to YOLOv12, so much of what Jocher et al. (2023) established applies directly to how the Toothwise models were trained and tested. YOLOv8 also serves as the direct comparison baseline against YOLOv12 in this study's model evaluation.

Accuracy in Diagnosing Caries in Young Permanent Molars Using Interproximal Radiographic Imaging and Validation by Artificial Intelligence. de Brito et al. (2025) applied YOLOv8 together with a pre-trained EfficientNet-B0 model to diagnose caries in interproximal radiographic images of young permanent molars, comparing the AI-based evaluation against diagnoses made by dental professionals. Their results showed strong inter-examiner agreement and confirmed that YOLOv8 performs reliably even on the small, low-contrast lesion types that are typically difficult to detect visually. This study provides direct evidence that YOLOv8 is a credible, well-validated baseline for detecting caries-type conditions, supporting its selection as the comparison model against YOLOv12 in this study.

Deep Learning in Oral Hygiene: Automated Dental Plaque Detection via YOLO Frameworks and Quantification Using the O'Leary Index. Ramírez-Pedraza et al. (2025) trained and compared several YOLO model variants, including YOLOv11, specifically for detecting dental plaque in varied imaging conditions, using the O'Leary Index as a clinical benchmark for plaque quantification. Their results showed that YOLO-based models could reliably distinguish plaque at different stages of maturity, though newer, less mineralized plaque was harder to detect than older, more visually distinct plaque, since it tends to blend with surrounding gingival tissue. This finding is directly relevant to Toothwise, since it identifies a specific detection challenge, distinguishing early plaque from gum tissue, that the model must account for during training and dataset annotation for the plaque class.

Diagnostic Accuracy of Dental Caries Detection Using Ensemble Deep Learning with Intraoral Camera Images. Kang et al. (2024) took a different approach from most dental AI studies by using an ensemble of deep learning models rather than a single model for caries detection. By combining multiple detectors, they were able to reduce false positives while keeping sensitivity high, which is especially important in a clinical setting where unnecessary alarms can make dentists lose trust in the system. Their work confirmed that photographic intraoral images, rather than radiographs, are a sufficient and practical input for reliable AI-based detection, validated against annotations made by

experienced dentists. This finding supports the still-image, standard-camera approach used in Toothwise, and gave the team confidence that deep learning for photographic dental screening is not just theoretically possible but has already been shown to work efficiently in a clinical environment similar to CodiSmile Dental Clinic.

AI-Assisted Detection of Interproximal, Occlusal, and Secondary Caries on Bite-Wing Radiographs Using YOLOv8. Karakus et al. (2024) applied YOLOv8 to the task of detecting different types of dental caries from bitewing radiographs. One thing that stood out in their findings was how well YOLOv8 handled small and overlapping lesions, which is a tricky problem in dental imaging because caries can be close together or partially hidden. Even though Toothwise uses standard photographic images rather than radiographs, the core detection challenge is similar. The model needs to find and correctly label multiple small, carious areas within a single image. The performance of YOLOv8 in this study, combined with the architectural improvements introduced in YOLOv12, provided the team with a strong technical basis for believing that YOLOv12 can handle the multi-condition detection demands of Toothwise.

YOLOv12: Attention-Centric Real-Time Object Detectors. Tian et al. (2025) introduced YOLOv12 in a preprint that gained significant attention in the computer vision community. The paper presents a new approach to the YOLO architecture centered around area attention mechanisms, which allow the model to focus on relevant regions of an image rather than processing everything equally. The results were impressive. The YOLOv12 N configuration achieved 40.6% mean average precision at an inference latency of just 1.64 milliseconds on a T4 GPU, outperforming comparable YOLOv10 and YOLOv11 models. For a system like Toothwise, where detection results need to be produced quickly enough to fit within a standard dental consultation, that combination of accuracy and speed is exactly what is needed. This paper is the primary technical reference for the detection model used in this project, and its findings directly justify choosing YOLOv12 for evaluation against YOLOv8 for deployment at CodiSmile Dental Clinic. At the time of this review, no widely published study had applied YOLOv12 specifically to multi-condition

dental screening from standard photographic images, positioning Toothwise among the earliest efforts to explore this architecture in that context.

Feedback Collection and Nearest-Neighbor Profiling for Recommendation Systems in Healthcare Scenarios. António et al. (2024) developed a physical activity recommendation system for patients in rehabilitation using a K-Nearest Neighbors sampling strategy combined with feedback collection, generating activity recommendations by identifying patients with the highest similarity to a given case based on their health profile. Their system was shown to outperform traditional collaborative filtering and purely human-constrained recommendation methods in producing recommendations suited to a patient's specific therapeutic needs, physical capacity, and stage of recovery. This is one of the few studies to apply a KNN-based approach directly to generating personalized, condition-specific health guidance rather than general product or content recommendations, and it provides a close technical parallel for how Toothwise's KNN-based feature classifies a screening into a risk tier and composes a prevention plan by matching a patient's detected findings against similar, dentist-validated reference cases.

Bridging the Gap or Widening the Divide: A Call for Capacity-Building in Artificial Intelligence for Healthcare in the Philippines. Gutierrez and Viacrusis (2023) took a critical look at AI readiness in Philippine healthcare. When the researchers tested AI generated responses to clinical queries relevant to Filipino patients, they found notable inaccuracies, including instances where the AI produced fabricated references and incorrect medical guidance. The paper used this to argue that Filipino healthcare professionals urgently need better AI literacy and that unsupervised use of AI clinical tools carries real risks. The call to action was for structured local validation and, most importantly, mandatory human oversight in any clinical AI application. This study speaks directly to a design decision in Toothwise. The system does not allow AI detection results to be shown to patients without going through a dentist first. At CodiSmile Dental Clinic, the dentist reviews every output before it becomes part of the clinical record or is communicated to the patient, which is exactly the kind of human oversight that Gutierrez and Viacrusis argued for.

2.3 Synthesis

The literature and studies reviewed in this chapter come together to build a clear and well-supported case for the approach that Toothwise has taken. Schwendicke et al. (2021) and Mohammad Rahimi et al. (2022) established through separate systematic reviews that deep learning models are genuinely reliable for dental caries detection, often performing at the level of trained clinicians or better. Moharrami et al. (2025) and Zhang (2024) extended this evidence beyond caries alone, showing that plaque, calculus, and gingivitis can also be reliably detected from standard intraoral photographs, with Zhang's work in particular demonstrating that multiple conditions can be detected from a single image source using one model, directly supporting the five-class detection approach used in Toothwise. Kang et al. (2024) and Askar et al. (2021) further confirmed that photographic images captured without specialized intraoral camera hardware are a sufficient and practical input modality for AI-based detection in an actual clinical setting, closely paralleling the still-image, standard-camera acquisition approach of Toothwise at CodiSmile Dental Clinic.

The most important gap identified across the reviewed works was stated most directly by Moharrami et al. (2024) in prior work on caries detection from photographs: existing dental AI systems detect conditions but stop there. None of the detection-focused studies reviewed generated recommendations or patient-facing guidance after detection. At the same time, Yue et al. (2021) and António et al. (2024) show that KNN-based recommendation approaches have been successfully applied in other healthcare domains, such as physical rehabilitation, but remain comparatively unexplored in dental screening specifically. Toothwise fills this gap by embedding a KNN-based risk classification and recommendation engine into a multi-condition detection pipeline, pairing YOLO-based

detection with an actionable, personalized prevention plan, a combination not found together in the literature reviewed.

On the technical side, Tian et al. (2025) established YOLOv12 as a highly accurate and efficient real-time object detection architecture, while Jocher et al. (2023), de Brito et al. (2025), Ramírez-Pedraza et al. (2025), and Karakus et al. (2024) documented the strong and well-validated performance of YOLOv8 specifically across caries and plaque detection tasks, supporting its selection as this study's comparison baseline. Askar et al. (2021) further showed that deep learning-based detection is achievable even under limited dataset conditions, supporting the training strategy used in this project.

The Philippine context literature from the WHO (2022) and Oliveros Villarico et al. (2023) grounds all of this in a concrete public health reality. Dental conditions are widespread among Filipinos, they often go undetected until they have already progressed, and access to professional dental care remains out of reach for many due to cost and geography. Alibudbud et al. (2025) and Gutierrez and Viacrusis (2023) shaped responsible design decisions around human oversight, data privacy, and clinical transparency. What no existing work reviewed has done is bring together YOLO-based multi-condition detection from standard photographic images, a KNN-based risk classification and personalized prevention plan generator, and a full-stack clinical system deployed at an actual Philippine dental clinic. That is the specific contribution of Toothwise.

CHAPTER III

METHODOLOGY

This chapter presents the methodology used in the development of Toothwise: A YOLO-Based Multi-Condition Dental Screening System with KNN-Based Recommendation for CodiSmile Dental Clinic. It describes the development methodology, the composition and responsibilities of the five-member development team, the project planning and requirements analysis process, the major system design diagrams required for the capstone examination, and the development, testing, and screenshot documentation of the current system.

3.1 Development Methodology

Toothwise was developed using the Agile Software Development Life Cycle (Agile SDLC). Agile was selected because the system integrates several interconnected components, including a React.js web interface, a FastAPI REST backend, a PostgreSQL database, a dual-model YOLOv8/YOLOv12 detection pipeline, a KNN-based risk classification and recommendation engine, authentication and authorization, audit logging, and a separate patient self-service portal. These components benefit from incremental development, continuous testing, regular review, and repeated refinement rather than a single sequential build.

Instead of completing the entire system in one long development cycle, the project was carried out through manageable iterations in which requirements were prioritized, features were implemented, tested, reviewed, and improved before the next set of features was developed. Each iteration followed the same general cycle of planning, requirements analysis, system design, development, testing, evaluation, and refinement, illustrated in Figure 1.

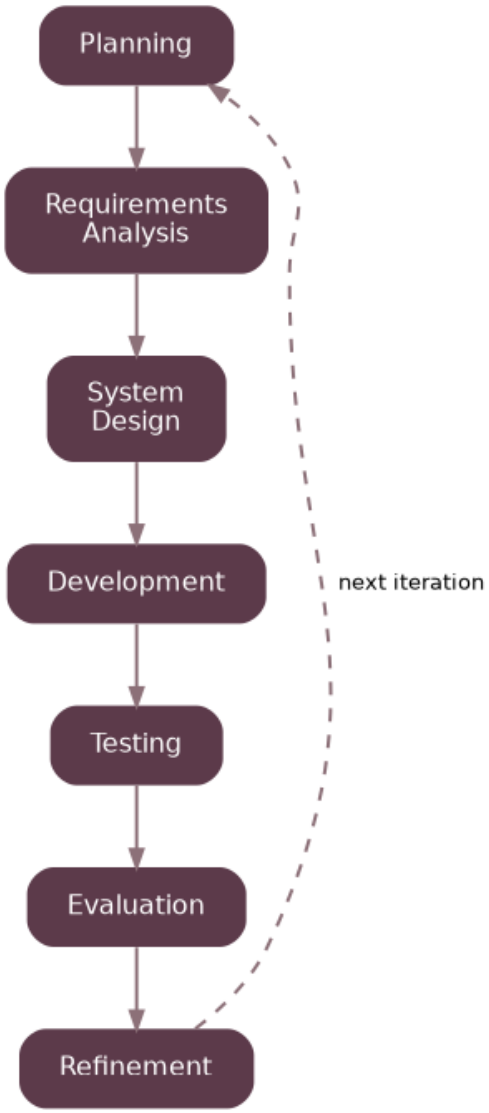


Figure 1 Agile Development Methodology

The cycle in Figure 1 was applied repeatedly across the project's documented iterations. Planning and requirements analysis identified the features required for the clinic workflow; system design translated those requirements into frontend routes, API endpoints, and database structures; development implemented the corresponding React, FastAPI, and machine-learning components; testing verified the implemented behavior;

evaluation compared the result against the intended requirements; and refinement corrected any issues found before the next iteration began.

3.2 Project Planning

The Toothwise project was planned around the screening and record-keeping workflow of CodiSmile Dental Clinic. The following considerations guided the planning of the system:

- **Project Objectives** — to provide CodiSmile Dental Clinic with an AI-assisted dental screening system that detects tartar, plaque, cavity, gum inflammation, and healthy tooth from a still photograph, estimates a risk level, and generates a recommendation/prevention plan, while also supporting day-to-day clinic record-keeping.
- **Project Scope** — the system covers staff authentication, patient management, dental screening and AI detection from still images, risk assessment, recommendation generation, clinical notes, PDF reporting, audit logging, and a self-service patient portal.
- **Target Users** — dentists and administrators at CodiSmile Dental Clinic, and the clinic's patients through the patient portal.
- **Clinic Workflow** — a patient is registered or already on file, a still intraoral photograph is captured through a standard device camera or uploaded during a screening, the system produces detection, risk, and recommendation results for the dentist to review, and the dentist may add clinical notes or generate a PDF report.
- **Functional Requirements** — detailed in Section 3.4.
- **Non-Functional Requirements** — detailed in Section 3.4, covering security, privacy, usability, reliability, performance efficiency, and data protection.

- **AI Screening Requirements** — the system was planned to use two independent detection models (YOLOv8 and YOLOv12) whose outputs are fused into a single clinical detection set, rather than relying on a single model.
- **Database Requirements** — a relational PostgreSQL schema was planned to store dentists, patients, screenings, detections, recommendations, clinical notes, and audit logs with referential integrity between related records.
- **Security Requirements** — JWT-based authentication, hashed passwords, rate-limited login endpoints, authenticated (non-public) image access, and audit logging were planned from the outset given the clinical and personal nature of the data involved.

The requirements and technical decisions behind this planning were drawn directly from the current Toothwise project package rather than from separate interview transcripts or survey forms, since the project does not maintain those as separate dated records. Table 1 summarizes the actual sources used.

Table 1 Sources of Requirements and Technical Data

Source	Use in the Study
System and code review	The frontend routes, pages, API client, FastAPI routers, schemas, security module, ML modules, and database models were reviewed to identify actual implemented functions and data flows.
Database and schema review	The PostgreSQL schema and migration scripts were reviewed to verify entities, relationships, enumerated values, detection sets, and audit logs.
Testing and debugging records	CHANGES.md documents API contract reconciliation, enum and image-path fixes, smoke testing, security validation, and feature verification.

AI evaluation procedure	The evaluate_models.py utility defines how precision, recall, F1, mAP50, mAP50-95, and average inference time can be obtained from a labeled test set. No real model comparison values are claimed here since the current package documents the use of mock inference until trained weights are installed.

3.3 Roles and Responsibilities of the Five-Member Development Team

Toothwise was developed by a five-member capstone development team. Table 2 presents each member's assigned role and corresponding responsibilities based on the work represented in the current project package. A member may carry more than one responsibility.

Table 2 Roles and Responsibilities of the Development Team

Group Member	Assigned Role	Responsibilities
Arreola, Reizen S.	Project Manager / System Analyst / Programmer	Coordinated the activities of the development team, monitored the project schedule, translated clinic workflow requirements into system requirements, and ensured that project deliverables remained aligned with the objectives of the Toothwise system; Developed and integrated the YOLOv8/YOLOv12 dental image detection pipeline, and the KNN recommendation engine
Magpantay, Rizelle Carla N.	Backend Developer	Implemented and integrated the FastAPI REST API, authentication and

		authorization logic, and database operations and connected the frontend with the backend API; resolved programming and integration issues identified during testing.
Sales, Keith Ashly V.	Frontend Developer	Implemented the React.js user interface, including the dashboard, patient and dentist management screens, screening workflow, patient portal, and the visual design and layout of the system.
Tubillas, Brayan Axel L.	Database Developer / Documentation Lead	Designed and maintained the PostgreSQL database schema and relationships; and conducted functional, security, and integration testing of the system; Responsible for organizing, compiling, and standardizing all documents into a cohesive capstone paper that complies with formatting, style, and academic guidelines.
Ponce, Robin C.	UI-UX Designer	Responsible for defining information architecture, and creating accessible design system for toothwise, to ensure the capstone deliverable delivers an intuitive, user-centered experience.

The specific division of responsibilities above reflects the major components of the current system — project coordination and requirements, backend/API development, frontend/interface development, database/documentation, and interface design work.

3.4 Requirements Analysis

The requirements of Toothwise were analyzed from the intended clinic workflow and the functions represented in the current system and its API. Toothwise has two authenticated user groups: clinic-side users (dentists and administrators) and patients using the self-service portal. Tables 3 and 4 present the functional and non-functional requirements that are currently implemented in the application.

Table 3 Functional Requirements

ID	Requirement	Specification
FR-01	Staff authentication	Authenticate dentists and administrators by email and password and issue a JWT bearer token for authenticated API access.
FR-02	Role-based access	Distinguish administrator, dentist, and patient access scopes and restrict protected functions accordingly.
FR-03	Patient management	Create, view, update, and delete patient records and view screening history and risk trends.
FR-04	Dental screening	Submit a still intraoral photograph, captured through a standard device camera or uploaded from an existing file, in a single multipart screening request and receive detection, risk, recommendation, and screening-record results.
FR-05	AI detection	Run the configured YOLOv8 and YOLOv12 detectors, retain raw model outputs, and produce a fused clinical detection set.

FR-06	Risk assessment	Calculate a 0-100 risk score and classify the screening as LOW, MEDIUM, or HIGH using the KNN algorithm.
FR-07	Recommendation generation	Use the KNN recommender and condition-specific rules to generate an ordered prevention plan.
FR-08	Screening review	View screening images, fused detections, risk information, recommendations, clinical notes, and model-comparison information.
FR-09	PDF report	Generate an authenticated PDF screening report containing patient/screening information, findings, prevention plan, and clinical notes.
FR-10	Patient portal	Patients register, log in, and view their own screening history and screening details.
FR-11	Audit logging	Record relevant authentication and create/update/delete actions for administrative review.
FR-12	Health monitoring	Expose a health endpoint that checks database connectivity without exposing internal errors.
FR-13	Database backup	Provide a pg_dump-based backup utility that creates timestamped, retained backups.

Table 4 Non-Functional Requirements

ID	Quality Area	Specification
NFR-01	Security	Passwords hashed with bcrypt; protected operations require JWT authentication; login/registration endpoints rate-limited to 5 requests per minute per IP.

NFR-02	Privacy	Screening images are not publicly exposed; authenticated and ownership-checked image endpoints are used instead of a public upload route.
NFR-03	Usability	Browser-based workflows for dashboard monitoring, patient records, screening, dentist management, audit review, and patient self-service.
NFR-04	Reliability	Backend returns a generic response for unhandled errors while recording diagnostic detail server-side.
NFR-05	Performance Efficiency	YOLOv8 and YOLOv12 inference times are recorded per screening; model-comparison views report runtime and detection statistics.
NFR-06	Data Protection	Relational constraints, enumerated status/risk values, foreign keys, and transactions preserve consistent, valid records.

3.5 System Design

This section presents the major design diagrams required for the capstone preliminary examination: the Use Case Diagram, the Entity Relationship Diagram, the Data Flow Diagrams at Level 0 and Level 1, and the development Gantt Chart. Each diagram was produced from the actual implemented system roles, database schema, and documented development sequence.

3.5.1 Use Case Diagram

Figure 2 presents the use case diagram of Toothwise. It represents the system's user roles — Administrator, Dentist, and Patient — and the functions each role can currently perform, verified against the implemented frontend routes and backend authorization checks. This diagram represents system user roles, not the development team presented in Section 3.3.

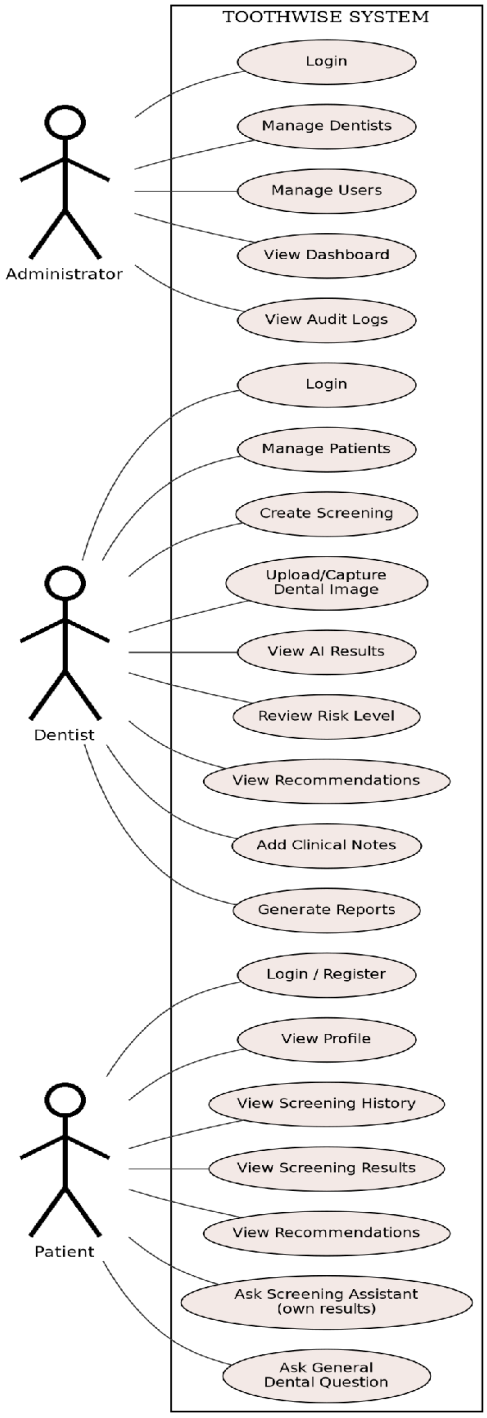


Figure 2 Use Case Diagram

3.5.2 Entity Relationship Diagram (ERD)

Figure 3 presents the Entity Relationship Diagram of the Toothwise PostgreSQL database, verified against the current database/schema.sql and backend/app/models.py. The database contains seven tables: dentists, patients, screenings, detections, recommendations, clinical_notes, and audit_logs, related through foreign keys as shown below. The complete field-level specification of each table is presented in Section 4.4.

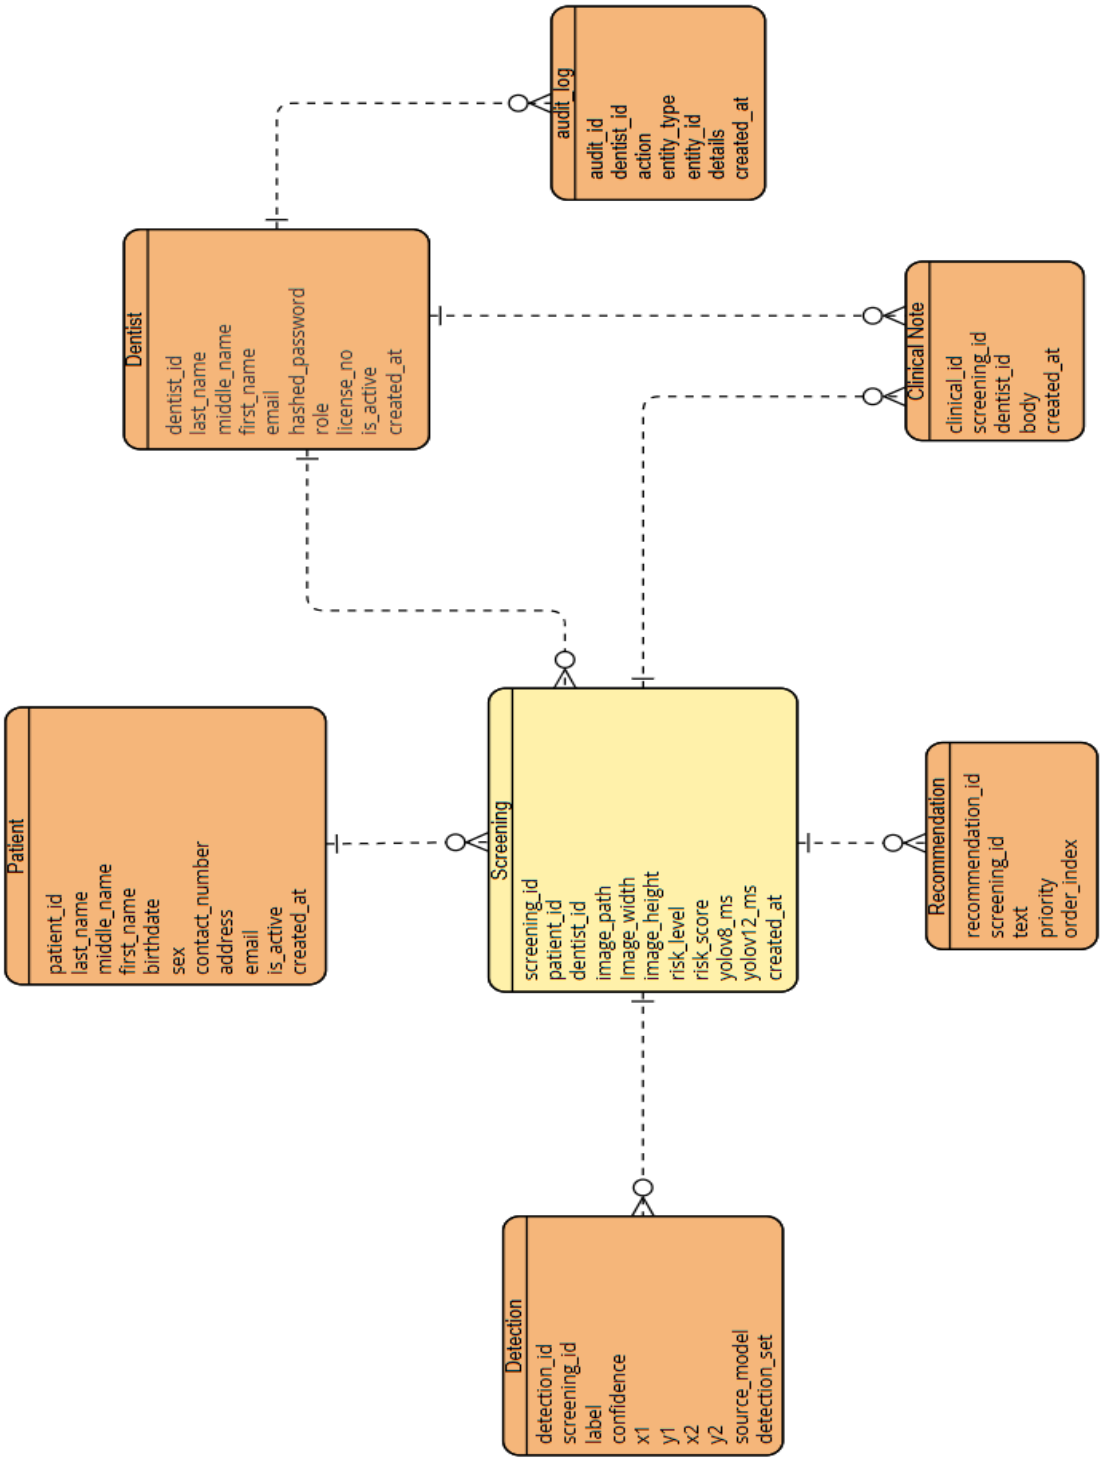


Figure 3 Entity Relationship Diagram (ERD)

Table 5 consolidates the seven core entities of the Toothwise database into a single database schema, summarizing each table's primary key, purpose, classification, and relationships to the other entities, ahead of the full field-level data dictionary presented in Section 4.4.

Table 5 Toothwise Database Schema Data Dictionary

Table	Primary Key	Role/ Classification	Purpose	Related Tables
dentists	id	Dimension Table	Stores staff accounts (dentist or admin role) used for staff login and role-based access control.	screenings, clinical_notes, audit_logs
patients	id	Dimension Table	Stores patient demographic information and, optionally, patient-portal login credentials.	screenings
screenings	id	Fact Table	Central record of one AI-assisted dental image screening event, including risk level and model timing.	patients, dentists, detections, recommendations, clinical_notes
detections	id	Fact Detail Table	Stores each bounding-box finding produced by the YOLOv8/YOLOv12 ensemble for a screening (5-label taxonomy).	screenings

recommendations	id	Fact Detail Table	Stores the risk classification and prevention/advice text generated by the KNN recommender for a screening.	screenings
clinical_notes	id	Fact Detail Table	Stores dentist-authored clinical notes attached to a screening.	screenings, dentists
audit_logs	id	System Table	Append-only record of sensitive staff actions, kept for administrative review.	dentists

3.5.3 Data Flow Diagram Level 0

Figure 4 presents the Level 0 (context) Data Flow Diagram of Toothwise. It shows the system as a single process exchanging data with its three external entities, Administrator, Dentist, and Patient, verified against the actual authenticated routes each role can reach.

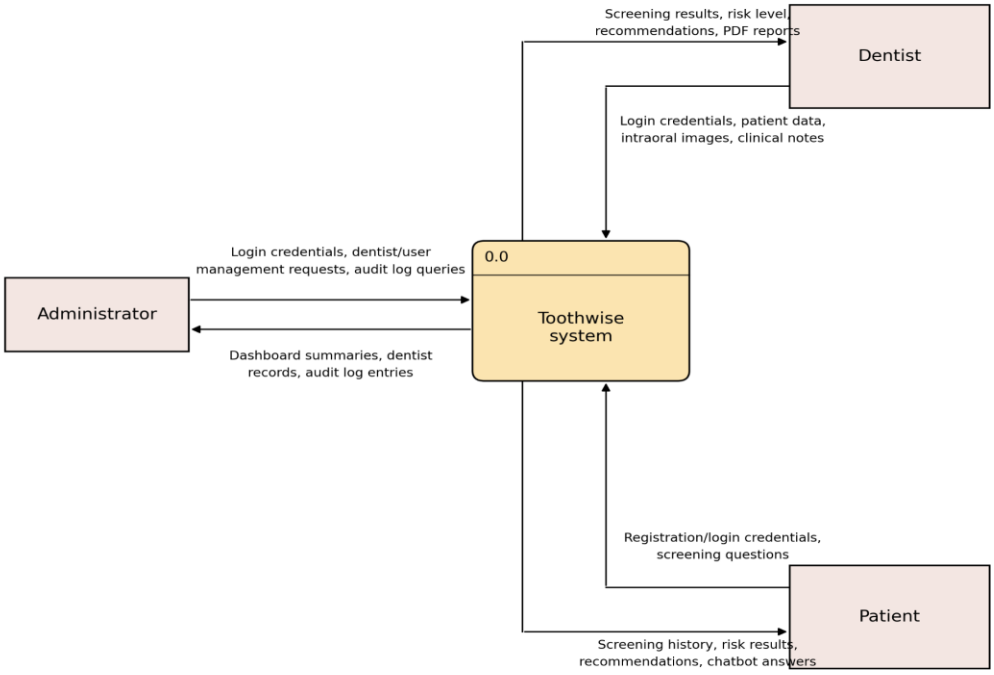


Figure 4 Data Flow Diagram Level 0

3.5.4 Data Flow Diagram Level 1

Figure 5 decomposes the Toothwise system into its nine major processes: Authentication and Authorization, Patient Management, Dental Screening, AI Detection, Risk Assessment, Recommendation Generation, Clinical Notes, Report Generation, and Audit Logging, together with their associated data stores. This diagram is consistent with the Level 0 diagram in Figure 4 and with the ERD in Figure 3.

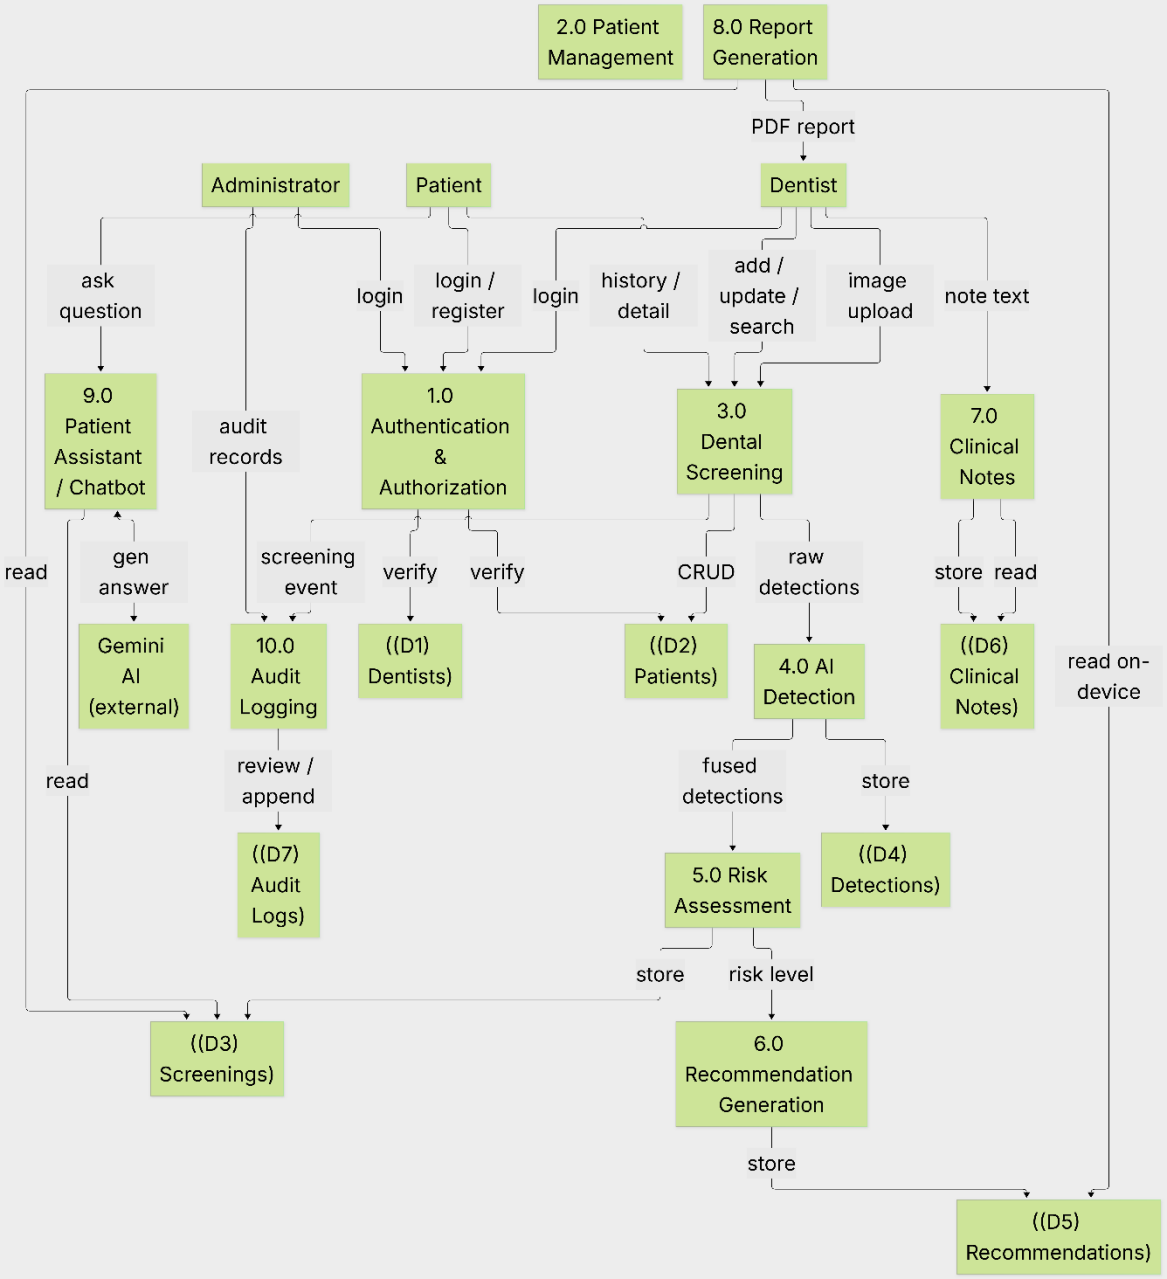


Figure 5 Data Flow Diagram Level 1

3.5.5 Gantt Chart

Figure 6 presents the Toothwise development schedule across five major phases: Planning, System Design, Development and AI Integration, Testing, and Documentation, spanning July 11 to October 30. Development and AI Integration is the longest and most substantial phase, since it covers the system's core functional components, from authentication and the relational database to the YOLOv8/YOLOv12 detection pipeline and the KNN risk classification and recommendation engine. Documentation ran in parallel with development, progressing through 30%, 60%, and 100% completion milestones alongside the corresponding system builds..

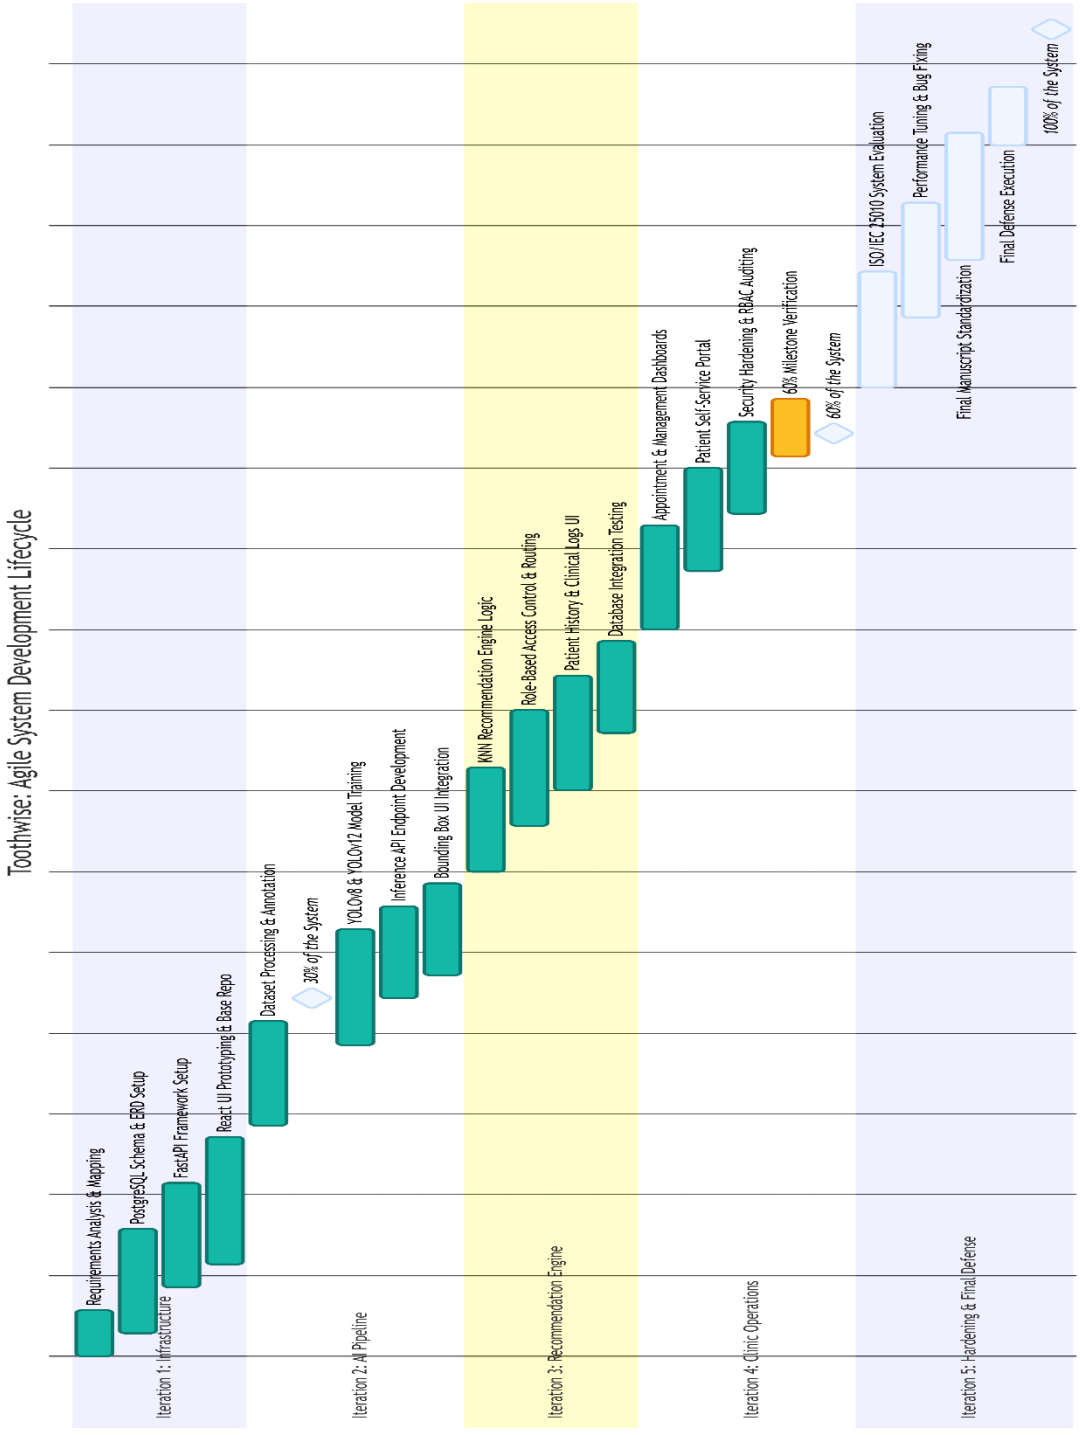


Figure 6 Toothwise Development Gantt Chart

3.6 Software Specifications

The software requirements of Toothwise cover the programming languages, frameworks, libraries, services, and development tools represented in the current project package, distinguishing development-environment resources (Table 6) from implementation/deployment-environment resources (Table 7).

Table 6. Software Development Requirements

Software	Specification	Description
Operating System	Windows 10 or 11, 64-bit	Primary development environment for the frontend, backend, database administration, and system testing.
Runtime / Backend Language	Python 3.11	Used for the FastAPI backend, AI inference, KNN-based recommendation engine, database integration, testing, and utility scripts.
Frontend Language	JavaScript (ES6+)	Used to implement the React.js user interface and client-side application logic.
Frontend Framework	React 18.2.0	Provides the component-based web interface for dentists, administrators, clinic staff, and the patient portal.
Routing	React Router DOM 6.21.1	Provides client-side routing for login, dashboard, patient records, screening, dentist management, audit logs, and the patient portal.

HTTP Client	Axios 1.6.5	Handles authenticated communication between the React frontend and the FastAPI REST API.
UI Feedback	React Hot Toast 2.4.1	Provides success and error notifications in the web interface.
Camera Integration	React Webcam 7.2.0	Supports still-image capture from a standard device camera (webcam) for uploading during a screening.
Charting	Recharts 2.10.3	Renders the patient risk trend and dashboard trend visualizations.
Backend Framework	FastAPI 0.111.0	Implements the REST API, request validation, authentication dependencies, routing, health checks, and integration with AI and PostgreSQL.
ORM / Database Driver	SQLAlchemy 2.0.30 and psycopg 3.1.19	Provide ORM-based database access and PostgreSQL connectivity.
Validation / Configuration	Pydantic 2.7.4 and pydantic-settings 2.3.4	Validate API payloads and load environment-based configuration.
AI Framework	Ultralytics YOLO with PyTorch	Runs the YOLOv8 and YOLOv12 object-detection models used in the dental screening pipeline.
Computer Vision	OpenCV 4.10.0.84, Pillow 10.3.0	Support image processing, validation, and preparation of uploaded still images.

Recommendation / ML	scikit-learn 1.5.0, joblib 1.4.2	Support the KNN risk classification and recommendation model and persistence of the trained recommender artifact.
Database	PostgreSQL 15	Stores patient, dentist, screening, detection, recommendation, clinical note, and audit-log records.
Authentication / Security	python-jose 3.3.0, bcrypt 4.2.0, SlowAPI 0.1.9	Provide JWT authentication, password hashing, and rate limiting.
PDF Reporting	ReportLab 4.2.2	Generates authenticated PDF screening reports.
Testing	pytest 8.2.2, HTTPX 0.27.0	Support automated and API-level testing of the FastAPI application.
Version Control	Git and GitHub	Used for source-code versioning, collaboration, and project traceability.
IDE	Visual Studio Code	Primary editor for frontend, backend, database, and AI development.
Package Managers	npm and pip	Manage JavaScript and Python project dependencies.

Table 7. Implementation Software Requirements

Software	Specification	Description
Server Operating System	Windows 10 or 11, 64-bit	Hosts the FastAPI backend, PostgreSQL database, and YOLO inference services in the

		documented local deployment setup.
Client Platform	Any operating system with a modern web browser	Dentists and clinic staff access the system without installing a separate desktop application.
Browser	Google Chrome 110+ or Mozilla Firefox 110+	Supported browser environment for the React.js interface and camera-enabled workflows.
Backend Runtime	Python 3.11 + FastAPI	Runs the REST API, authentication, screening workflow, and supporting services.
Database	PostgreSQL 15+	Stores persistent clinical and administrative records.
Authentication	JWT Bearer tokens	Separates staff and patient access scopes and protects authenticated API routes.
AI Runtime	YOLOv8 + YOLOv12 ensemble	Runs two detection models independently and fuses their outputs for the clinical detection set. The current package defaults to deterministic mock inference until trained weights are installed.
Recommendation Runtime	KNN, K=5	Classifies the screening into a risk tier used to compose the prevention plan.
Upload Limit	15 MB maximum configured upload size	Limits the size of uploaded screening images before processing.

API Documentation	FastAPI Swagger UI / ReDoc	Provides developer-facing documentation for the REST API.

3.7 Hardware Specifications

Table 8. Hardware Development Requirements

Hardware	Specification	Description
Processor	Intel Core i7 or AMD Ryzen 7, 8 cores, 2.5 GHz or above	Recommended for concurrent frontend, backend, image-processing, and model-development workloads.
RAM	16 GB minimum; 32 GB recommended	Supports model training, dataset preprocessing, database services, and development tools.
GPU	NVIDIA GPU with CUDA support, at least 6 GB VRAM	Recommended for practical YOLO training and accelerated inference during model development.
Storage	512 GB SSD or higher	Stores source code, datasets, model weights, Python environments, PostgreSQL data, and generated artifacts.
Display	1920 × 1080 or higher	Supports development, image review, detection

		overlays, and responsive interface testing.
Network	Stable Internet connection, 10 Mbps or higher	Supports dependency installation, dataset access, repository synchronization, and remote development resources.

Table 9. Implementation Hardware Requirements

Hardware	Specification	Description
Server Machine	Intel Core i5 / AMD Ryzen 5, 8 GB RAM, 256 GB SSD or better	Runs the FastAPI API, PostgreSQL database, and YOLO inference services.
Client Workstation	Computer or tablet with at least 4 GB RAM and a modern browser	Used by dentists and clinic staff to access the web application.
Device Camera	Smartphone camera or webcam capable of capturing still images	Captures the patient's still intraoral photograph used as input to the screening process; no specialized intraoral camera hardware is required.
Local Area Network	Wired or Wi-Fi router, 100 Mbps or higher	Connects the server and authorized client devices for local access.

Uninterruptible Power Supply	600 VA minimum recommended	Helps protect the server and preserve availability during short power interruptions.

3.8 System Development

Toothwise was developed iteratively, following the six development phases in Table 10, applied repeatedly across the five documented iterations in Table 11 rather than as a single sequential build.

Table 10 Development Phase Activities

Phase	Activities
1. Planning and Requirements	Define the clinic workflow, user roles, screening objectives, AI detection needs, risk assessment, recommendation output, patient portal functions, security requirements, and reporting needs.
2. System Analysis	Translate the requirements into frontend routes, FastAPI endpoints, database entities, AI modules, authentication dependencies, and data flows.
3. System Design	Design the three-layer architecture, relational schema, JWT authentication model, detection-set structure, risk classification and recommendation flow, patient ownership checks, and local network deployment.
4. Development	Implement React pages and API services, FastAPI routers and schemas, SQLAlchemy models, PostgreSQL schema, YOLO detector integration, KNN risk classifier and recommender, reports, audit logs, and patient portal.
5. Integration and Testing	Reconcile frontend requests with the actual backend contract, validate database enum labels and image paths, smoke-test API workflows, test security controls, and verify the frontend build.
6. Deployment and Maintenance	Prepare the PostgreSQL database, backend environment, frontend build, model-weight configuration, local network access, backups, and operational monitoring. Continue maintenance through versioned changes and database migration scripts.

Table 11 Agile Iteration Summary

Iteration / Focus	Major Activities	Output
Iteration 1 – Core System Foundation	Initial React.js frontend, FastAPI backend, PostgreSQL schema, authentication, and core patient/dentist structures.	Working application foundation and authenticated API.
Iteration 2 – AI Screening	Integration of still-image upload, YOLO detection workflow, detection storage, risk scoring, and screening results.	End-to-end dental screening workflow.
Iteration 3 – Recommendations and Clinical Workflow	KNN risk classification and recommendation logic, prevention plans, clinical notes, screening details, and PDF reporting.	Expanded clinical screening and recommendation workflow.
Iteration 4 – Clinic Operations	Dashboard functions, dentist management, audit logs, and patient portal features.	Broader clinic and patient workflows.
Iteration 5 – Testing and Hardening	API contract reconciliation, database and enum corrections, image access controls, rate limiting, validation, security headers, and frontend integration fixes.	Hardened and verified system behavior.

Each iteration in Table 10 was carried through the full cycle shown earlier in Figure 1: the corresponding requirements were planned and analyzed, the affected parts of the system were designed, the frontend and backend components were developed and

integrated, the iteration's features were tested, the result was evaluated against the intended requirement, and any issue found was corrected before the next iteration began. The Gantt chart in Figure 6 visualizes this same sequence.

The system was implemented as a three-tier web application: a React.js single-page frontend, a FastAPI REST backend exposing versioned JSON endpoints, and a PostgreSQL relational database. The AI detection pipeline (YOLOv8 and YOLOv12) and the KNN-based risk classification and recommendation engine were developed as backend modules invoked synchronously by the screening endpoint. The complete technical specification of this architecture is presented in Chapter IV.

3.9 System Testing

Toothwise was tested through functional, integration, and security testing conducted against the running backend using FastAPI's TestClient, supplemented by manual and code-review verification for checks that are not easily exercised through an automated client alone.

To confirm that Toothwise satisfies its intended functional requirements, a systematic software testing phase was conducted. This process evaluated module-level correctness, system-wide component interaction, and user acceptance, while specifically validating the accuracy of the YOLOv8/YOLOv12 detection pipeline and the KNN risk classification and recommendation engine in producing appropriate detections, risk classifications, and prevention plans based on patient screening data.

The system was evaluated using three structured testing tiers: Unit Testing (utilizing a white-box approach), Integration Testing (utilizing a gray-box approach), and User Acceptance Testing (utilizing a black-box approach). Structured test case matrices recorded the function under test, input parameters, expected behavior, actual results, and overall pass/fail status.

3.10 Full Current System Screenshots

This section documents the current Toothwise system as implemented, one screen per numbered figure, in the order a user would encounter them: the public landing page, staff login, the dashboard, patient management, the screening workflow and results, dentist management, the audit log, and the patient portal.

Captured screenshots require a running instance of the system, a live PostgreSQL database with representative seeded data, and the FastAPI backend serving the React frontend, which is outside what can be produced from the static project files alone.

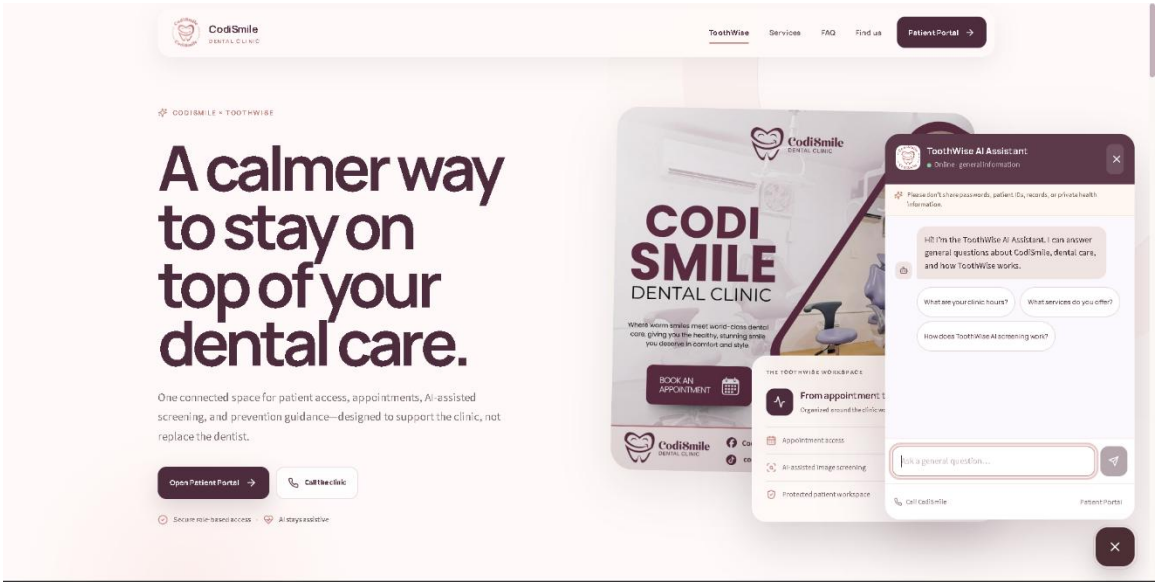


Figure 7 Public Landing Page. *This is the first page guests see when they visit the site. It introduces what the platform does and gives new users a quick way to sign up or log in.*

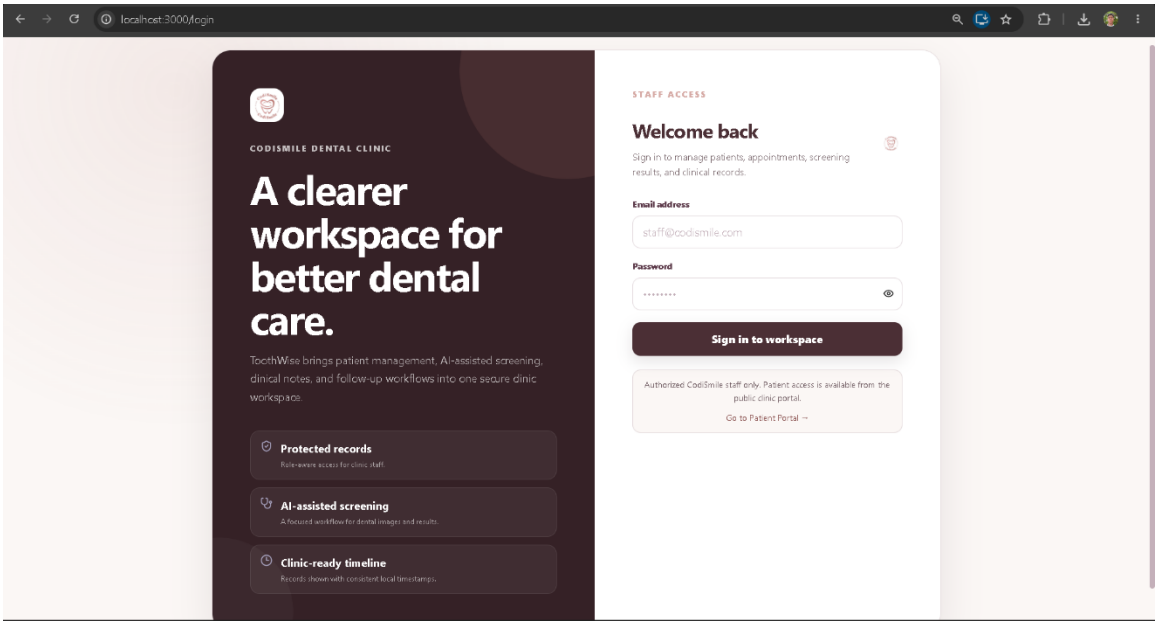


Figure 8 Staff Login Page. *This is the login portal reserved exclusively for staff. It includes inputting an email address along with a password for a secure login.*

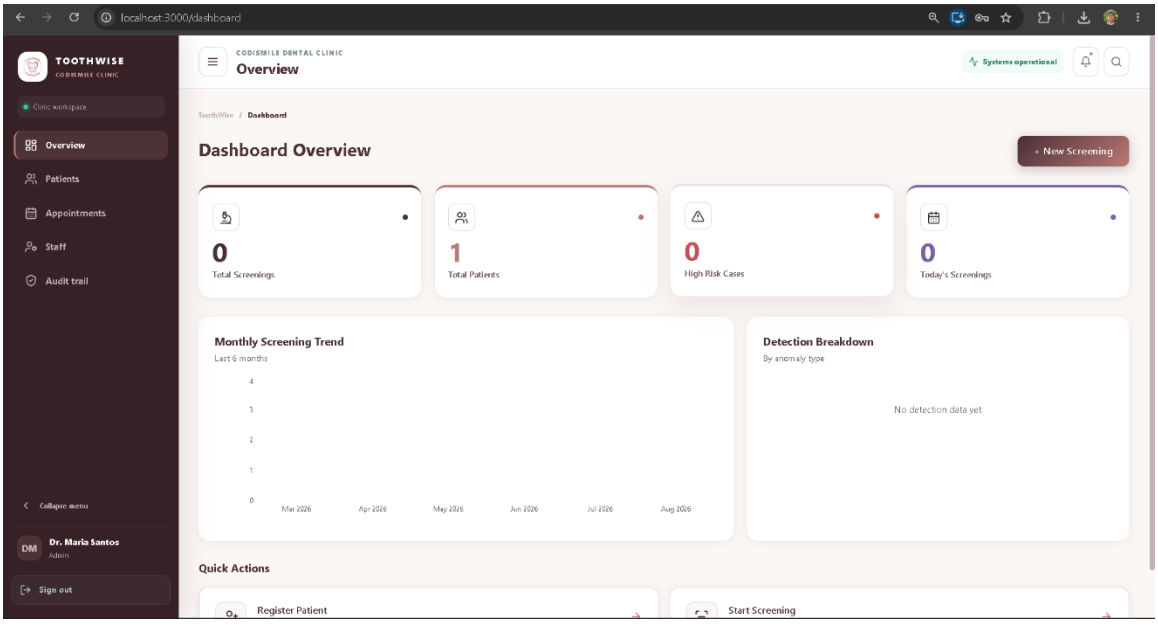


Figure 9 Dashboard. *This is the main screen staff members see right after logging in. It gives them quick summary screenings of patients and cases as well as monthly reports for the system.*

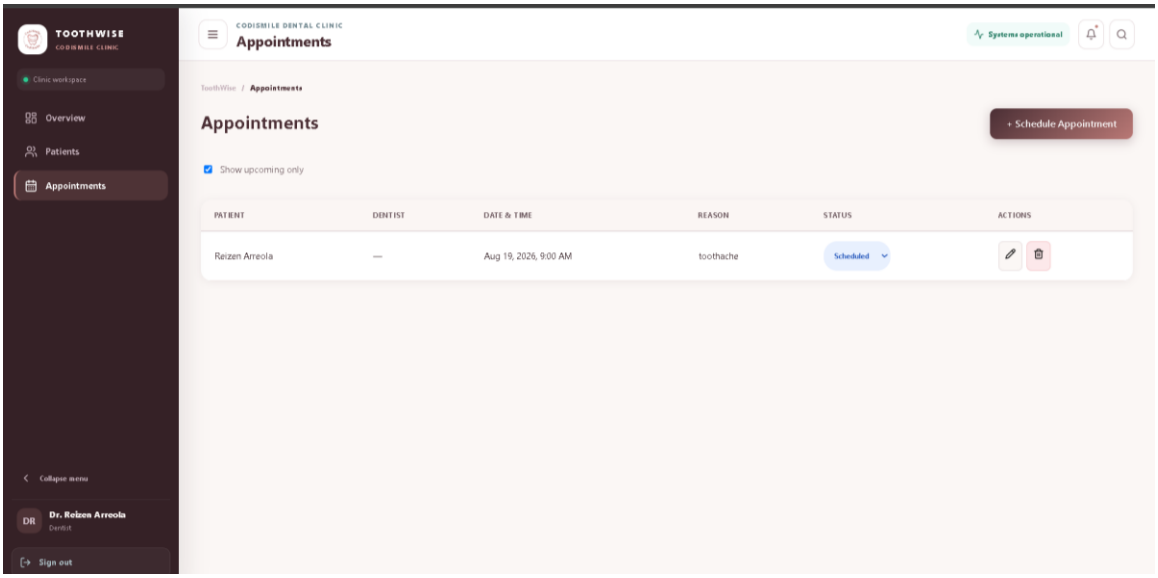


Figure 10 Patient List. *Displays all registered patient in the system. Let staff quickly view basic details, and click on a patient to open their full medical profile. Staff may also register patients manually.*

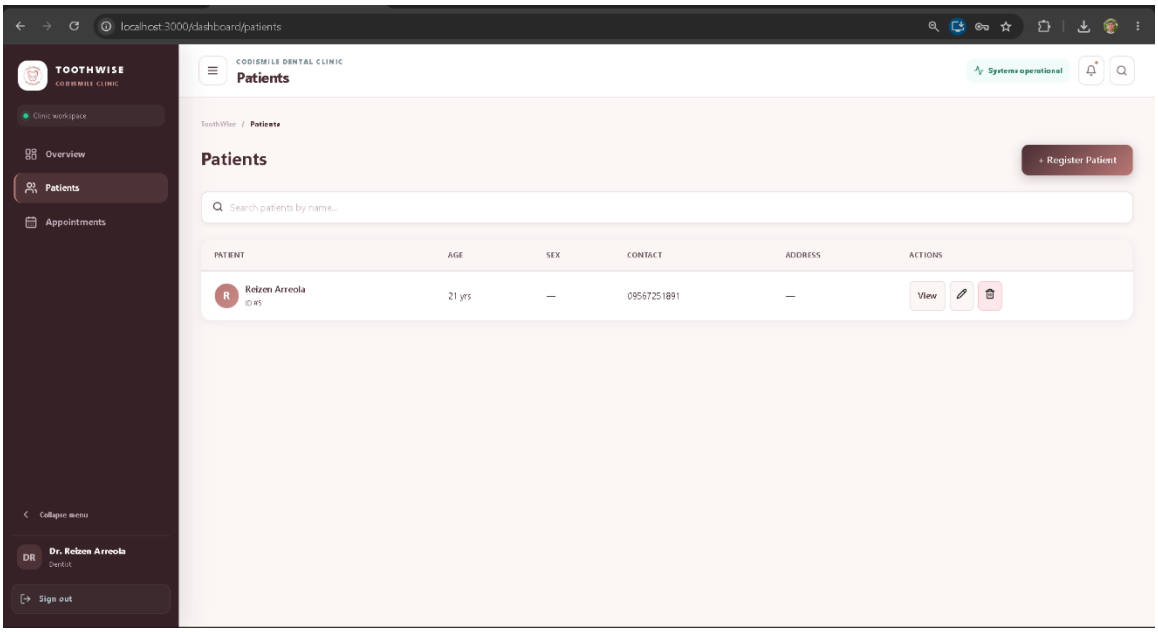


Figure 11 Patient Details. *This screen shows the complete record for an individual patient. It organizes their medical history, current treatment plans, and personal contact information all in one place.*

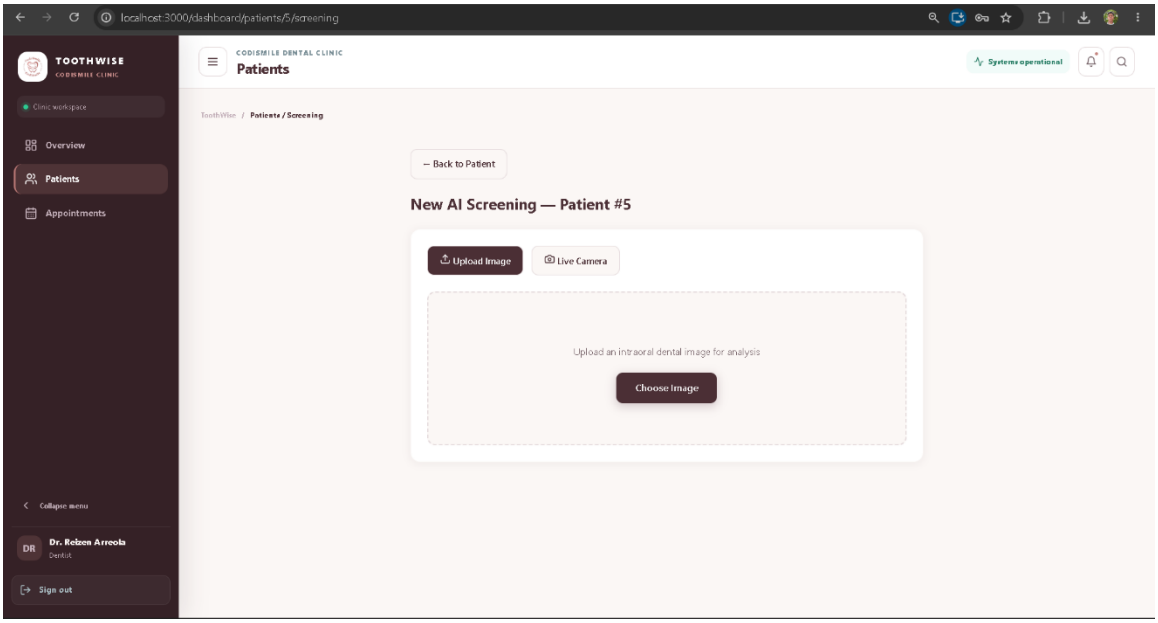


Figure 12 New Screening (Capture/Upload). *This screen allows staff to capture a photo live using the system's built-in camera feature or upload an existing file directly from their device.*

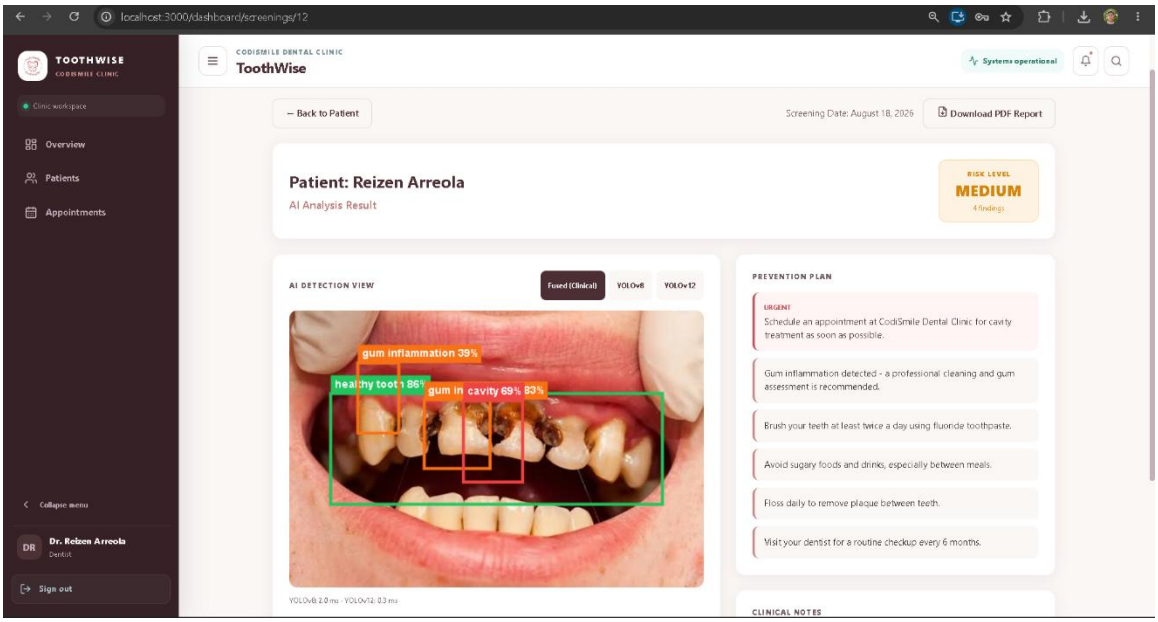


Figure 13 Screening Results. *This screen displays the system's analysis of the uploaded image, highlighting any detected abnormalities, calculating an overall risk level, and providing a recommended prevention plan.*

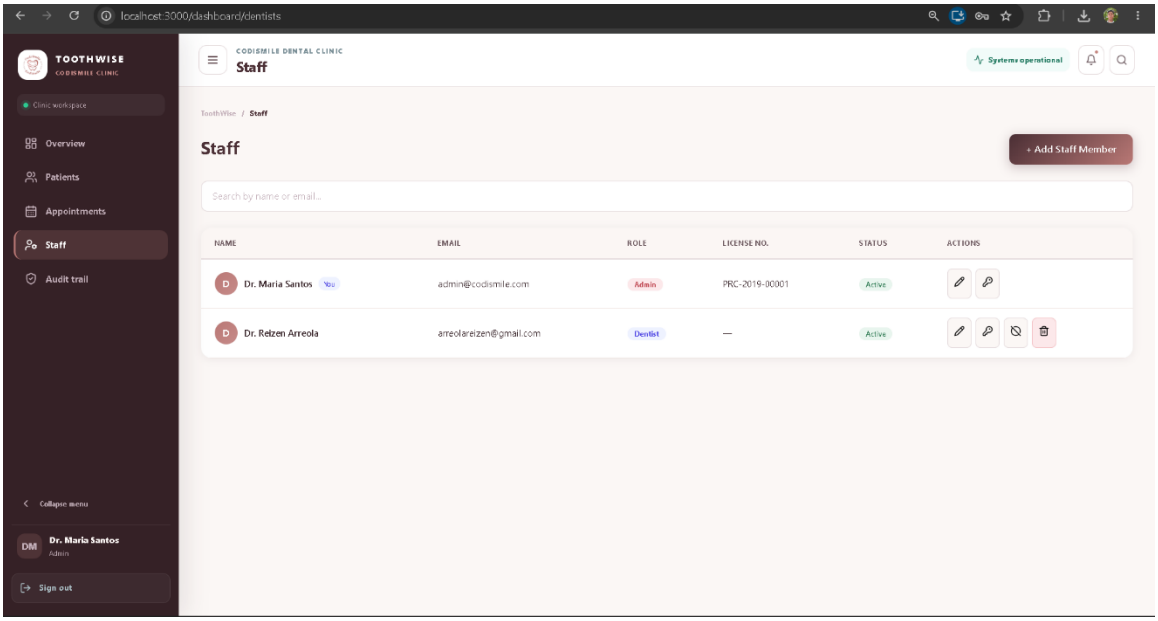


Figure 14 Staff Management (Admin). *This dashboard lets system administrators manage the workforce by adding new employees, updating staff details, and setting specific access permissions.*

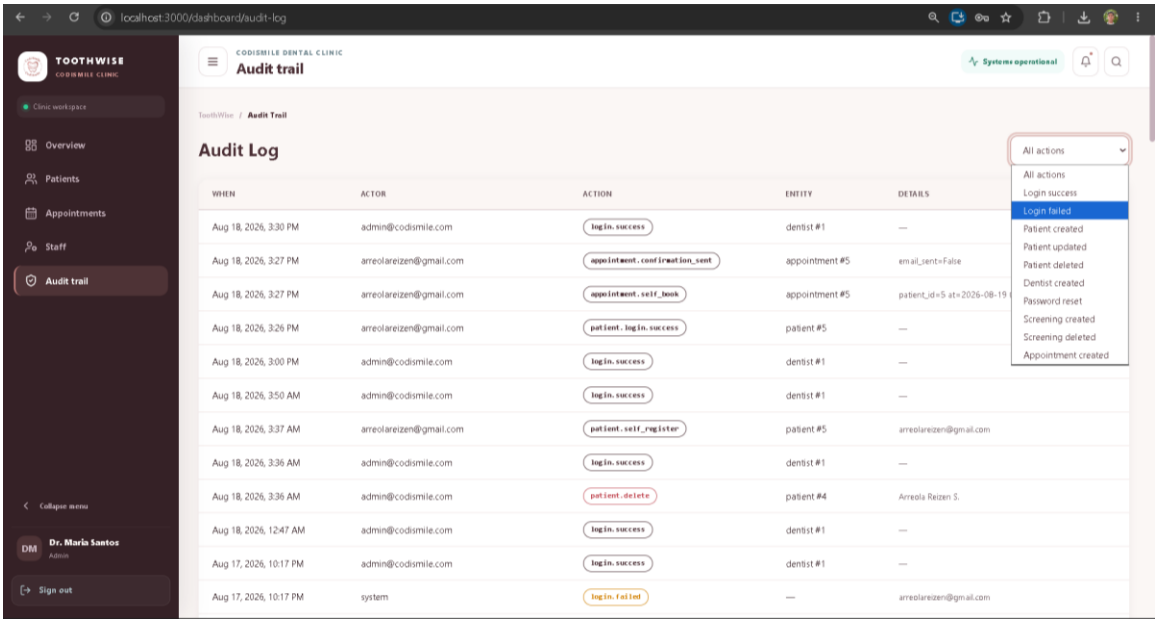


Figure 15 Audit Log (Admin). *This administrative panel tracks all user actions across the system, creating a permanent timeline of who logged in, what data was changed, and when it happened.*

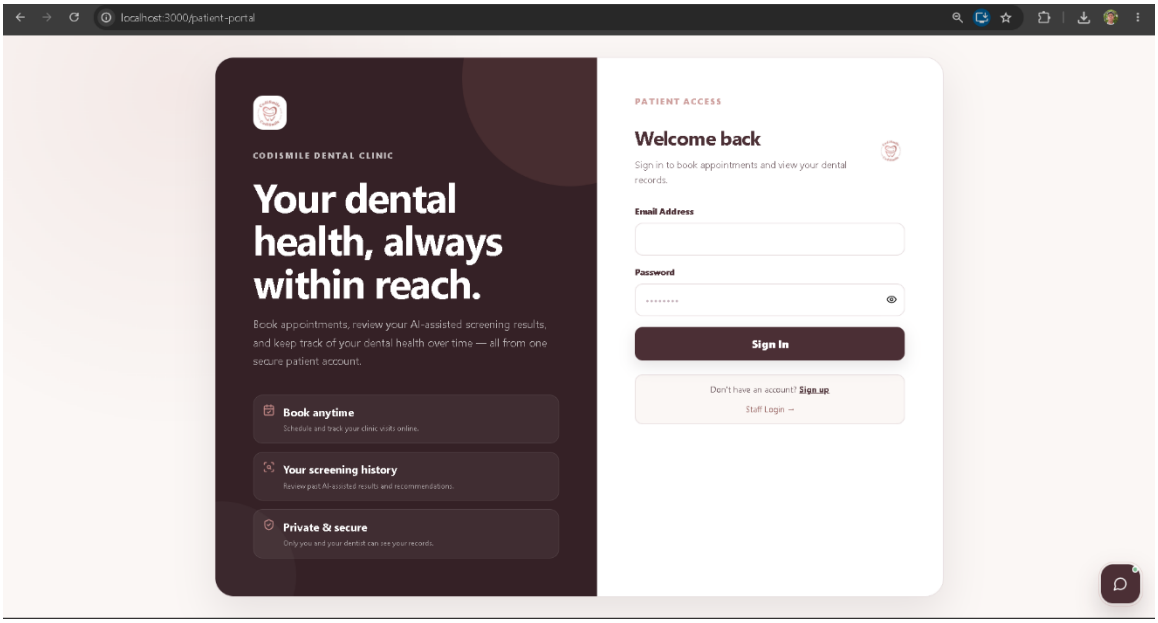


Figure 16 Patient Portal—Login. *This is the login portal for patients, allowing them to securely access their own accounts using their personal credentials.*

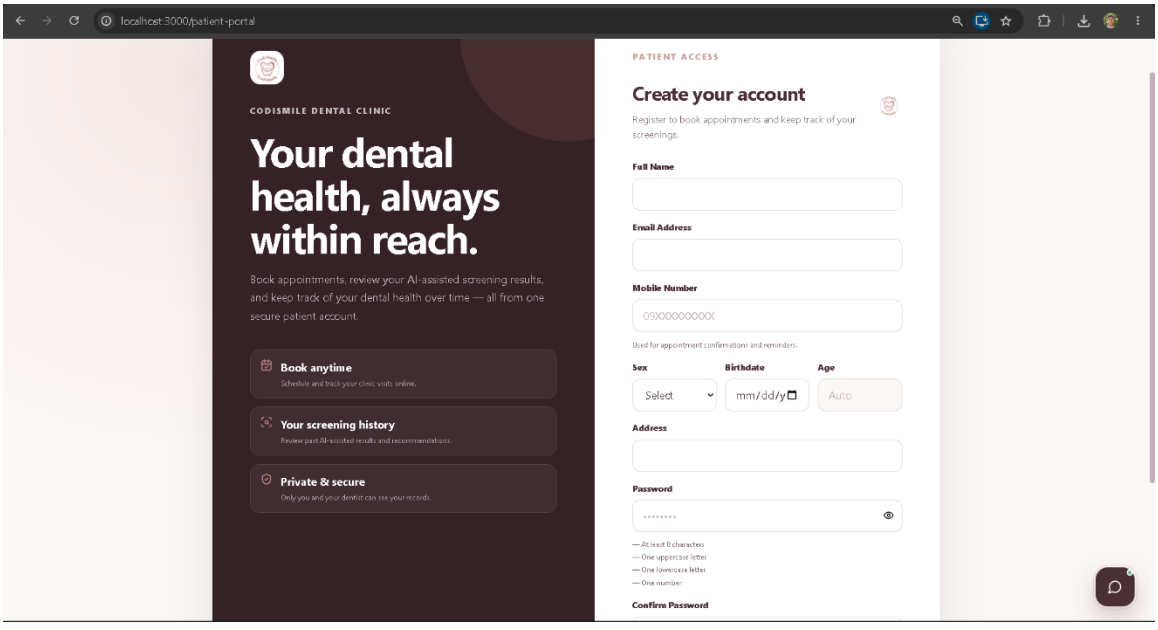


Figure 17 Patient Portal — Register. *This page allows new patients to create an account by filling out their basic personal information, contact details, and password.*

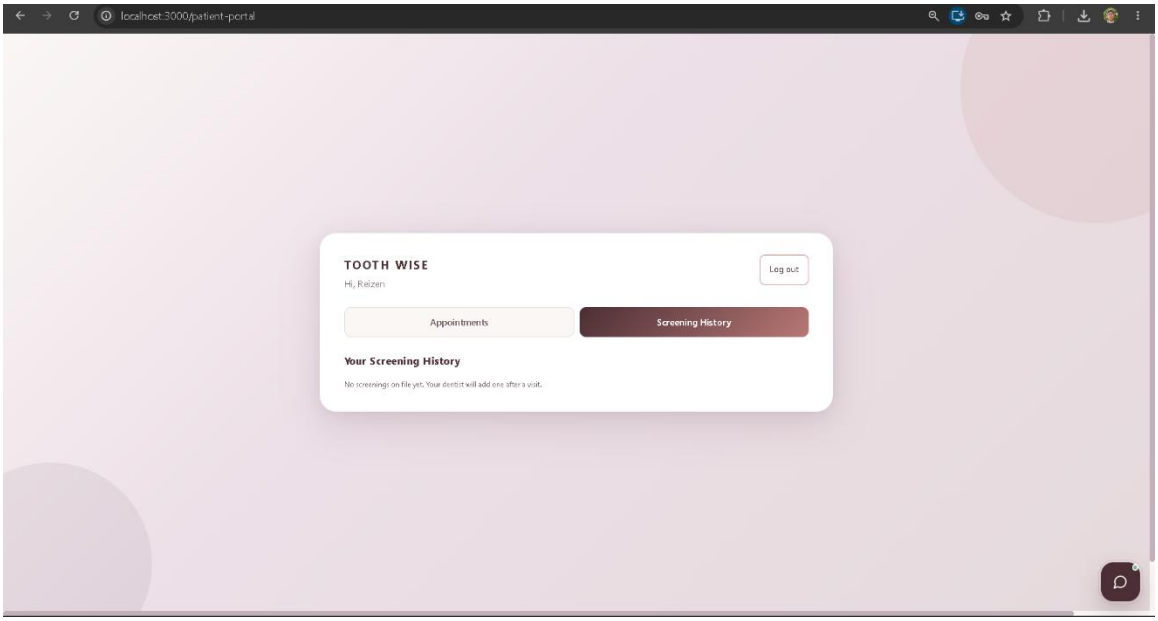


Figure 18 Patient Portal — Screening History. *This screen gives patients a private dashboard where they can see a complete list of their past evaluations, tracking their previous results and dates over time.*

REFERENCES

Alibudbud, R. C., et al. (2025). Guiding responsible AI in healthcare in the Philippines. npj Digital Medicine. https://doi.org/10.1038/s41746-025-01755-3

Askar, H., Krois, J., Gilen, M., Buttner, M., Paris, S., & Schwendicke, F. (2021). Detecting white spot lesions on dental photography using deep learning: A pilot study. Journal of Dentistry, 107, 103615. https://doi.org/10.1016/j.jdent.2021.103615

Jocher, G., Chaurasia, A., & Qiu, J. (2023). Ultralytics YOLO (Version 8.0.0) [Software]. Ultralytics. https://github.com/ultralytics/ultralytics

Kang, S., Shon, B., Park, E. Y., Jeong, S., & Kim, E. K. (2024). Diagnostic accuracy of dental caries detection using ensemble techniques in deep learning with intraoral camera images. PLOS ONE, 19(9), e0310004. https://doi.org/10.1371/journal.pone.0310004

Karakus, R., Ozic, M. U., & Tassoker, M. (2024). AI assisted detection of interproximal, occlusal, and secondary caries on bite wing radiographs using YOLOv8. Journal of Imaging Informatics in Medicine, 37(6), 3146 to 3159. https://doi.org/10.1007/s10278-024-01113-x

Mohammad Rahimi, H., Nadimi, M., Ghalyanchi Langeroudi, A., Taheri, M., & Ghafouri Fard, S. (2022). Application of deep learning in diagnosis of dental caries: A systematic review. Journal of Dentistry, 122, 104115. https://doi.org/10.1016/j.jdent.2021.104115

Moharrami, M., Schmidt, J. C., Amoo, C. F., & Mohd Salleh, N. (2024). Detecting dental caries on oral photographs using artificial intelligence: A systematic review. Oral Diseases, 30(4), 1765 to 1783. https://doi.org/10.1111/odi.14496

Oliveros Villarico, M., Flores, K. K. R., de Guzman, J. C. Z., Manrique, A. L. V., Millo, G. A., & Estrera, K. R. P. (2023). Pediatrician knowledge, attitude and practices on children's oral health in a tertiary public hospital in the Philippines. Acta Medica Philippina, 57(7), 56 to 63. https://doi.org/10.47895/amp.vi0.5323

Park, E. Y., Cho, H., Kang, S., Jeong, S., & Kim, E. K. (2022). Caries detection with tooth surface segmentation on intraoral photographic images using deep learning. BMC Oral Health, 22(1), 573. https://doi.org/10.1186/s12903-022-02589-1

PhilHealth and Department of Health. (2025). Cost effectiveness analysis of oral health care package of services within a comprehensive PhilHealth benefit package. PMC. https://pmc.ncbi.nlm.nih.gov/articles/PMC12572879/

Schwendicke, F., Golla, T., Dreher, M., & Krois, J. (2021). Convolutional neural networks for dental image diagnostics: A systematic review. Journal of Dentistry, 91, 103226. https://doi.org/10.1016/j.jdent.2019.103226

The Journal of Evidence Based Dental Practice. (2025). Variability in manual dental caries detection: Observer consistency in clinical screenings. The Journal of Evidence Based Dental Practice, 25(1), 1 to 12. https://doi.org/10.1016/j.jebdp.2024.101950

Tian, Y., Ye, Q., & Doermann, D. (2025). YOLOv12: Attention centric real time object detectors. arXiv preprint arXiv:2502.12524. https://arxiv.org/abs/2502.12524

University of Santo Tomas. (2023). Bridging the gap or widening the divide: A call for capacity building in artificial intelligence for healthcare in the Philippines. Journal of Medicine of the University of Santo Tomas, 7. https://doi.org/10.35460/2546-1621.2023-0081

World Health Organization. (2022). Oral health Philippines country profile. WHO. https://www.who.int/publications/m/item/oral-health-phl-2022-country-profile

World Health Organization. (2024). Oral health: Key facts. WHO. https://www.who.int/news-room/fact-sheets/detail/oral-health

TOOTHWISE

A YOLO Based Multi Condition Dental Screening System with KNN Based Recommendation for CodiSmile Dental Clinic

CHAPTER IV AND CHAPTER V

Results Discussion Summary Conclusions and Recommendations

Capstone Project Documentation

CHAPTER IV

RESULTS AND DISCUSSION

This chapter presents the verified technical results of the ToothWise prototype as of the current development stage. It reports results that can be reproduced from the available source code, model files, and evaluation harnesses. The ISO/IEC 25010 user evaluation is intentionally reported as pending because the planned 100 respondents have not yet completed the survey. Accordingly, this chapter does not claim user-acceptability scores, clinical accuracy, consultation-time reduction, or patient outcomes that have not been measured.

4.1 Scope of the Reported Results

The system combines a dual YOLO image-screening pipeline, an auditable risk-scoring module, a K-Nearest Neighbors (KNN) recommendation component, authenticated patient and dentist workflows, and a web-based interface. The project directory contains trained YOLOv8 and YOLOv12 weight files with saved validation metrics. The retained metrics support a model-to-model comparison, but the test images, annotation files, confusion matrices, and per-class results are not available in the project directory. Therefore, the results are reported as saved validation performance rather than as an independent clinical diagnostic-accuracy study.

Evidence source	Verified result	Interpretation
YOLO model assets	yolov8_dental.pt and yolov12_dental.pt are present.	The system can be configured for real-model inference.
Supported labels	5 configured labels: healthy tooth, cavity, gum inflammation, plaque, and tartar.	The class order matches both trained checkpoints and the application label map.

Evidence source	Verified result	Interpretation
KNN evaluation harness	Synthetic holdout and clinician-confirmed case evaluation are implemented.	The recommendation tier can be checked separately from image detection.
ISO/IEC 25010 instrument	A survey instrument and response template are prepared.	Acceptability results remain pending until 100 completed responses are collected.

4.2 Trained Image Detection Models and Dataset Readiness

The image-analysis component uses two object-detection checkpoints: YOLOv8 and YOLOv12. Both models are intended to process the same intraoral image. Their detections are then compared and fused by label and intersection-over-union (IoU). When matching boxes from both models overlap at or above the configured threshold, the system produces an ensemble output. This architecture allows the application to preserve the separate outputs while presenting a single fused result to the screening workflow.

Model artifact	File size	Configured role	Reported metric status
YOLOv8 dental checkpoint	52,040,210 bytes	First detector in the ensemble	Precision 53.96%; recall 49.73%; mAP50 46.83%; mAP50-95 26.80%
YOLOv12 dental checkpoint	39,817,646 bytes	Second detector in the ensemble	Precision 49.43%; recall 49.05%; mAP50 45.02%; mAP50-95 25.87%
Ensemble fusion	Software implementation	Fuses same-label overlapping boxes; configured IoU threshold is 0.50	Requires evaluation against labeled held-out images

Both checkpoints were trained for 60 epochs at an image size of 640 using the same five-class dataset configuration. YOLOv8 achieved the stronger saved validation values across all four aggregate metrics: it exceeded YOLOv12 by 4.53 percentage points in precision, 0.67 in recall, 1.81 in mAP50, and 0.93 in mAP50-95. These saved values demonstrate comparative performance on the validation split used during training. They do not establish clinical generalizability or replace professional judgment. A valid clinical-performance claim still requires an independent patient-disjoint labeled test set.

4.3 Comparison of the Saved Validation Results

The model checkpoints retain the aggregate validation results produced during training. Table 2 compares the two models. Precision reflects the proportion of predicted boxes that are correct, recall reflects the proportion of labeled boxes detected, mAP50 measures average precision at an IoU threshold of 0.50, and mAP50-95 evaluates detection quality across stricter IoU thresholds from 0.50 to 0.95.

Model	Precision	Recall	mAP50	mAP50-95
YOLOv8	53.96%	49.73%	46.83%	26.80%
YOLOv12	49.43%	49.05%	45.02%	25.87%

YOLOv8 is the stronger individual checkpoint under the retained validation results. The advantage is modest in recall, mAP50, and mAP50-95, but it is consistent across all reported aggregate measures. The

retained files do not include the validation image count, class distribution, per-class average precision, or confusion matrix. A fresh independent evaluation should therefore be completed on a patient-disjoint held-out test set before making a clinical accuracy claim or selecting a permanent deployment threshold.

4.4 Risk Scoring and Recommendation Results

ToothWise converts detected findings into a transparent risk score. Healthy-only findings produce a low-risk score of zero. For non-healthy findings, the highest confidence-weighted severity is used as the base score. A small capped increment is added for additional findings. The final score is limited to 100 and is categorized as LOW (0 to 34), MEDIUM (35 to 69), or HIGH (70 to 100). This deterministic rule makes the reason for each tier inspectable during testing and clinical review.

Risk-tier component	Verified implementation behavior
Severity calculation	Uses the confidence-weighted severity of the most serious non-healthy detection.
Multiple findings	Adds a capped six-point increment for each additional concerning finding.
Risk thresholds	LOW: 0-34; MEDIUM: 35-69; HIGH: 70-100.
Recommendation output	Prioritizes urgent advice for high risk, then condition-specific and general oral-health guidance.

4.5 KNN Recommendation Evaluation

The KNN recommender was evaluated with its built-in synthetic holdout procedure. The holdout contains 2,000 deterministically generated feature vectors created with a different random seed from the training bootstrap. The measured accuracy was 0.9275 (92.75%), while the weighted F1 score was 0.9251 (92.51%). These results show that the KNN model largely reproduces the risk-tier pattern represented in the generated data.

This result must be interpreted narrowly. Because the bootstrap and holdout data are generated from the same rule-based severity framework, the metric demonstrates internal consistency rather than clinical effectiveness. At the time of evaluation, the clinician-confirmed case journal contained zero real labeled cases. Therefore, no claim can be made that the recommender accurately predicts risk tiers in a real patient population.

Evaluation measure	Observed result	Appropriate interpretation
Synthetic holdout size	2,000 cases	A reproducible software-level evaluation set
Accuracy	0.9275 or 92.75%	Agreement with the generated tier pattern
Weighted F1 score	0.9251 or 92.51%	Balanced synthetic classification performance
Real clinician-confirmed cases	0 cases	Clinical validation has not started

4.6 Functional Verification of the System

The backend test suite contains 73 collected automated tests covering authentication, role and token scoping, dentist and patient records, screening creation and retrieval, label normalization, risk scoring, KNN recommendation behaviors, and selected security conditions. The suite provides useful regression coverage for the system functions that can be tested without a live clinical study. It should be run in the deployment environment before final demonstration, with the final pass/fail summary recorded in the manuscript.

Area covered	Examples of verified behaviors
Authentication and access control	Login, incorrect credentials, token requirements, and separation of dentist and patient access.
Record management	Create, retrieve, update, search, and delete patient and dentist records.
Screening workflow	Create and retrieve screenings, store detections and recommendations, and add clinical notes.
Decision logic	Risk boundaries, score cap, label aliases, recommendation ordering, and tier confirmation.

4.7 ISO IEC 25010 Survey Status

The planned ISO/IEC 25010 evaluation will use 100 respondents and a five-point Likert-scale instrument. Until the questionnaires are completed and encoded, weighted means, standard deviations, verbal interpretations, and an overall acceptability rating cannot be reported. Stating a survey result before data collection would undermine the validity of the study.

When the survey is completed, the analysis should report the respondent profile; the number of valid responses; the mean and standard deviation for each quality characteristic; the overall weighted mean; the interpretation scale; and selected qualitative comments. The results should distinguish software quality perceptions from clinical diagnostic effectiveness, which requires a separate validated clinical dataset.

4.8 Discussion

The present evidence supports the conclusion that ToothWise is an implemented, testable prototype with trained image-model artifacts, an auditable risk engine, and a KNN-supported recommendation workflow. The saved validation results show that YOLOv8 is the stronger individual model on the available aggregate measures, while YOLOv12 provides a second independent detection output for the deployed ensemble. The KNN synthetic holdout result is useful for verifying recommendation-pipeline consistency. Neither result substitutes for an independent clinical validation study.

This distinction is important for responsible deployment. ToothWise should be presented as a clinical decision-support and screening system. A licensed dentist remains responsible for assessment, diagnosis, and treatment decisions. The next phase of the study should supply the missing empirical evidence: a locked, labeled detection test set; clinician-confirmed recommendation outcomes; and the 100-user quality survey.

CHAPTER V

SUMMARY CONCLUSIONS AND RECOMMENDATIONS

5.1 Summary of Findings

The development of ToothWise produced a web-based multi-condition dental screening prototype for CodiSmile Dental Clinic. The system integrates two trained YOLO checkpoint files, an ensemble-fusion process, a deterministic risk engine, a KNN-based recommendation component, authenticated portals, and data-management functions for patient screening records.

The five configured screening labels are healthy tooth, cavity, gum inflammation, plaque, and tartar. The class order in the trained YOLOv8 and YOLOv12 checkpoints matches the application label map. The project has the technical capability to run real model inference when the trained checkpoint paths and runtime environment are configured. The saved validation results favor YOLOv8: precision 53.96%, recall 49.73%, mAP50 46.83%, and mAP50-95 26.80%, compared with YOLOv12 at 49.43%, 49.05%, 45.02%, and 25.87%, respectively. These figures are not independent clinical test results.

The KNN recommendation component achieved 92.75% accuracy and 92.51% weighted F1 on a 2,000-case synthetic holdout. Since the data were generated from the same rule-based framework used in the prototype, these figures show internal consistency only. No clinician-confirmed cases were available at the time of evaluation. The backend also contains 73 collected automated tests covering important functional and access-control behaviors.

The ISO/IEC 25010 evaluation instrument is prepared, but the planned 100-user survey has not yet been conducted. Consequently, this study does not report a weighted mean, acceptability rating, or stakeholder satisfaction result.

5.2 Conclusions

Based on the verified development evidence, ToothWise has achieved its prototype implementation objective. It provides an end-to-end workflow for storing screening information, processing multi-condition detections, assigning an understandable risk tier, and generating prevention guidance. Its rule-based scoring logic and automated tests make key software decisions inspectable and repeatable.

The current evidence is not sufficient to conclude that ToothWise has a particular clinical diagnostic accuracy, improves consultation time, or is highly acceptable to users. Those conclusions require measured evidence from a held-out annotated dataset, clinician review, workflow timing, and the completed ISO/IEC 25010 survey. The system should therefore be framed as a decision-support prototype rather than an autonomous diagnostic tool.

The synthetic KNN result indicates that the recommendation pipeline consistently maps generated detection features into risk tiers. Its clinical usefulness remains to be established through clinician-confirmed outcomes. The saved YOLO validation metrics support YOLOv8 as the stronger individual model, while an independent patient-disjoint test set remains necessary for a final clinical performance claim.

5.3 Recommendations

5.3.1 Complete the Model Evaluation

• Freeze a patient-disjoint, labeled test set and record its image count and class distribution.

• Run the trained YOLOv8, YOLOv12, and fused ensemble on that test set, then report precision, recall, mAP50, mAP50-95, per-class average precision, and a confusion matrix.

• Have a licensed dentist review a sample of false positives and false negatives to identify annotation ambiguity and clinically important failure modes.

5.3.2 Conduct the 100 User ISO IEC 25010 Survey

• Collect the 100 responses before assigning any acceptability rating.

• Encode the results accurately, remove duplicate or incomplete entries according to a documented rule, and report the valid response count.

• Compute the weighted mean and standard deviation for each quality characteristic and present the interpretation scale used by the study.

5.3.3 Validate in a Controlled Clinical Workflow

• Use ToothWise alongside, not in place of, professional dental assessment during a supervised pilot.

• Record real clinician-confirmed risk tiers so the KNN evaluator can measure performance on non-synthetic cases.

• Measure workflow time prospectively using defined start and end points if time-efficiency is a research objective.

5.3.4 Improve Dataset Governance and Generalizability

• Document data provenance, consent, de-identification, image quality criteria, annotation protocol, and annotator agreement.

• Expand the dataset across varied lighting, cameras, age groups, dental conditions, and clinic settings while monitoring class imbalance.

• Retain training logs, data configuration files, metrics, and confusion matrices with the final model weights for reproducibility.

5.3.5 Maintain Safety and Privacy Controls

• Keep the dentist in the clinical decision loop and display that screening outputs are decision support, not final diagnoses.

• Use strong credentials, role-based authorization, audit logs, encrypted deployment configuration, and controlled access to patient images.

• Perform regression testing after model, dependency, or workflow changes before deployment.

5.4 Proposed Next Steps for the Manuscript

After the survey and model evaluation are completed, Chapter IV should be updated by replacing the pending-result fields with measured values and adding the final survey tables, model-metric tables, and interpretations. Chapter V should then be revised so its conclusions correspond only to the completed evidence. This sequence will keep the manuscript defensible during oral presentation and avoid overstating the capabilities of the prototype.